

CSE 120: Principles of Operating Systems

Lecture 2: Processes

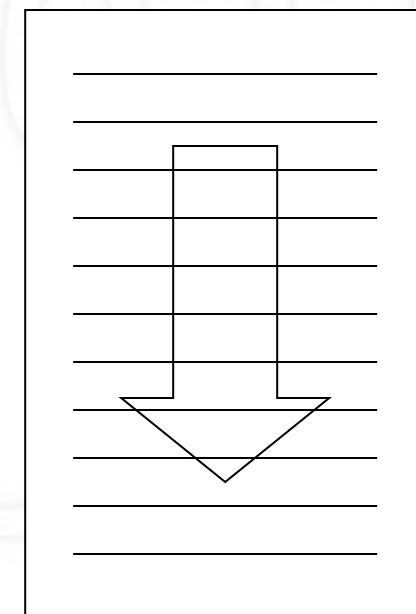
Prof. Joseph Pasquale
University of California, San Diego
January 11, 2023

Introduction

- Most basic kernel function: run a program
- Users wants ability to run multiple programs
- How to achieve given single CPU + memory?

Process

- Abstraction of a running program
 - “a program in execution”
- Dynamic: has state, changes
 - Whereas a program is static



Basic Resources for Processes

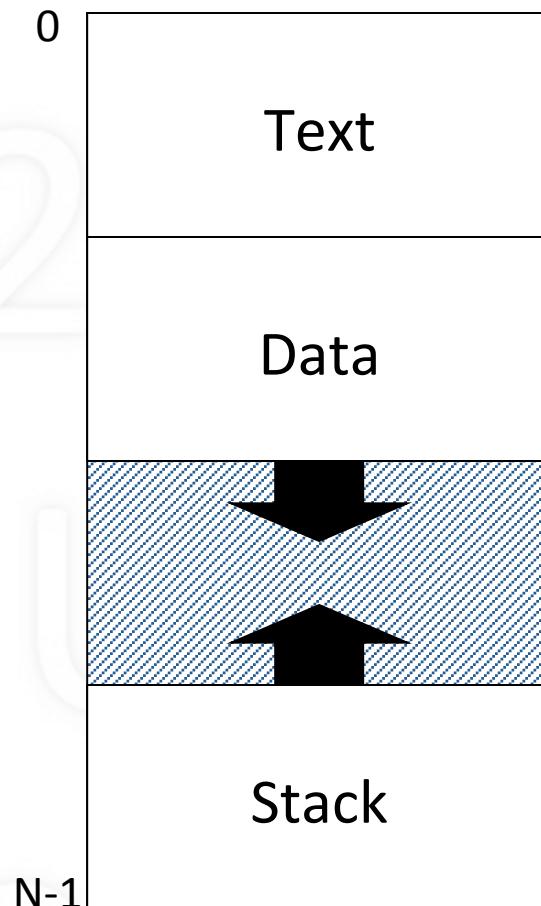
- CPU
 - Processing cycles (time)
 - To execute instructions
- Memory
 - Bytes or words (space)
 - To maintain state
- Other resources (e.g., I/O)

Context of a Process

- Context: machine and kernel-related state
- CPU context: values of registers
 - PC (program counter)
 - SP (stack pointer), FP (frame pointer), GP (general)
- Memory context: pointers to memory areas
 - Code, static variables (init, uninit), heap, shared, ...
 - Stack of activation records
- Other (kernel-related state, ...)

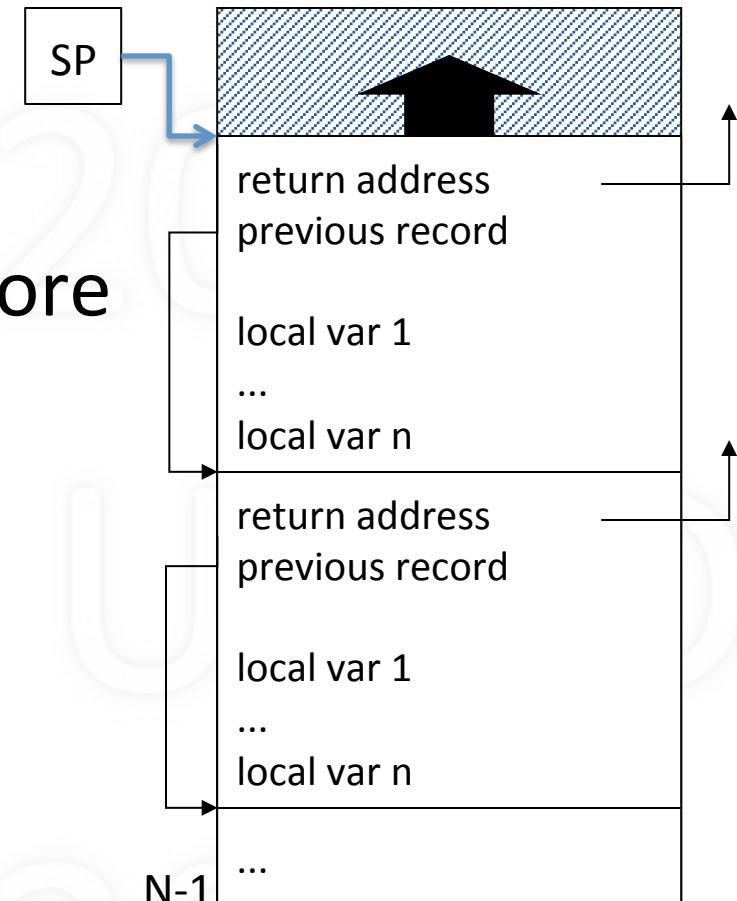
Process Memory Structure

- Text
 - Code: program instructions
- Data
 - Global variables
 - Heap (dynamic allocation)
- Stack
 - Activation records
 - Automatic growth/shrinkage

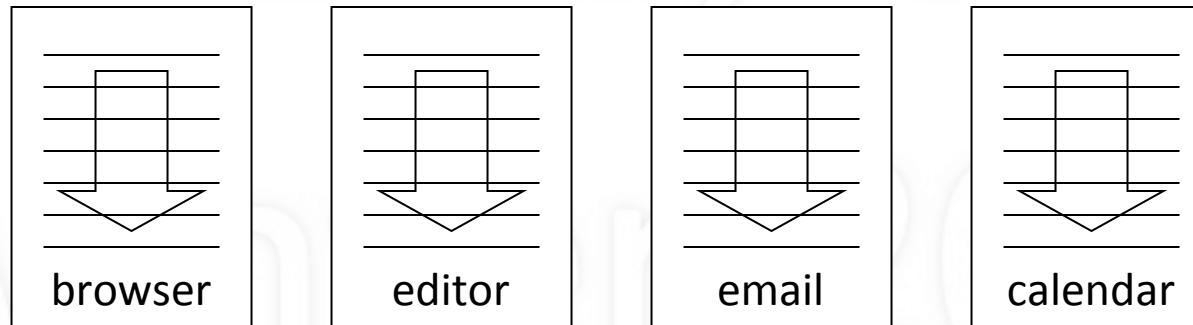


Process Stack

- Stack of activation records
 - One per pending procedure
- An activation record may store
 - where to return to
 - link to previous record
 - automatic (local) variables
 - other (e.g., register values)
- Stack pointer points to top



Goal: Support Multiple Processes



- Users would like to run multiple programs
 - “simultaneously”
 - Not all actively using the CPU
 - Some waiting for input, devices (e.g., disk), ...
- How to do this given single CPU?

Multiprogramming

- Given a running process
 - At some point, it needs a resource, e.g., I/O device
 - Say resource is busy, process can't proceed
 - So, “voluntarily” gives up CPU to another process
- Yield (p)
 - Let process p run (voluntarily give up CPU to p)
 - Requires context switching

Context Switching

- Allocating CPU from one process to another
 - First, save context of currently running process
 - Next, restore (load) context of next process to run
- Loading the context
 - Load general registers, stack pointer, etc.
 - Load program counter (must be last instruction!)

Simple Context Switching

- Two processes: A and B
- A calls Yield(B) to voluntarily give up CPU to B
- Save and restore registers
 - General-purpose, stack pointer, program counter
- Switch text and data
- Switch stacks
 - Note that PC is in the middle of Yield!

The “magic” of Yield

```
magic = 0;           // local variable
save A's context:   // current process
    asm save GP;     // general purpose registers;
    asm save SP;     // stack pointer
    asm save PC;     // program counter, note value!
if (magic == 1) return;
else magic = 1;
restore B's context: // process being yielded to
asm restore GP;
asm restore SP;
asm restore PC;     // must be last!
```

Example

```
int me;

Main () /* process A */
{
    me = Getpid ();
    Yield (B);
    Yield (A);
}

Yield (p)
{   int magic;

    magic = 0;
    SaveContext (me);
    if magic == 1 return;
    else magic = 1;
    RestoreContext (p);
}
```

```
int me;

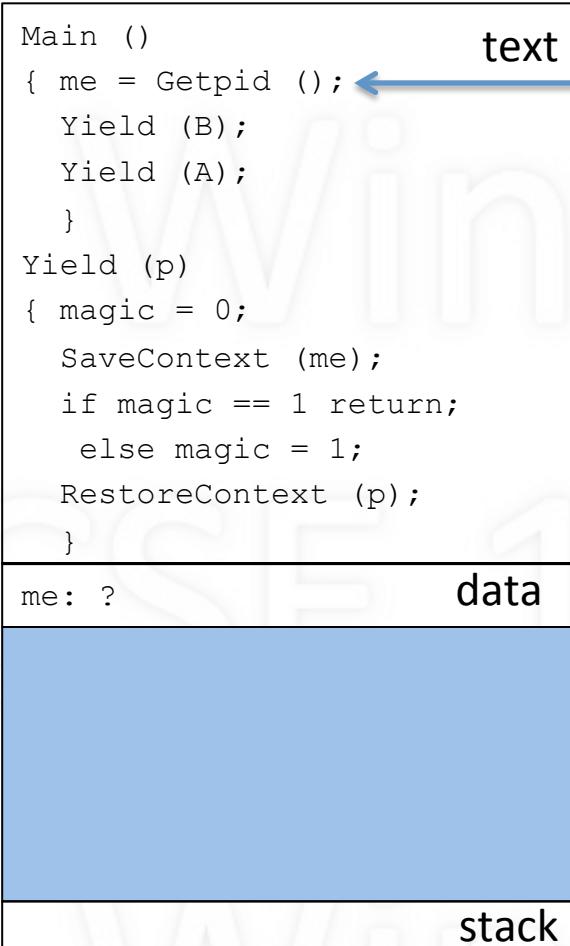
Main () /* process B */
{
    me = Getpid ();
    Yield (A);
    Yield (A);
}

Yield (p)
{   int magic;

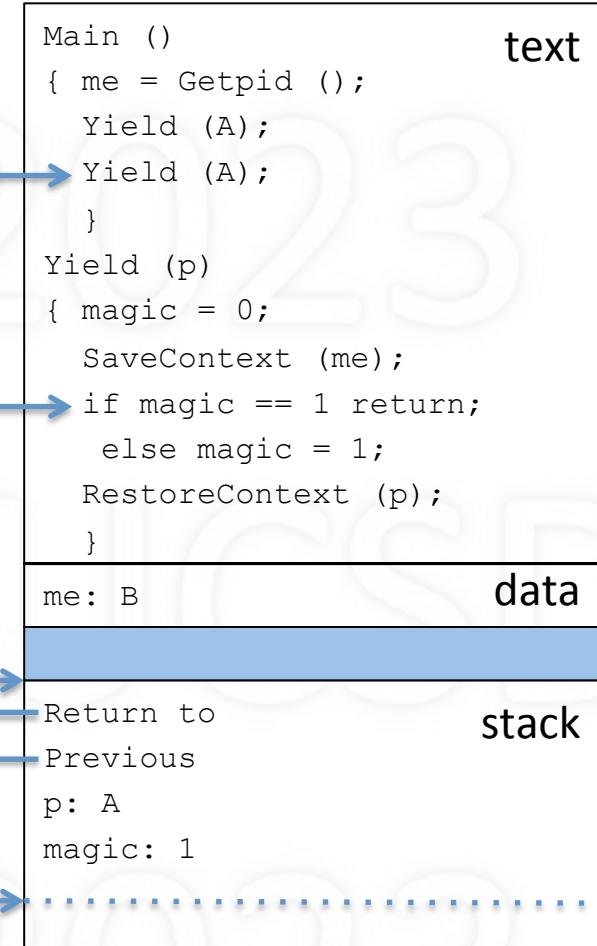
    magic = 0;
    SaveContext (me);
    if magic == 1 return;
    else magic = 1;
    RestoreContext (p);
}
```

In this example, A is about to set me to its process ID, and yield to B. B had already yielded to A: note B's saved PC and SP.

Process A

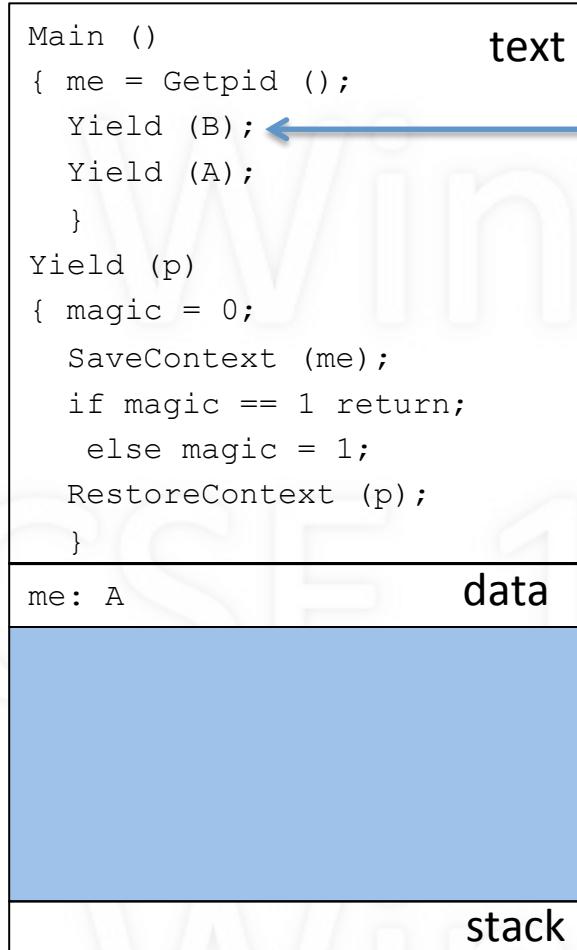


Process B



Process A has just set me to 'A' and is about to call Yield. The PC always points to the next instruction to be executed.

Process A



PC
SP

shared
memory

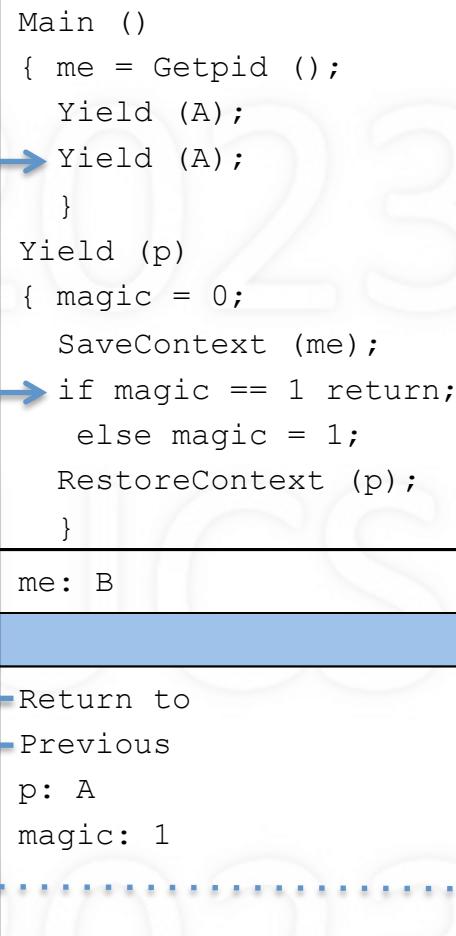
A.context

PC
SP

B.context

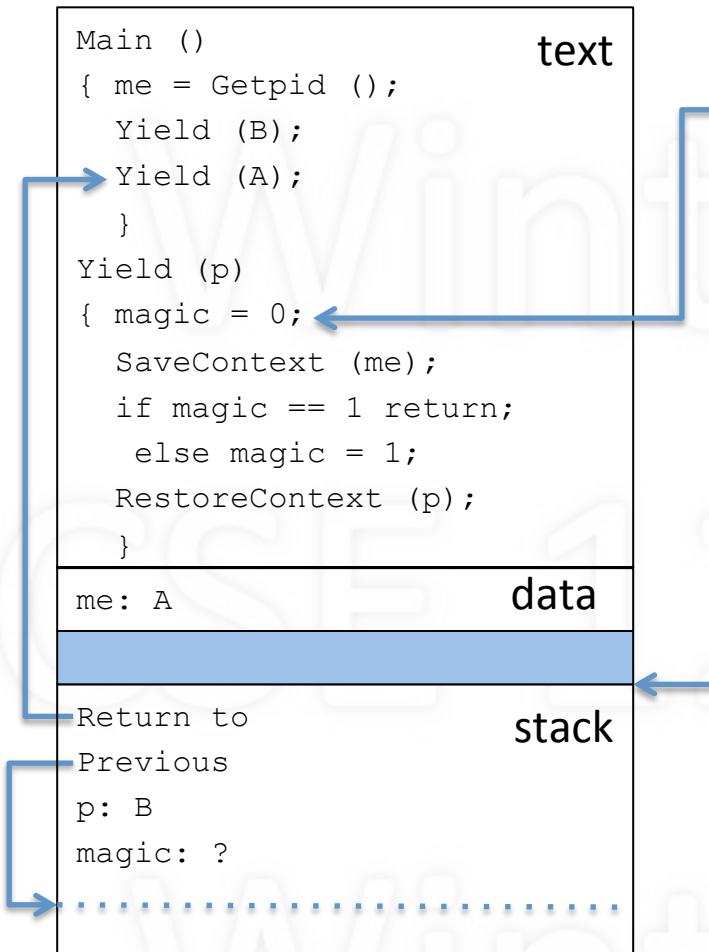
PC
SP

Process B

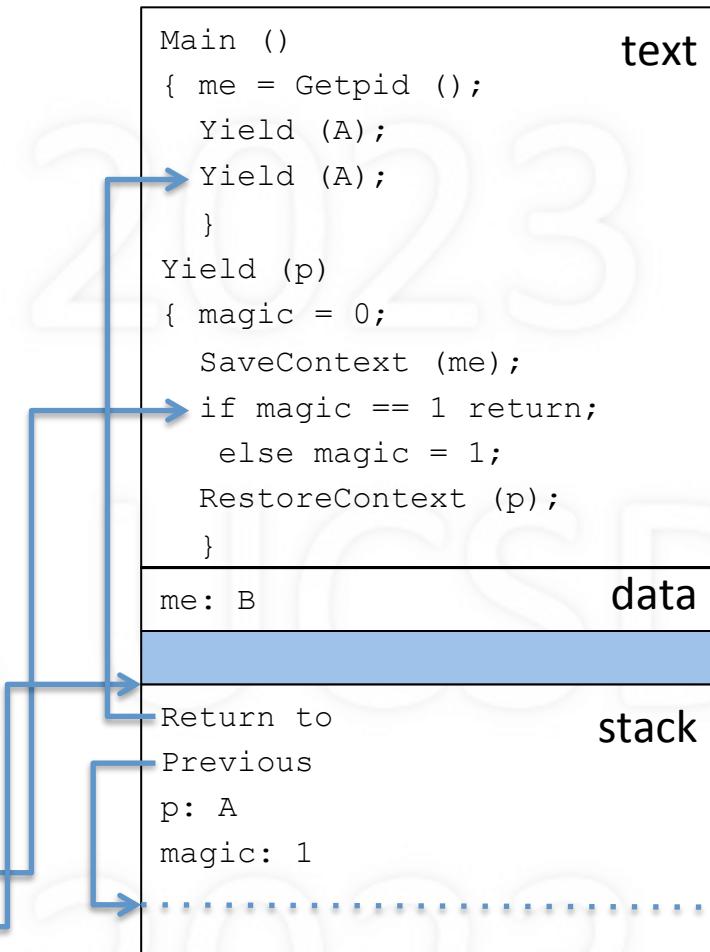


Upon entering Yield, an activation record is pushed on the stack.
It contains links, and local variables p and magic.

Process A

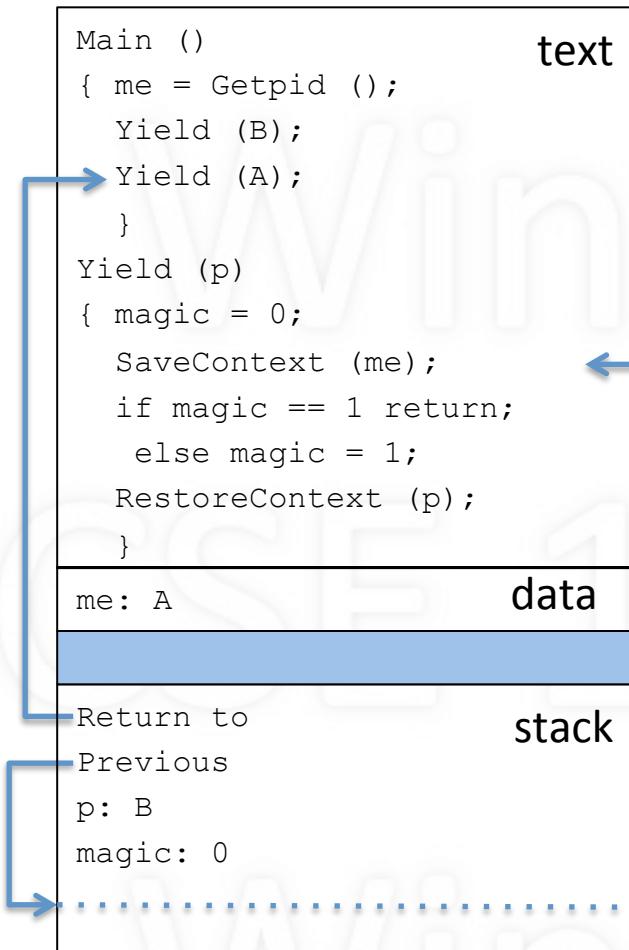


Process B

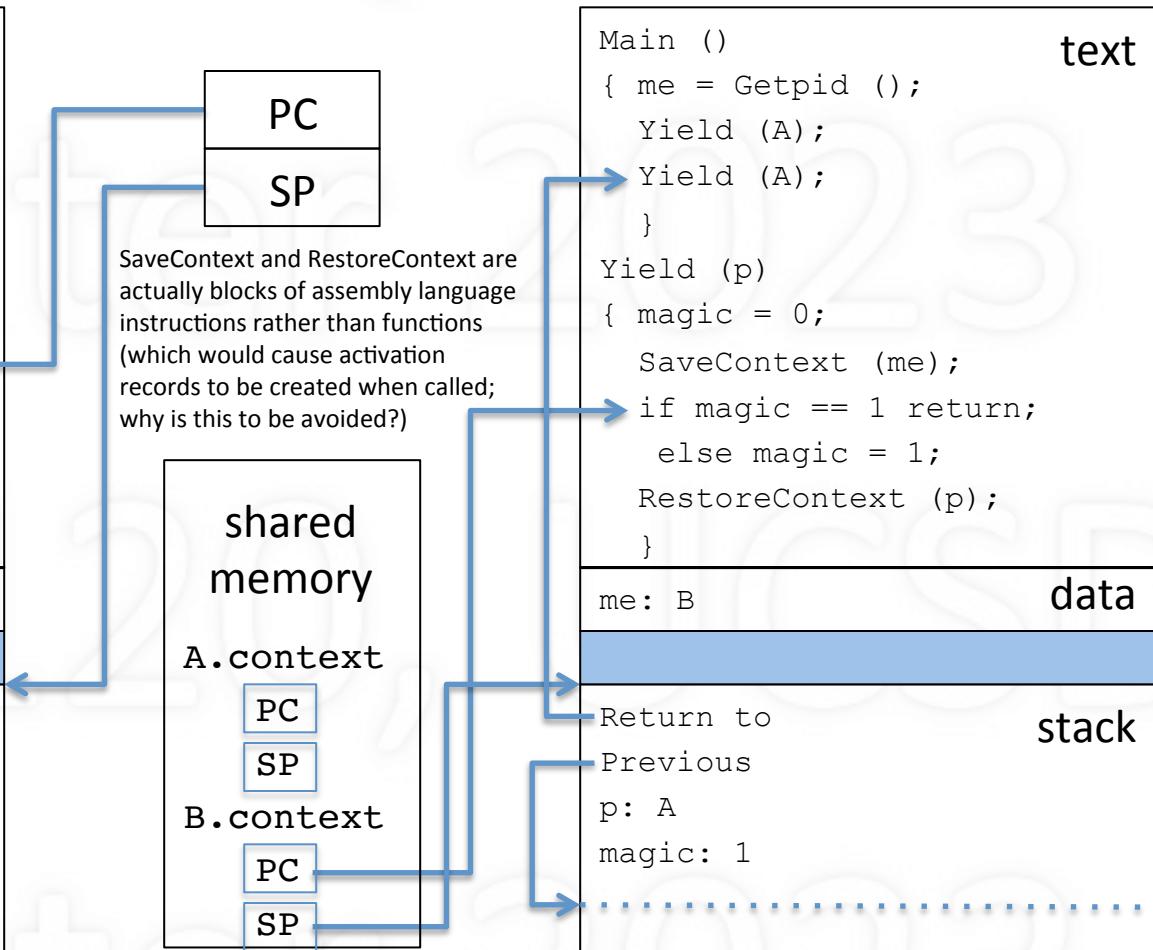


Variable magic is set to 0. It is an automatic variable, dynamically allocated on the stack. Next: SaveContext.

Process A

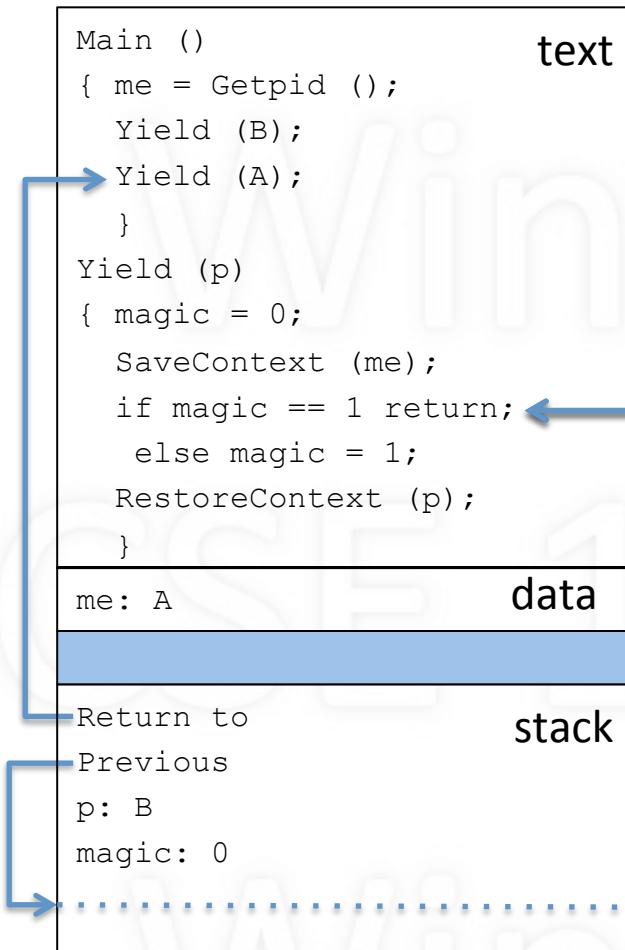


Process B

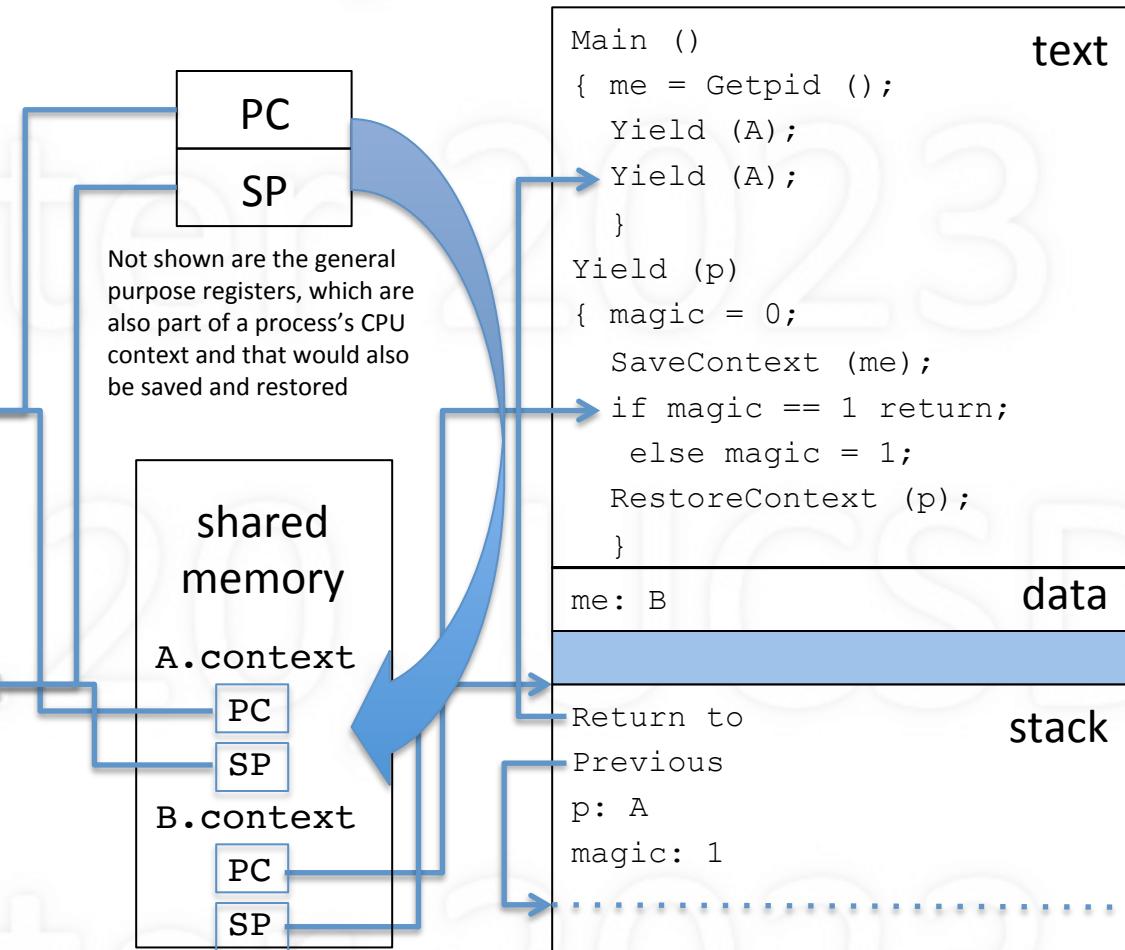


A's context is now saved. The saved PC points to just after the call to SaveContext. Compare this to B's saved context.

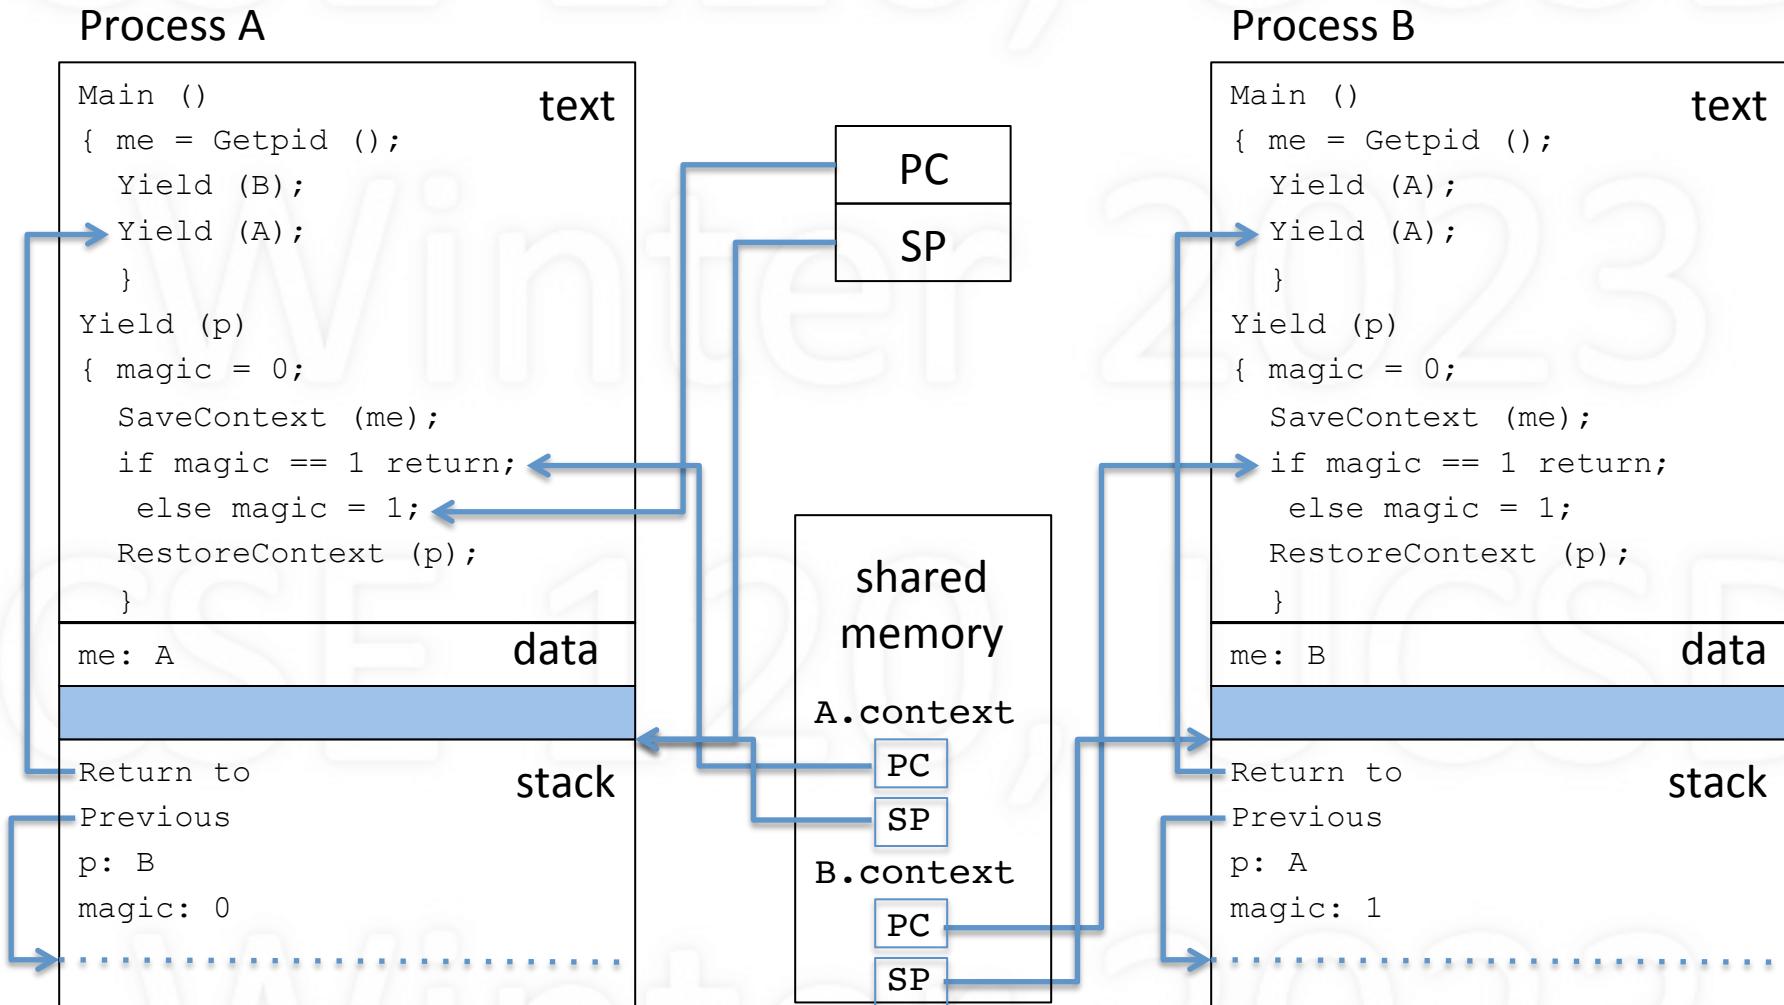
Process A



Process B

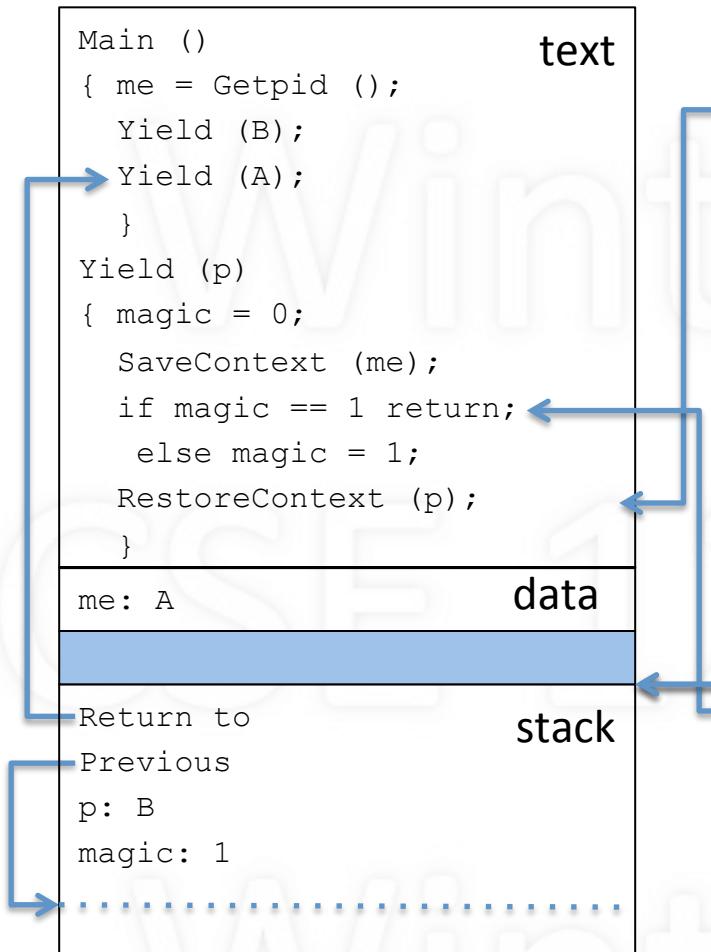


A just checked whether magic equals 1, which was false, and so, on to the else clause to set magic to 1.

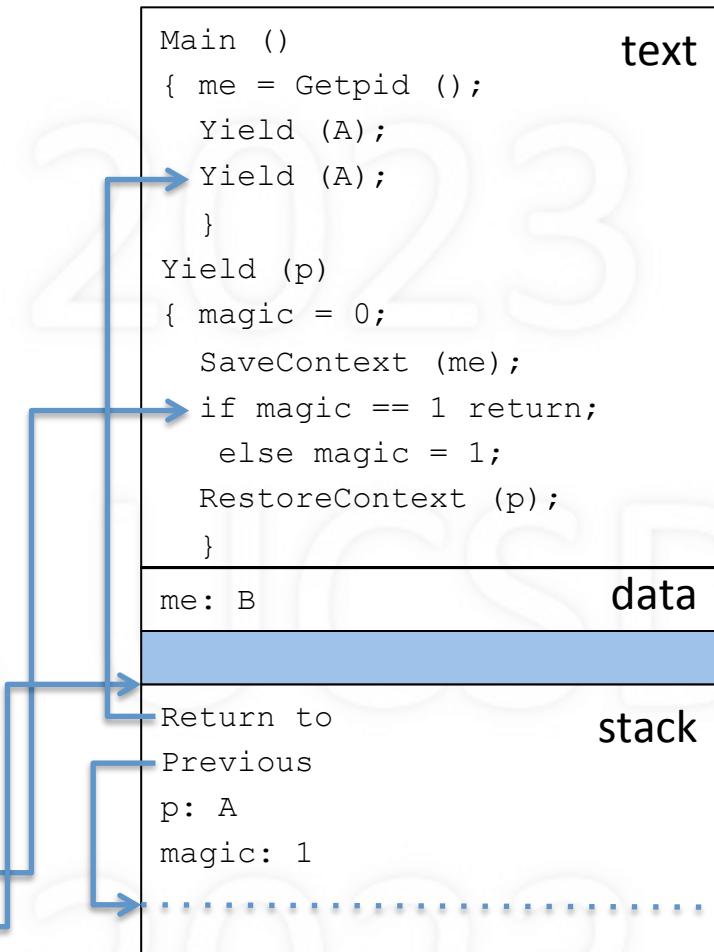


A sets magic to 1, and is about to restore B's context (just like B's situation in the past when it restored A's context).

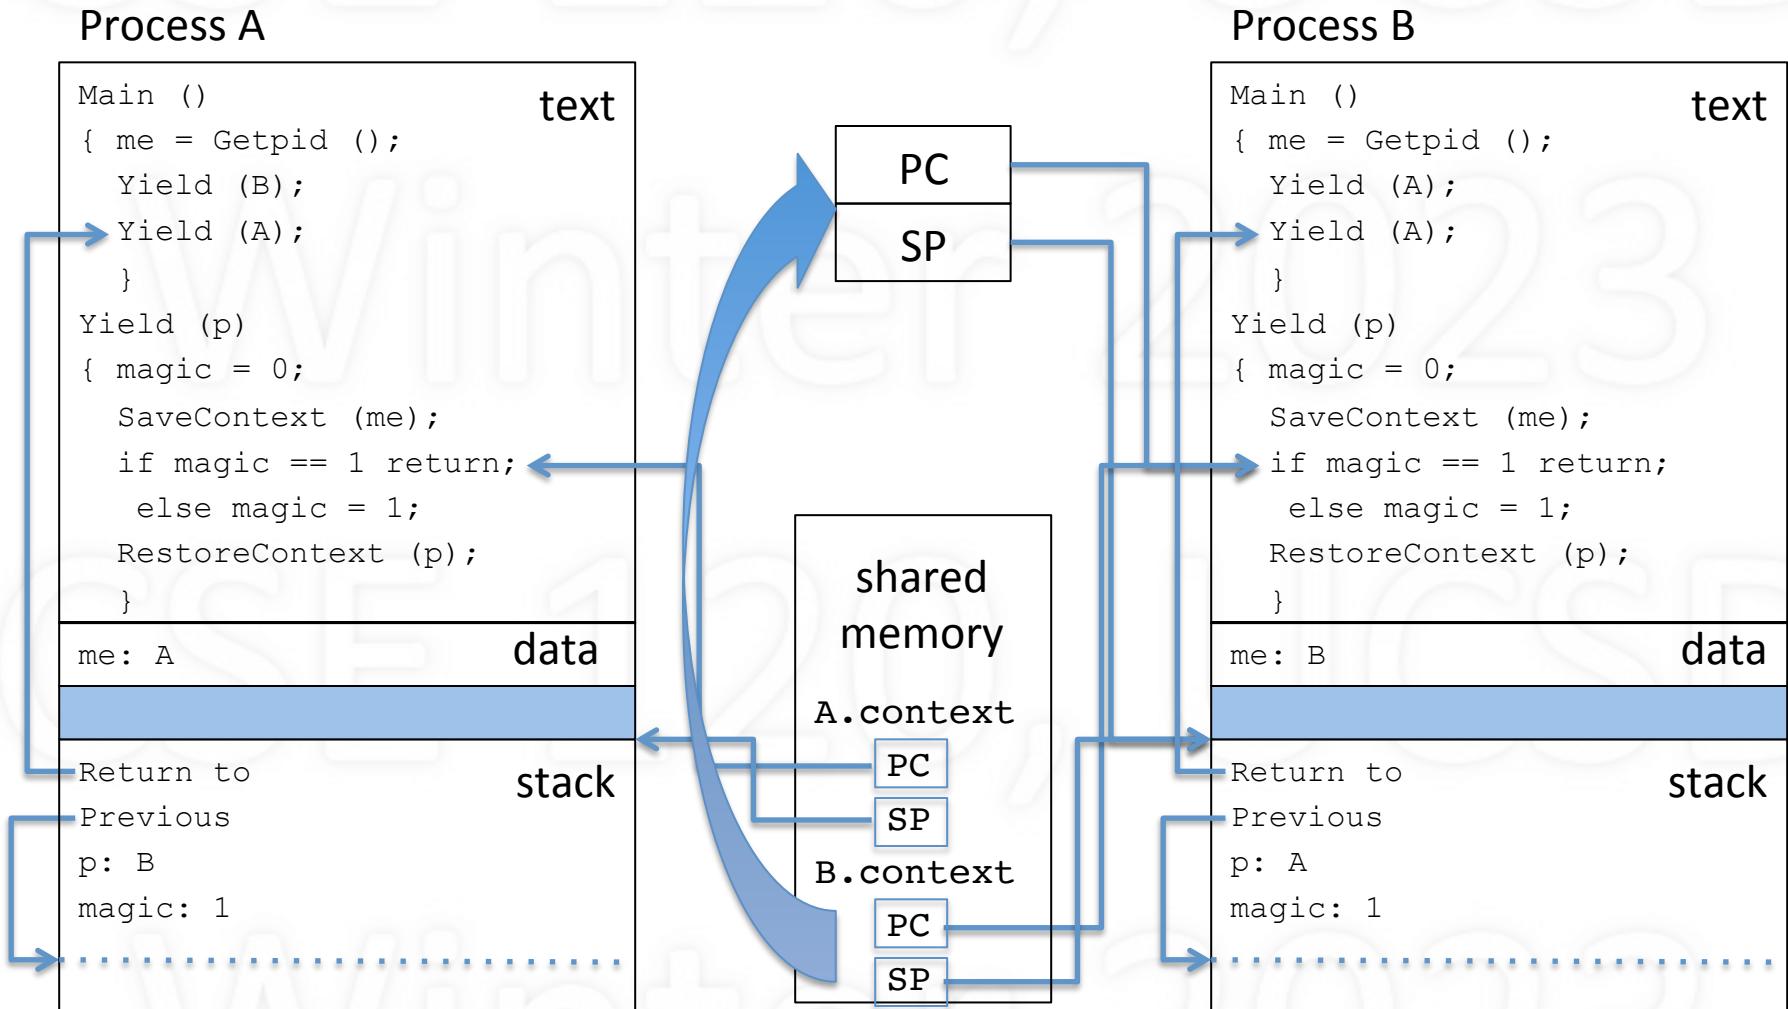
Process A



Process B

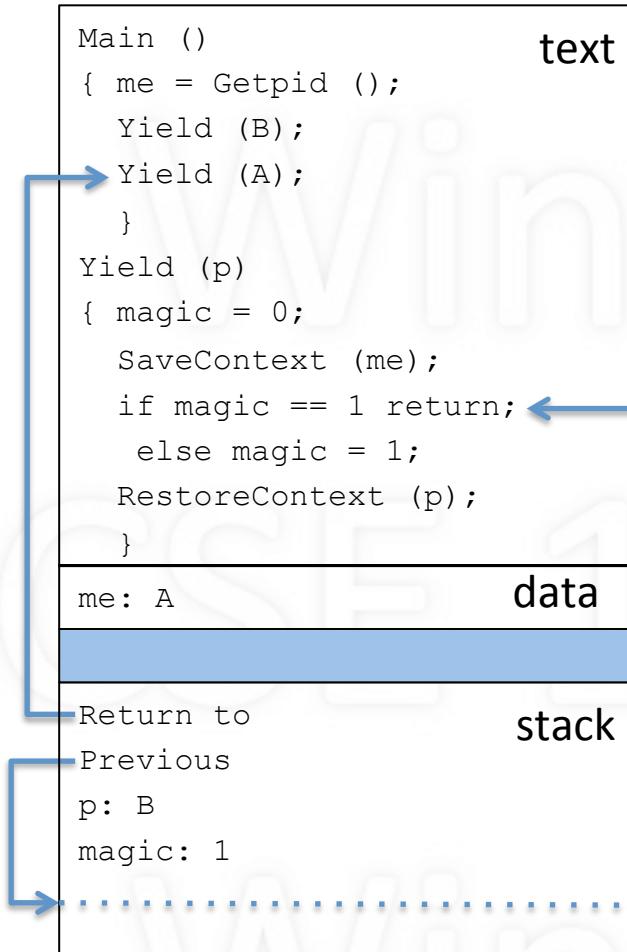


B's context is restored (i.e., the machine state is now that of B).
 The PC points to the if statement.

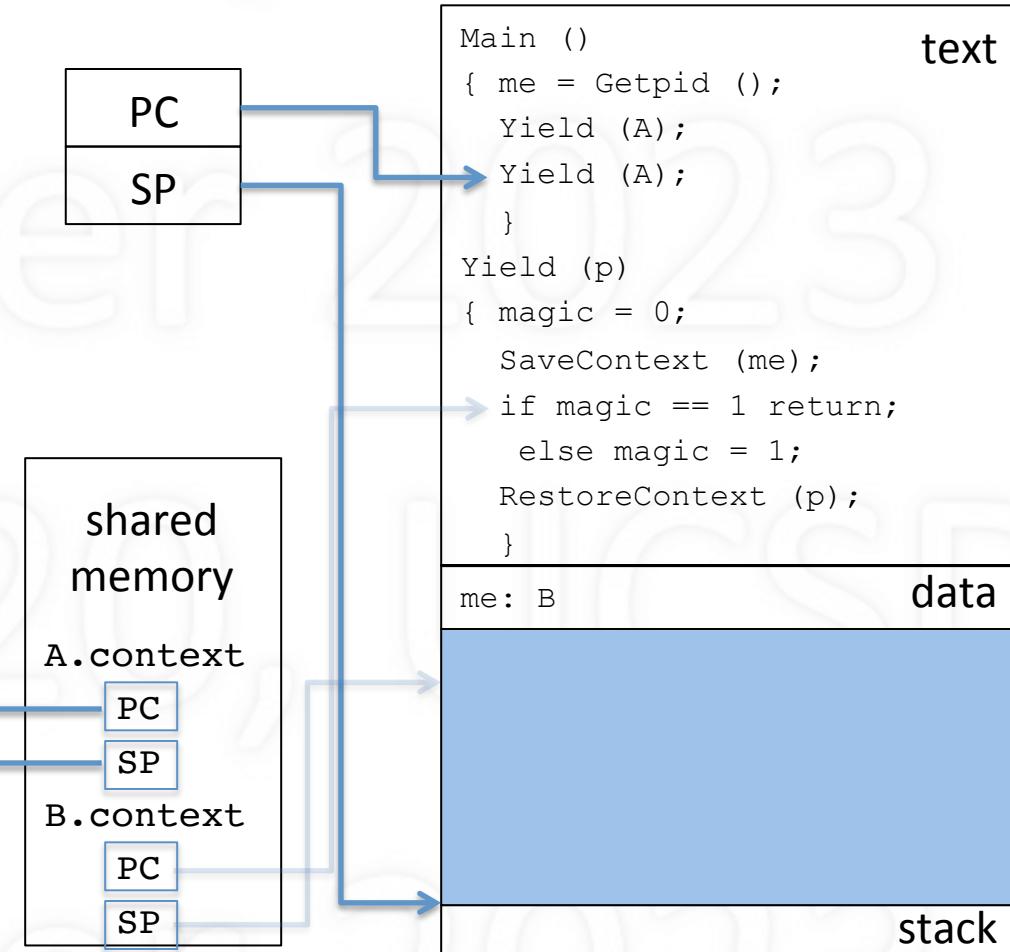


Since magic equals 1, B returns from Yield (unlike last time when magic equaled 0). No wonder it's called magic!

Process A

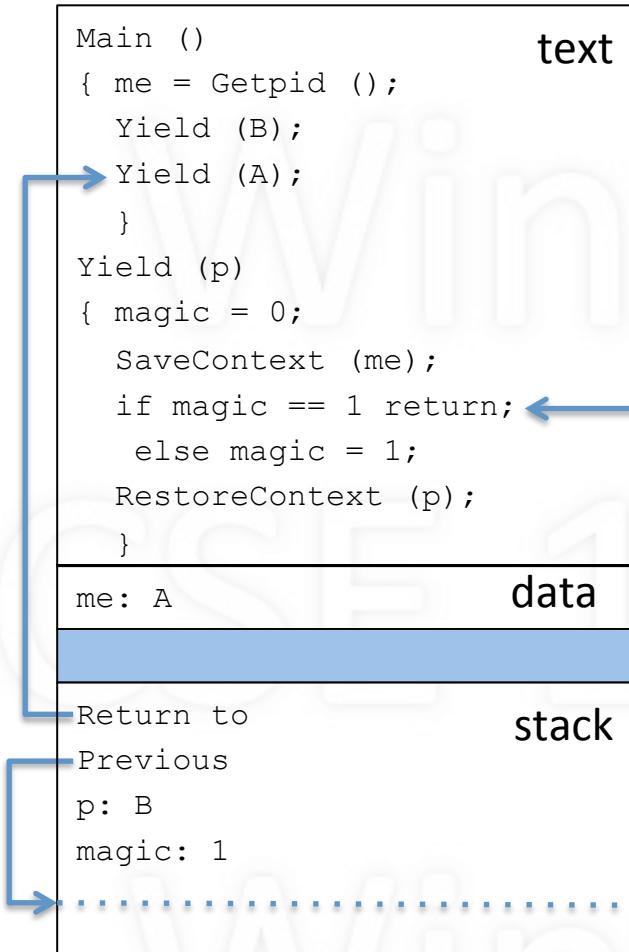


Process B

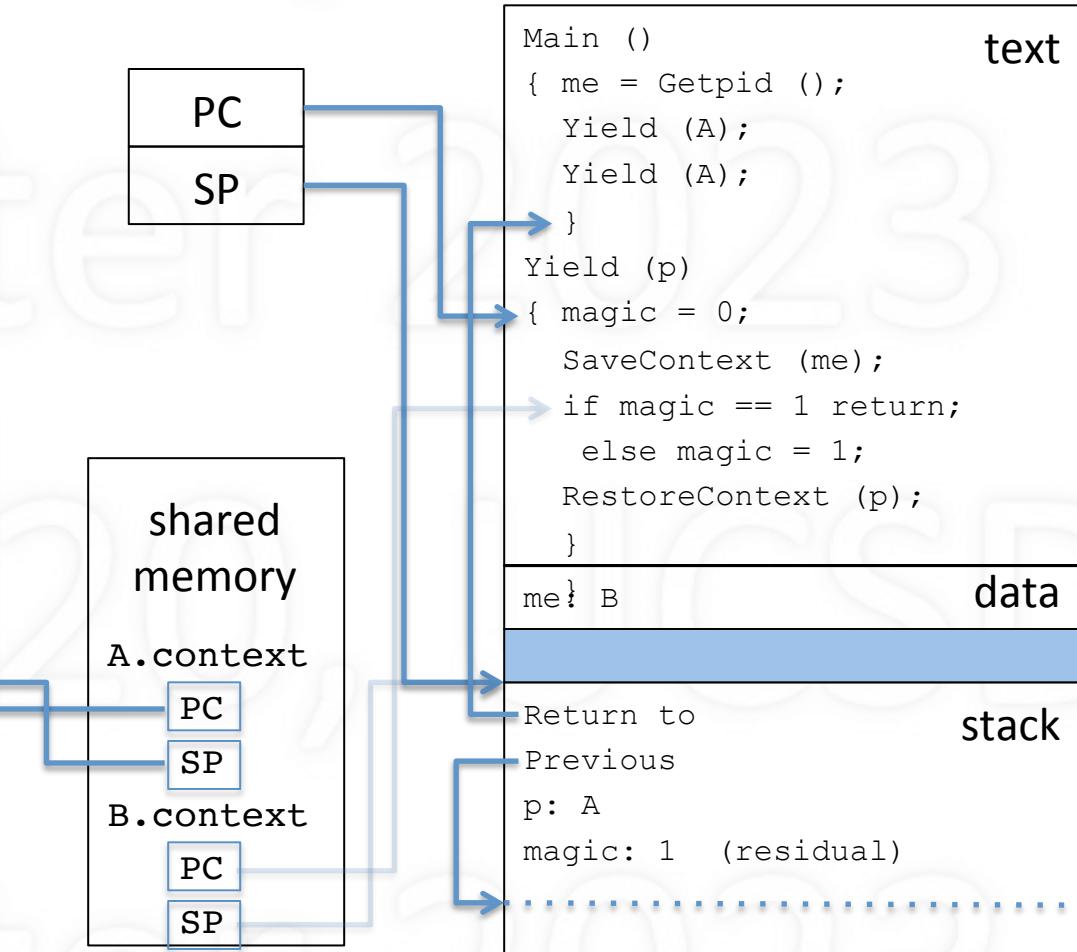


See if you can figure out how the rest of this works.

Process A

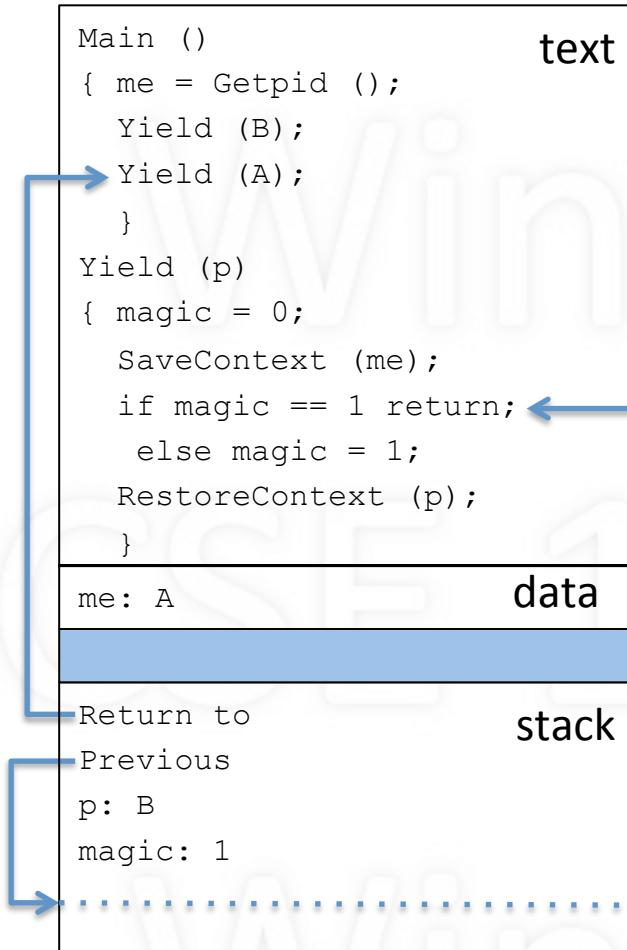


Process B

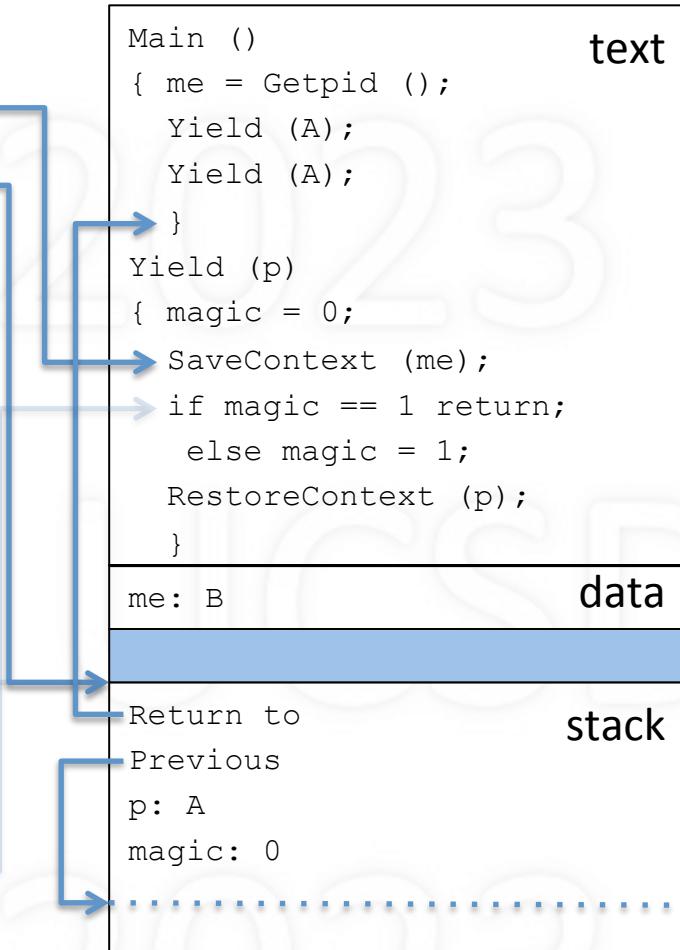


B is yielding to A.

Process A

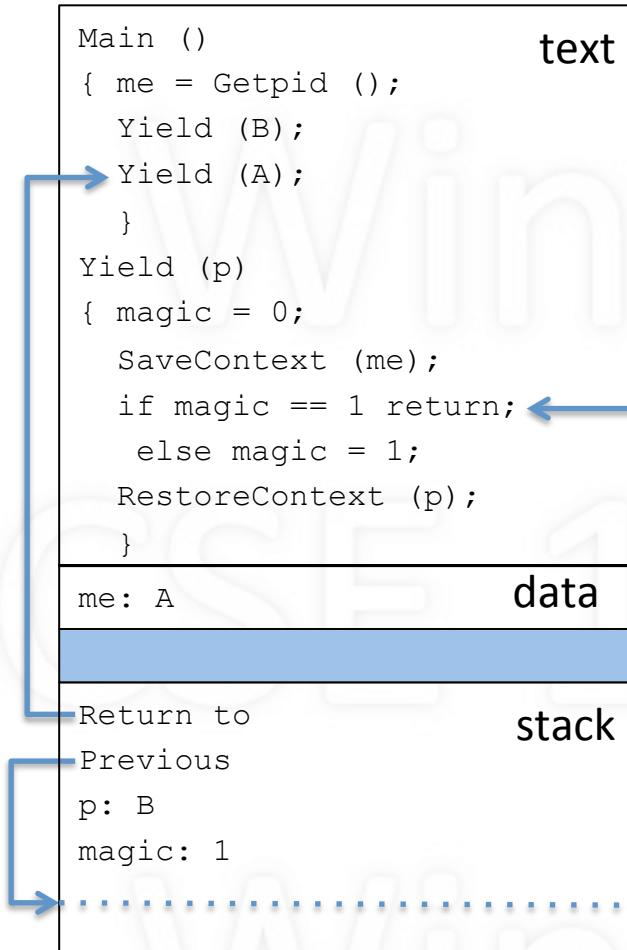


Process B

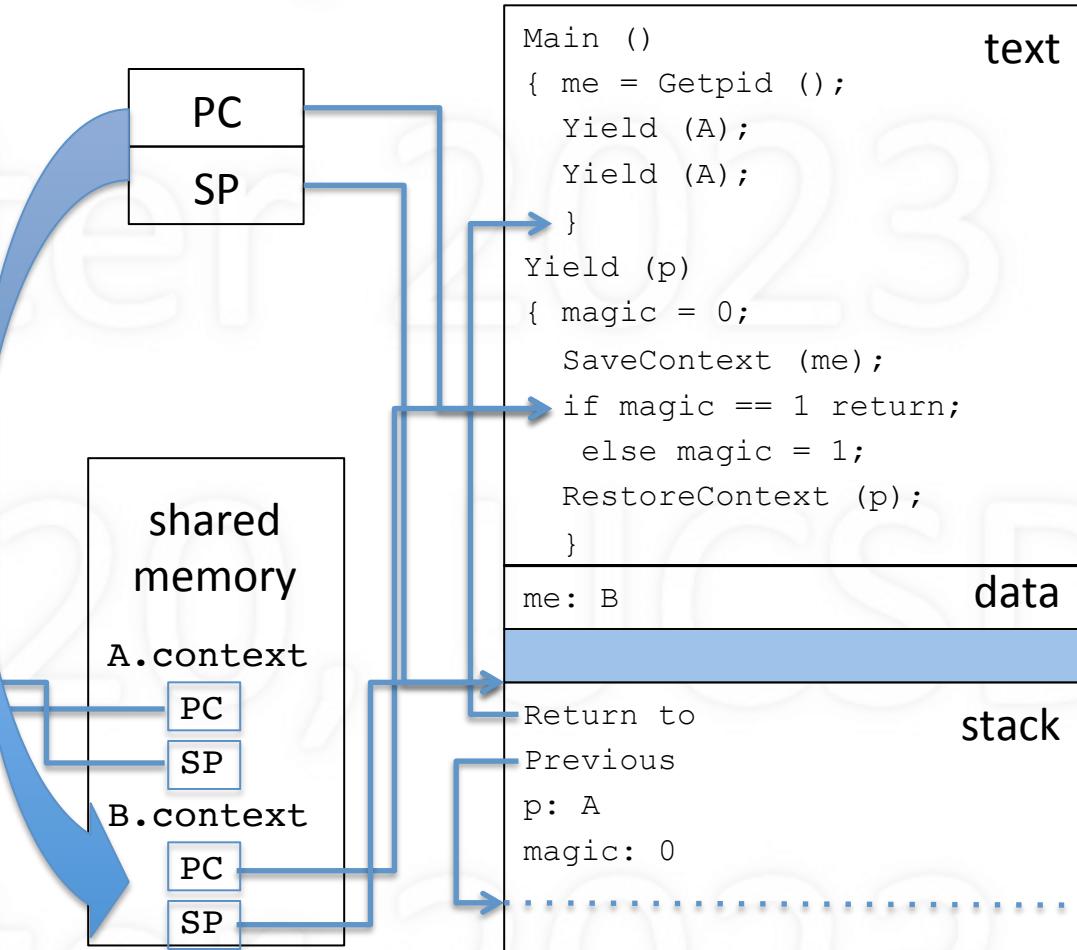


Goes through the same steps as we saw for A.

Process A

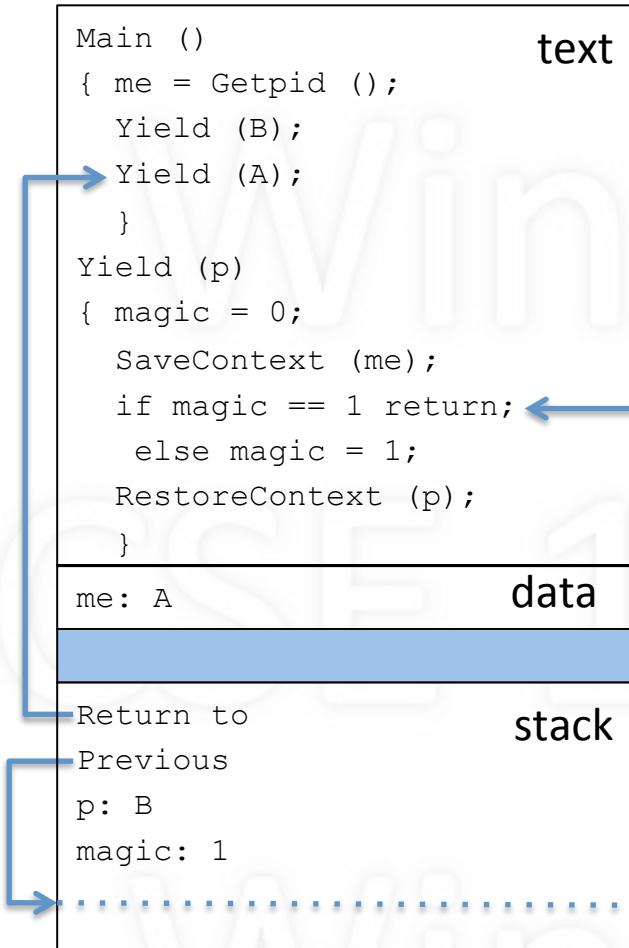


Process B

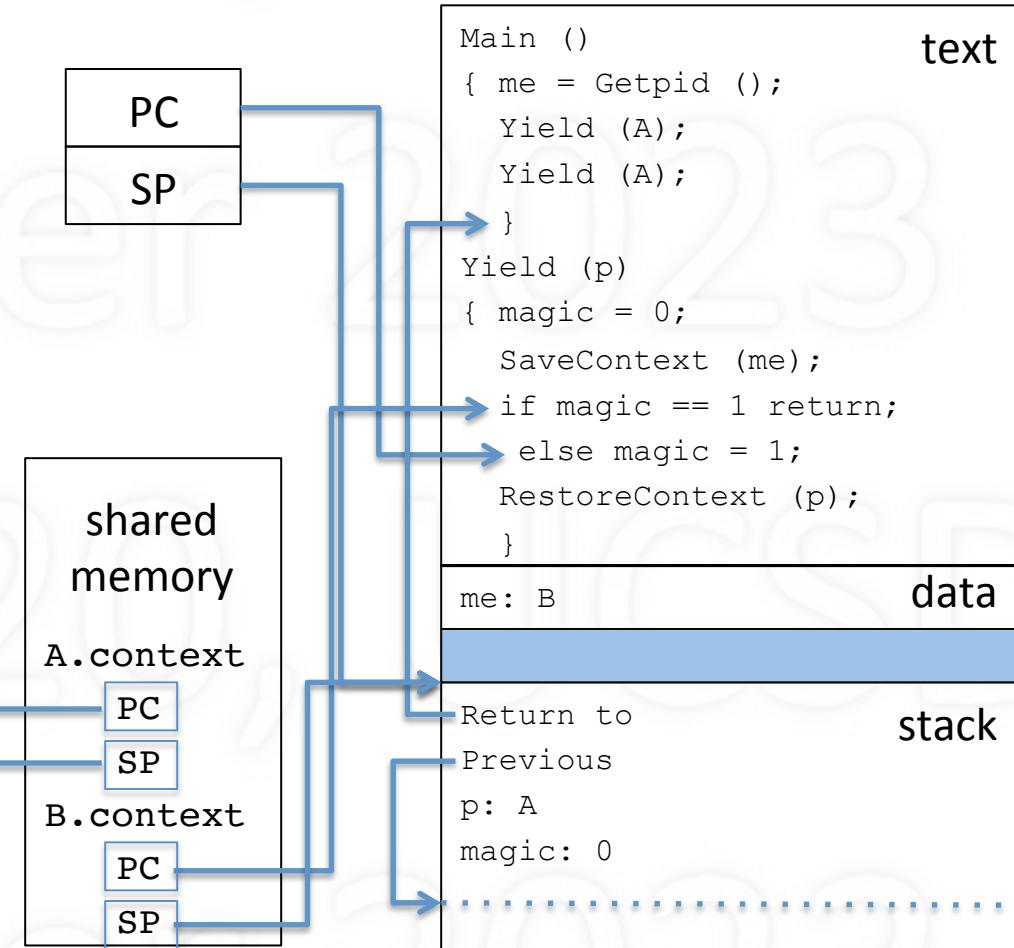


Note value of magic.

Process A

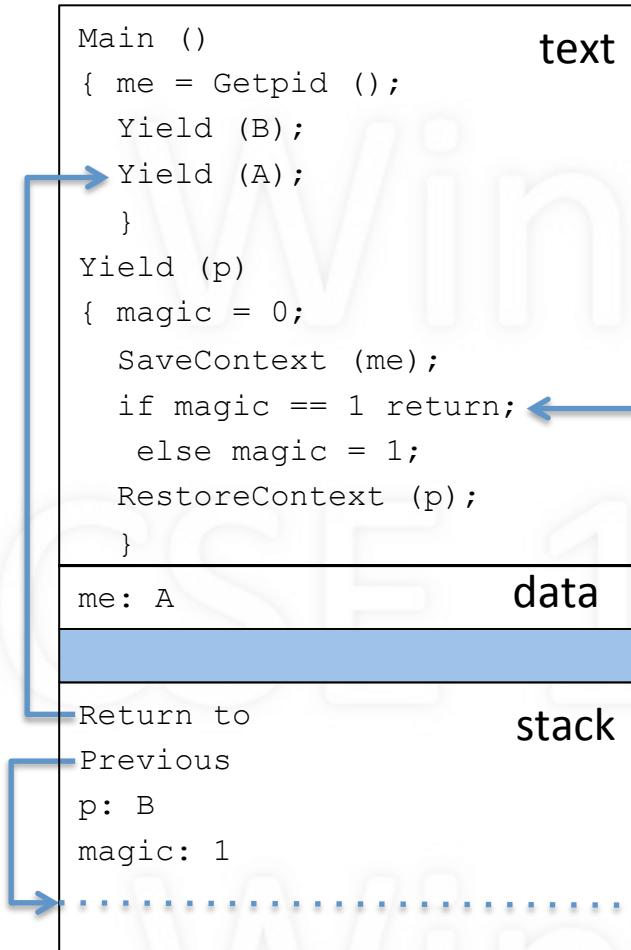


Process B

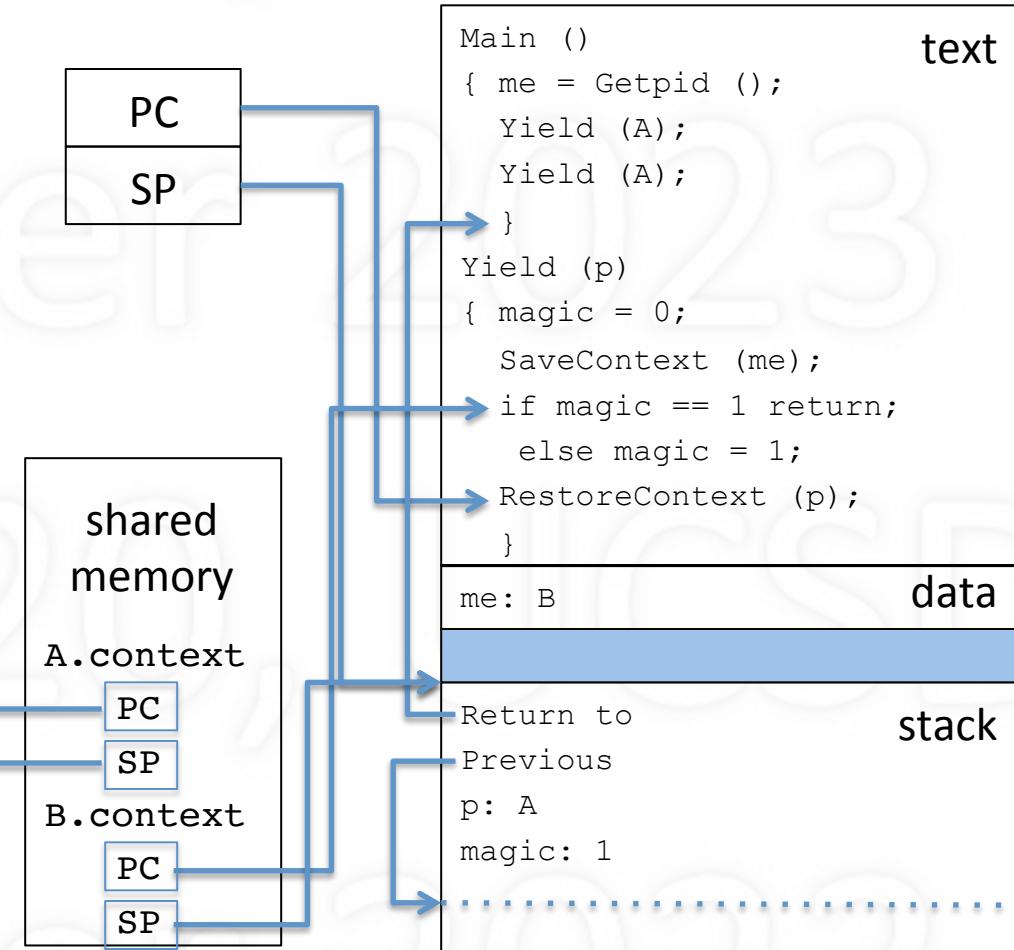


About to restore context of A.

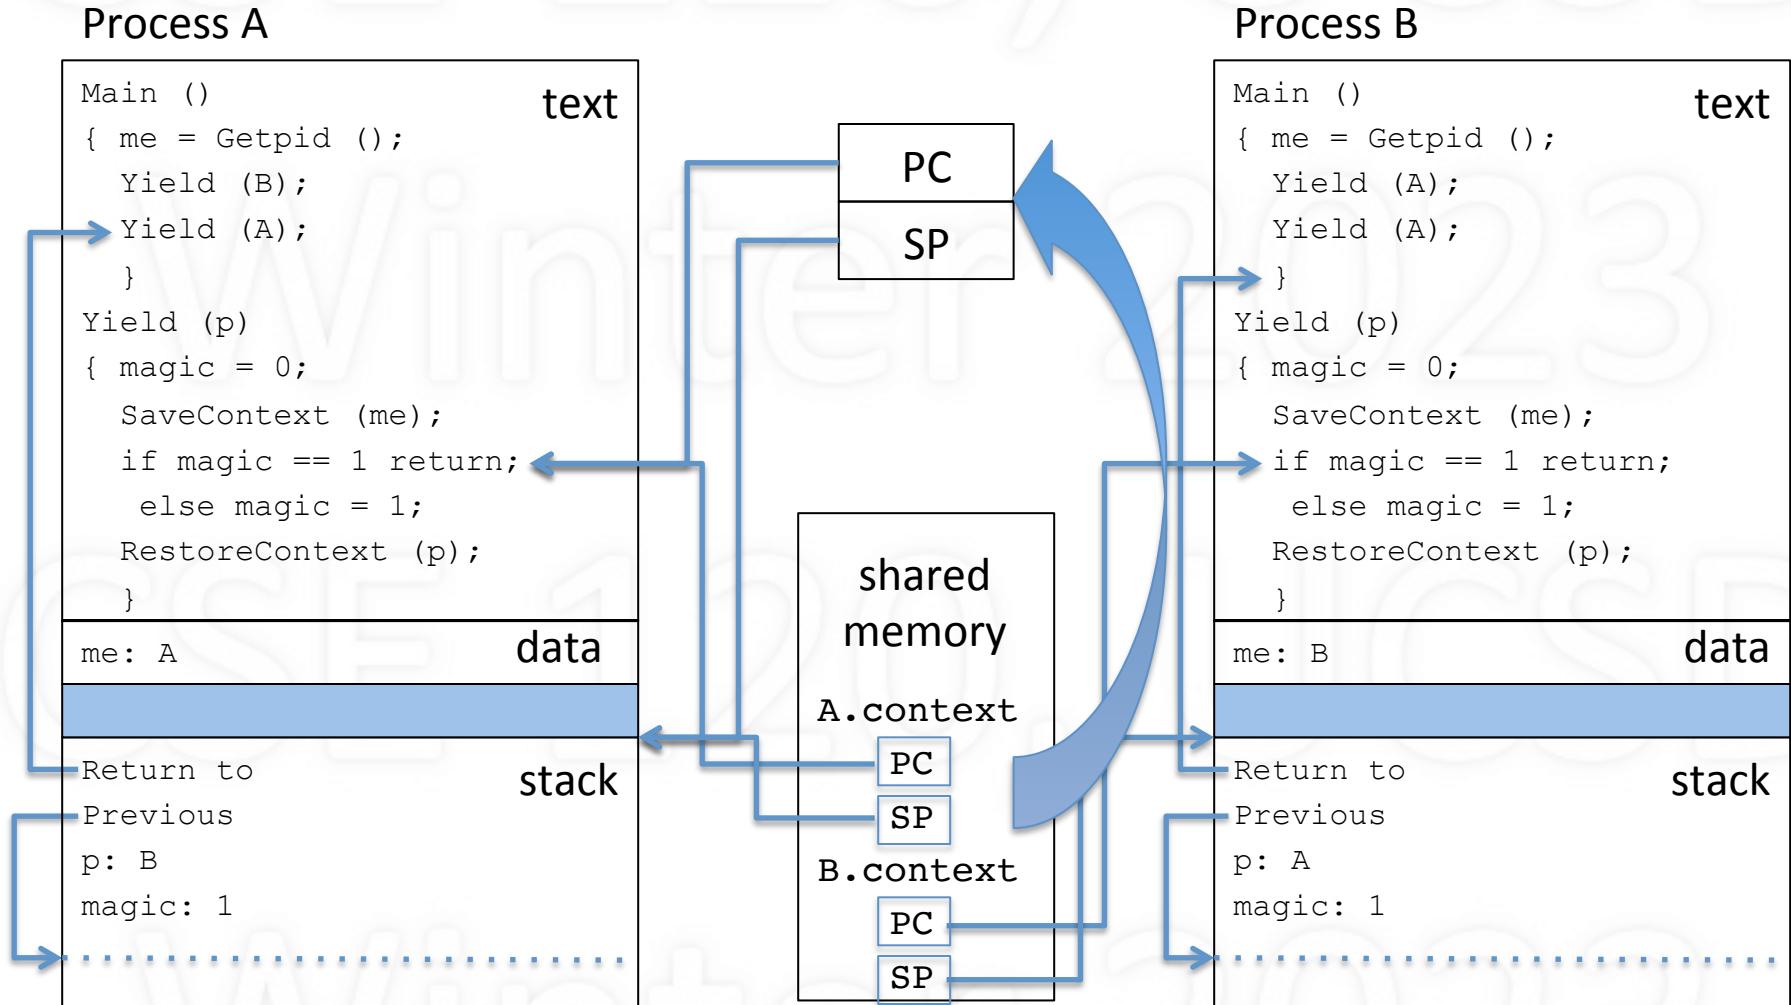
Process A



Process B

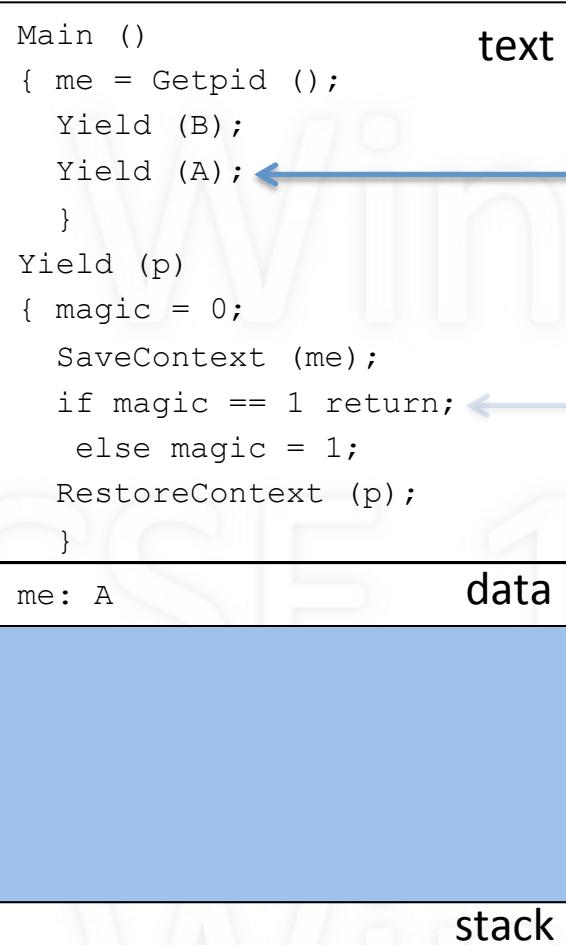


A resumes. Note value of magic.

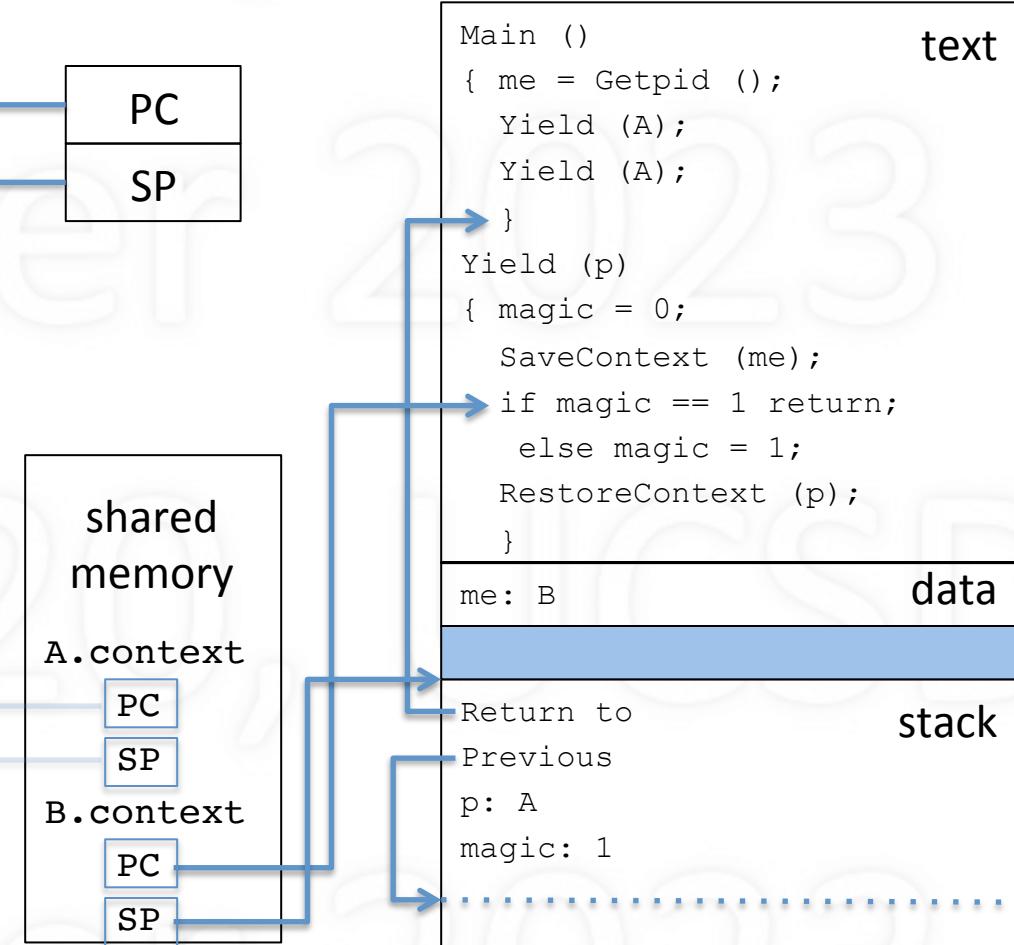


Returns from Yield, and now about to call Yield again, but to itself!

Process A

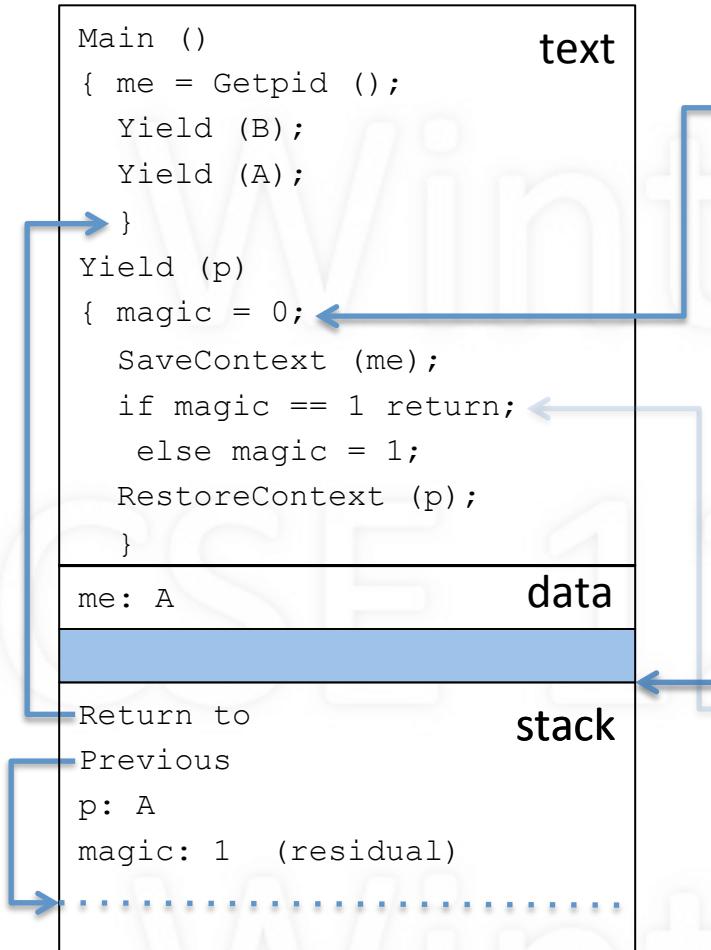


Process B

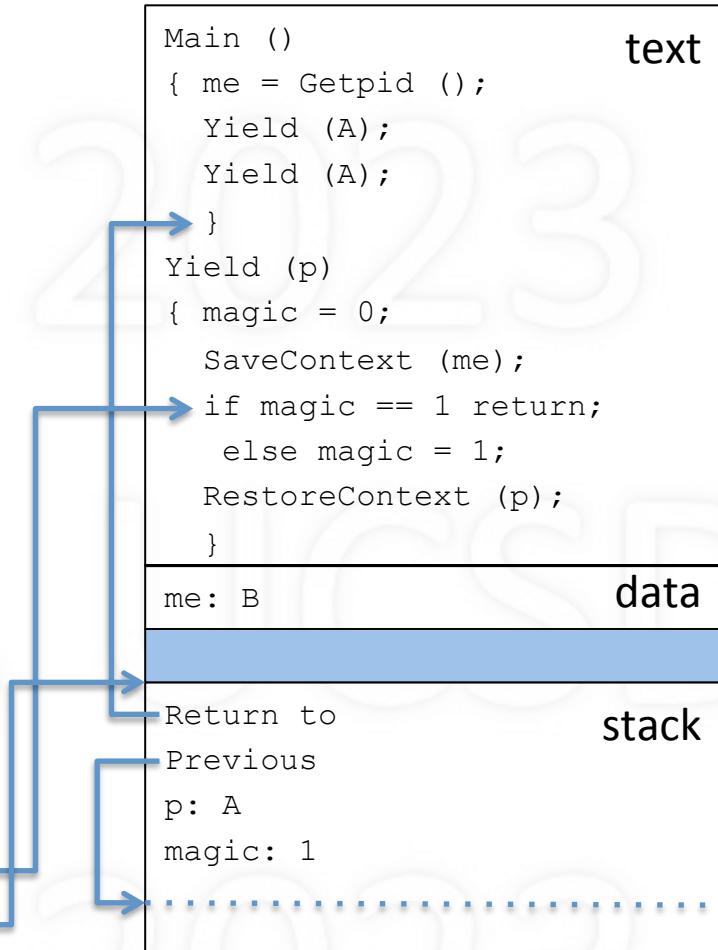


Inside Yield. Note new activation record on stack.

Process A

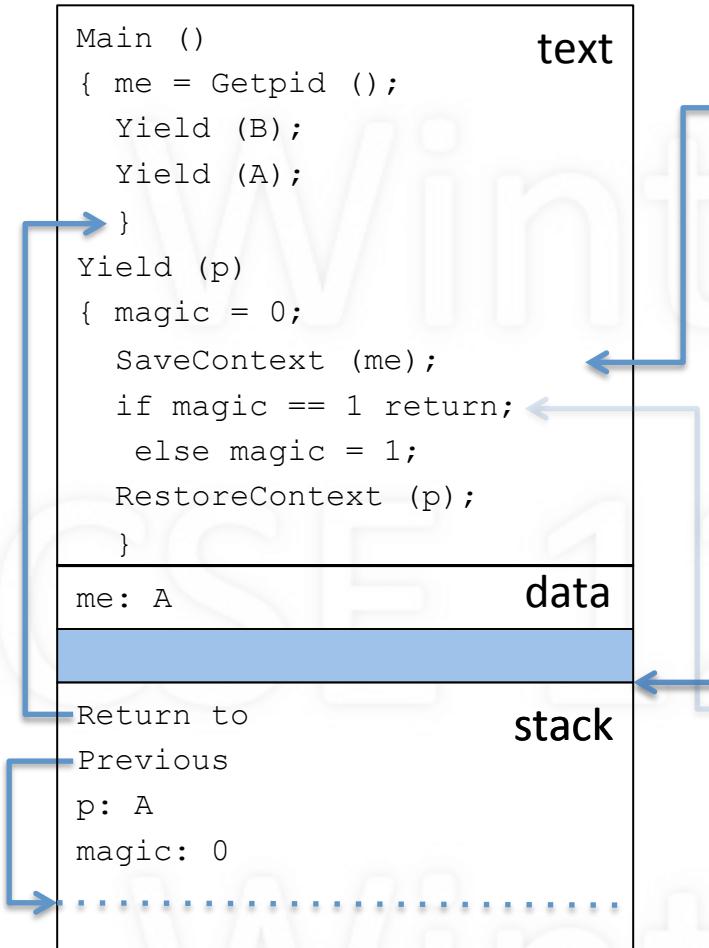


Process B

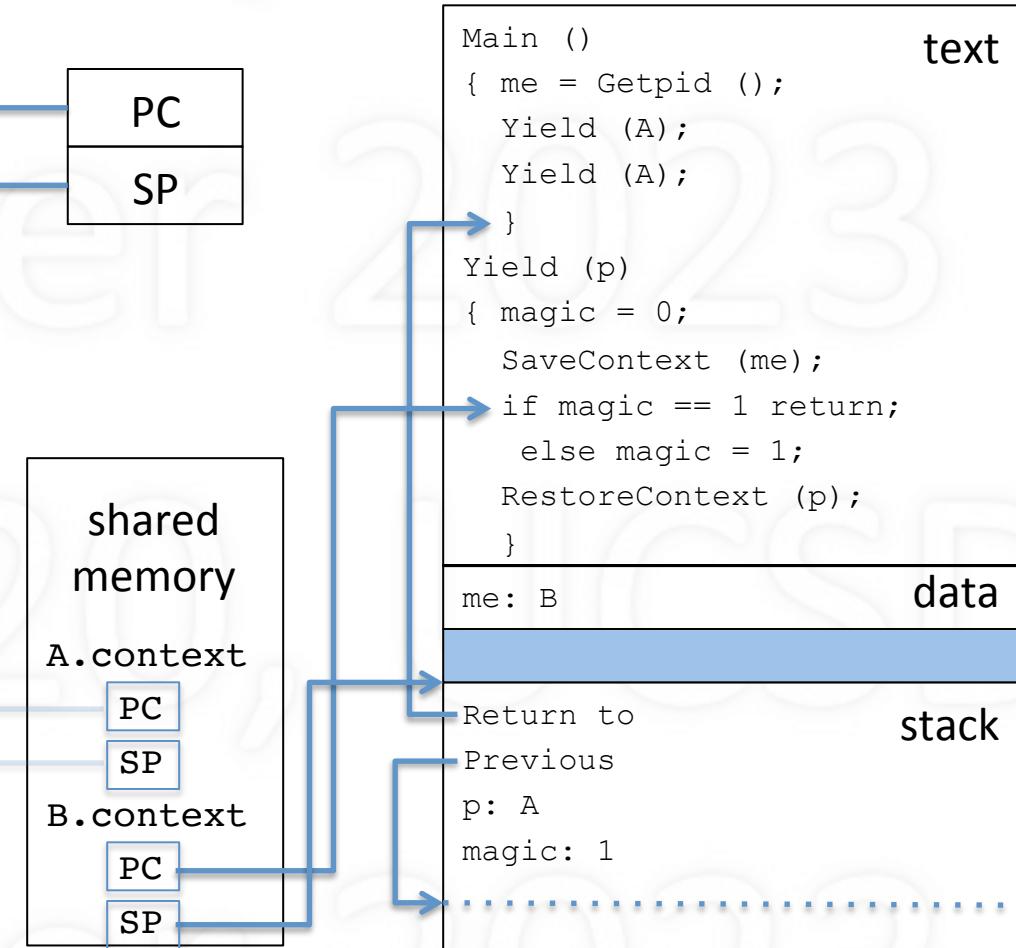


About to save context of A.

Process A

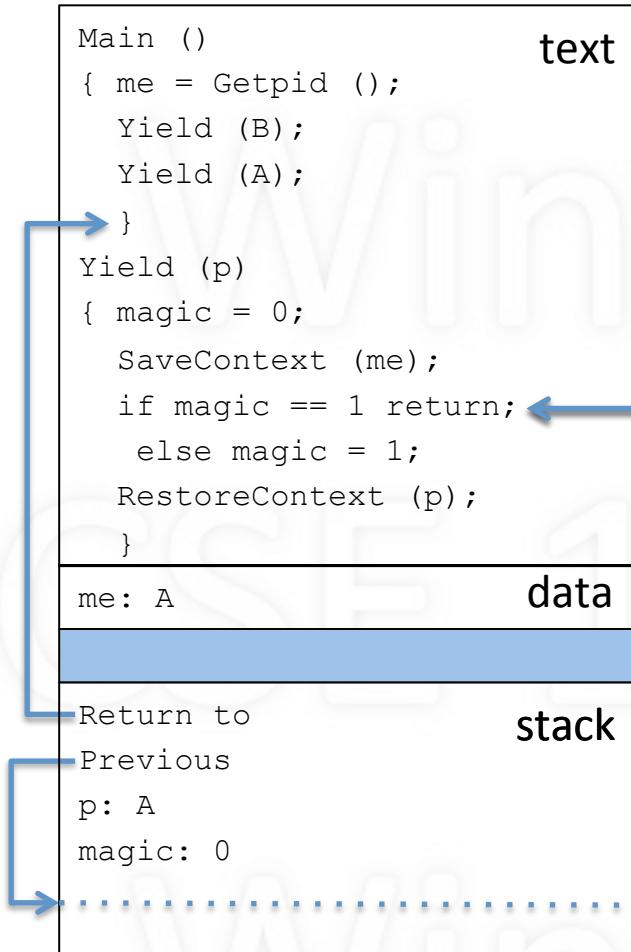


Process B

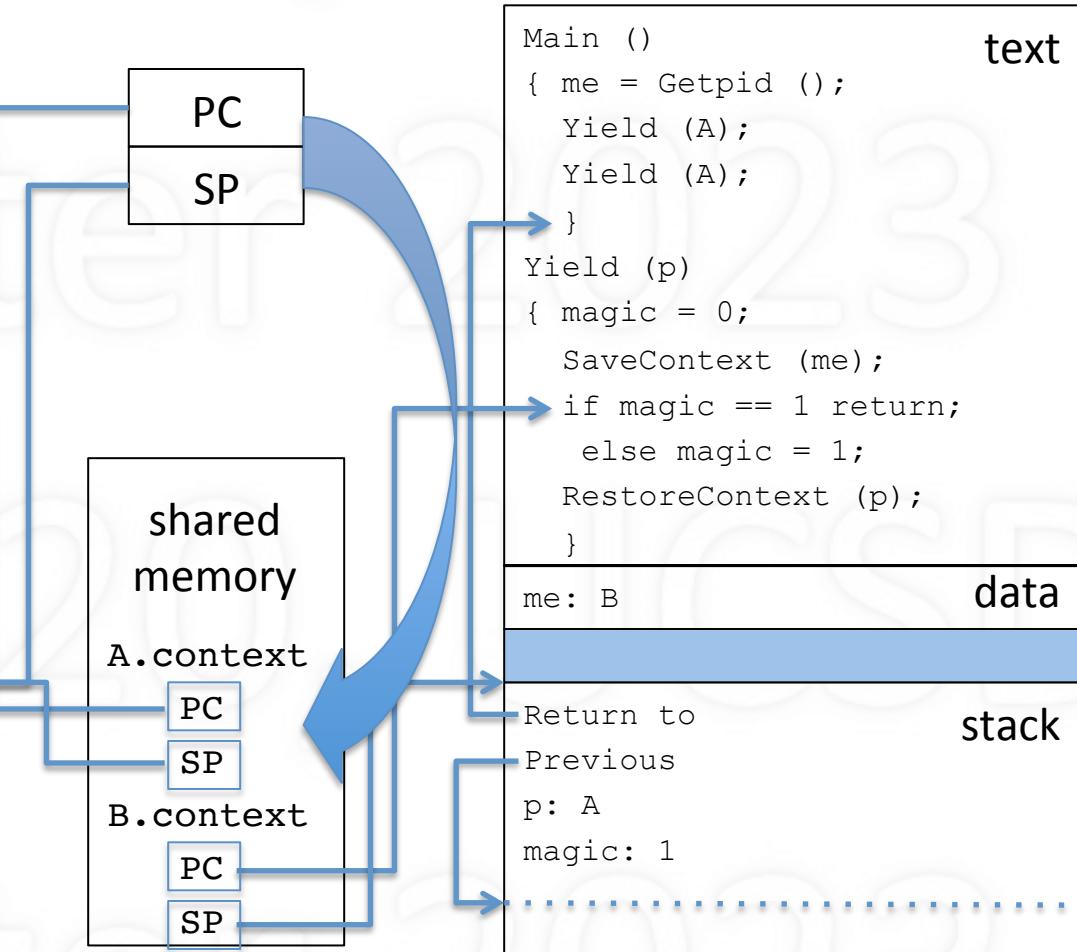


Note value of magic.

Process A

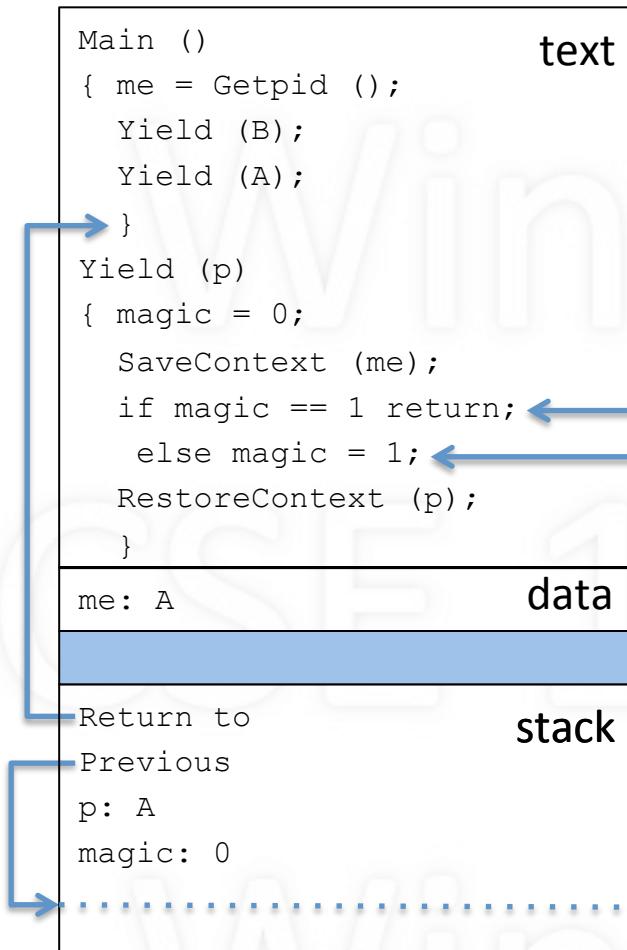


Process B

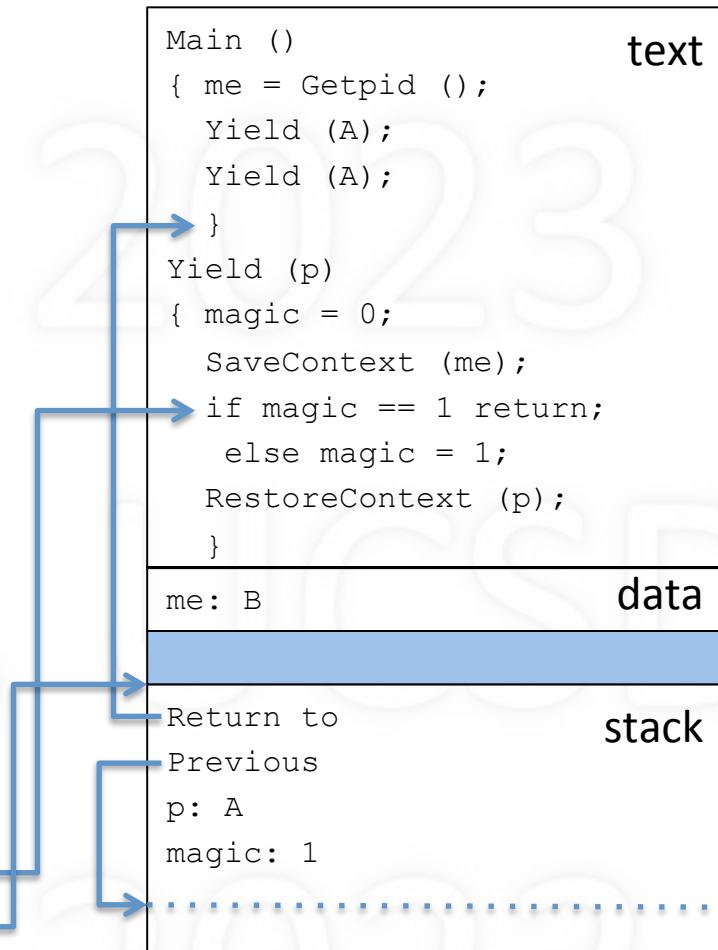


Set magic.

Process A

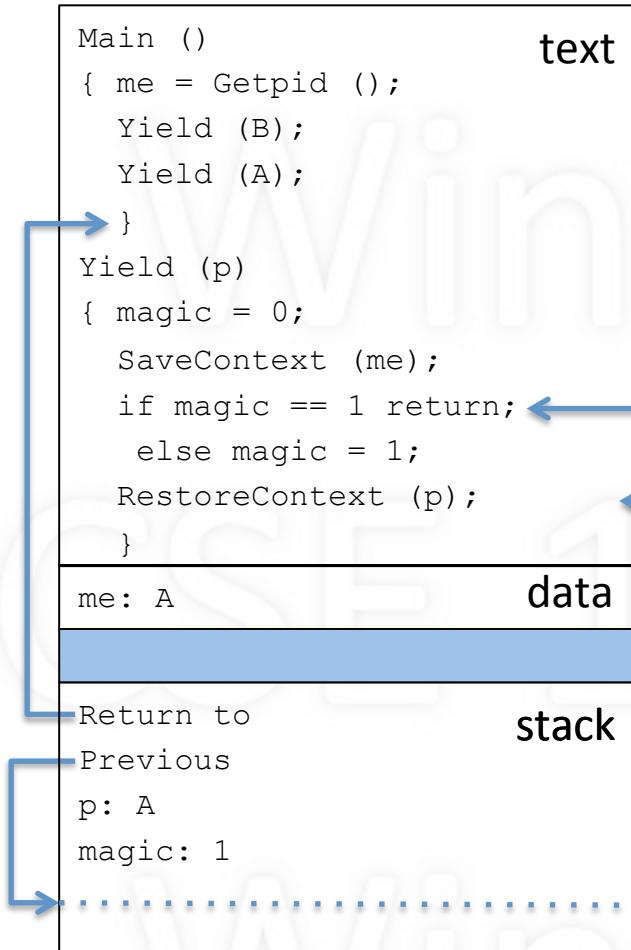


Process B

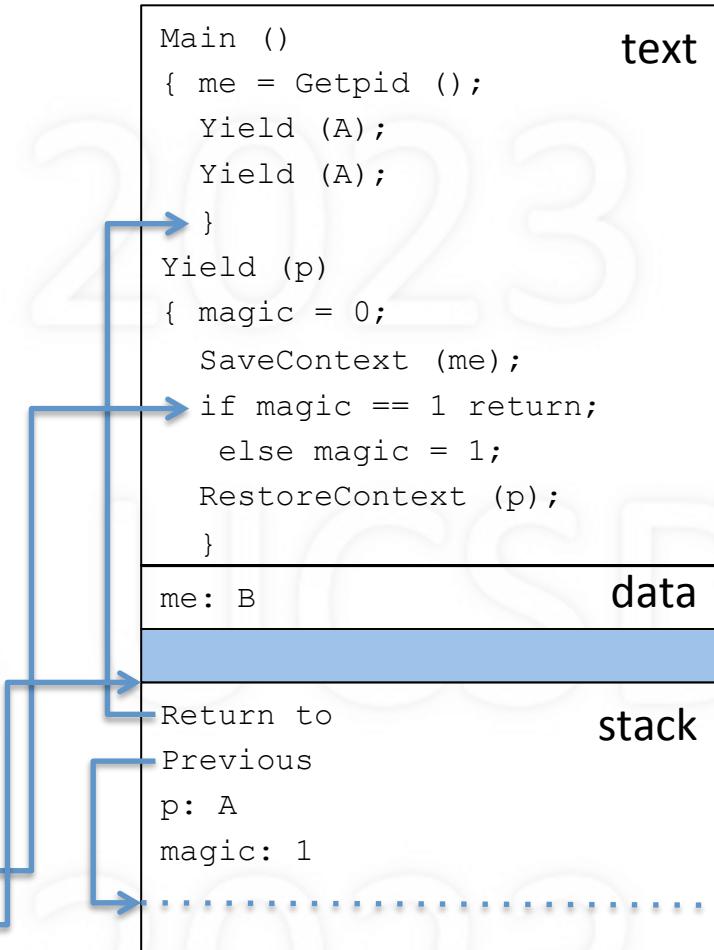


Restore context of A!

Process A

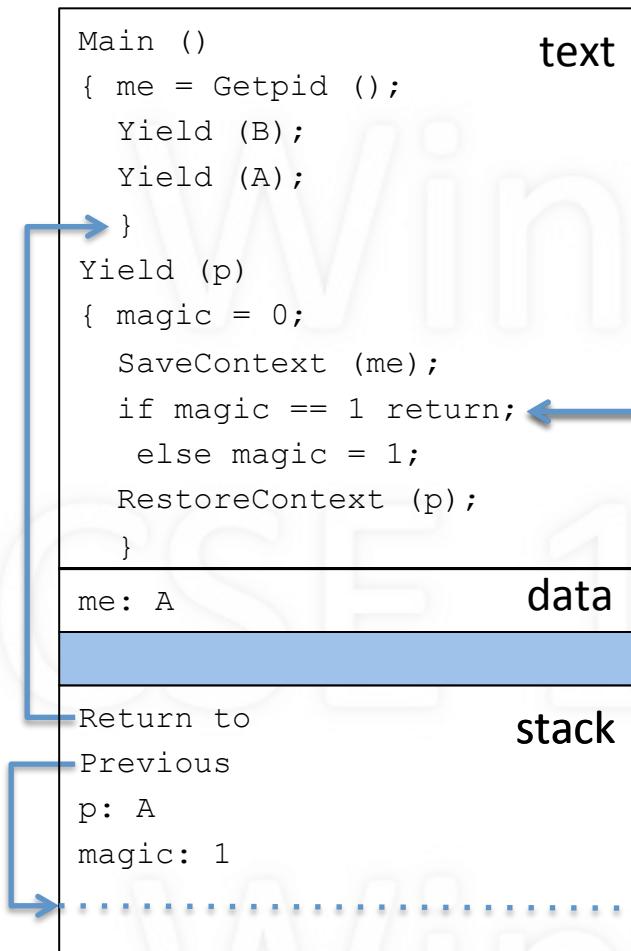


Process B

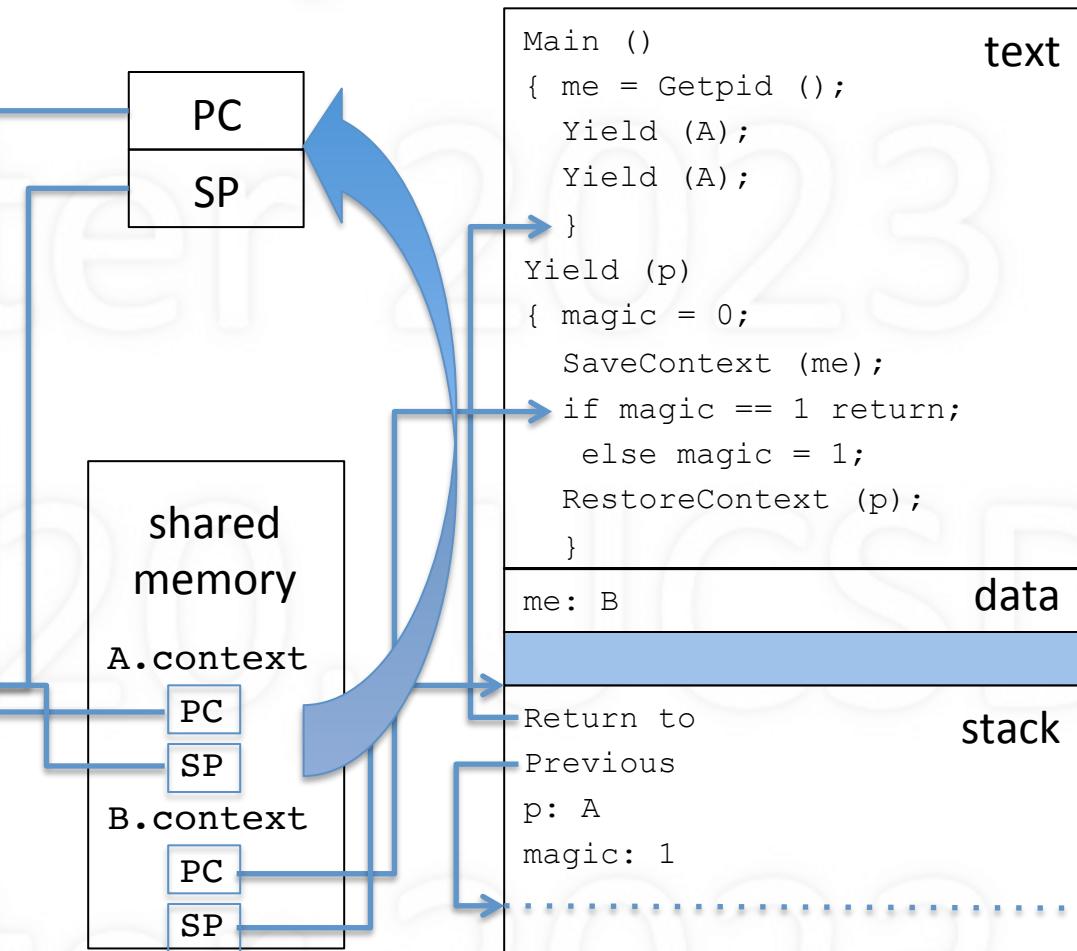


Note where A executes.

Process A



Process B



Return from Yield.

Process A

```
Main ()
{ me = Getpid ();
  Yield (B);
  Yield (A);
}
Yield (p)
{ magic = 0;
  SaveContext (me);
  if magic == 1 return;
  else magic = 1;
  RestoreContext (p);
}
```

text

me: A

data

stack

PC
SP

shared
memory

A.context

PC
SP

B.context

PC
SP

Process B

```
Main ()
{ me = Getpid ();
  Yield (A);
  Yield (A);
}
Yield (p)
{ magic = 0;
  SaveContext (me);
  if magic == 1 return;
  else magic = 1;
  RestoreContext (p);
}
```

text

me: B

data

stack

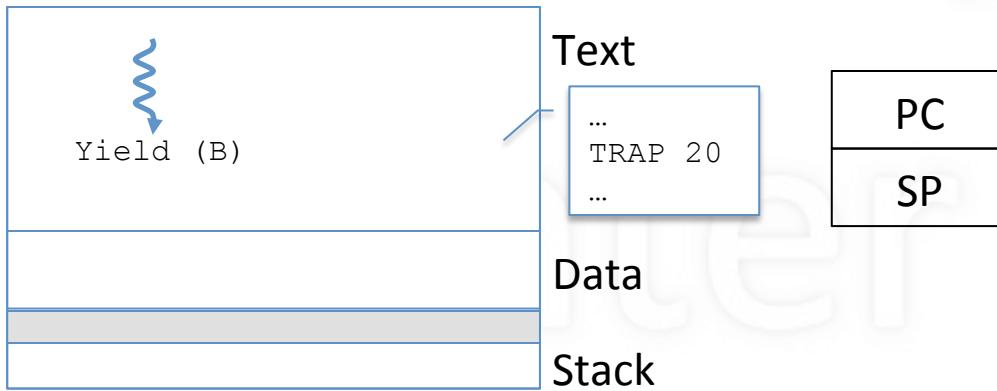
Return to
Previous
p: A
magic: 1

Yielding via the Kernel

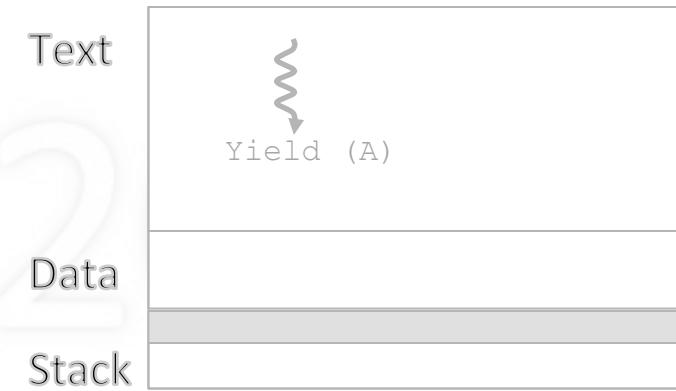
- Yield routine is common code: put in kernel
- Process contexts are also in the kernel
 - This way they are protected
 - Only needed by Yield routine anyway
- But what is the kernel?
 - code that supports processes
 - runs as an extension of current process
- Has text, data, and multiple stacks

Yielding via the Kernel

Process A

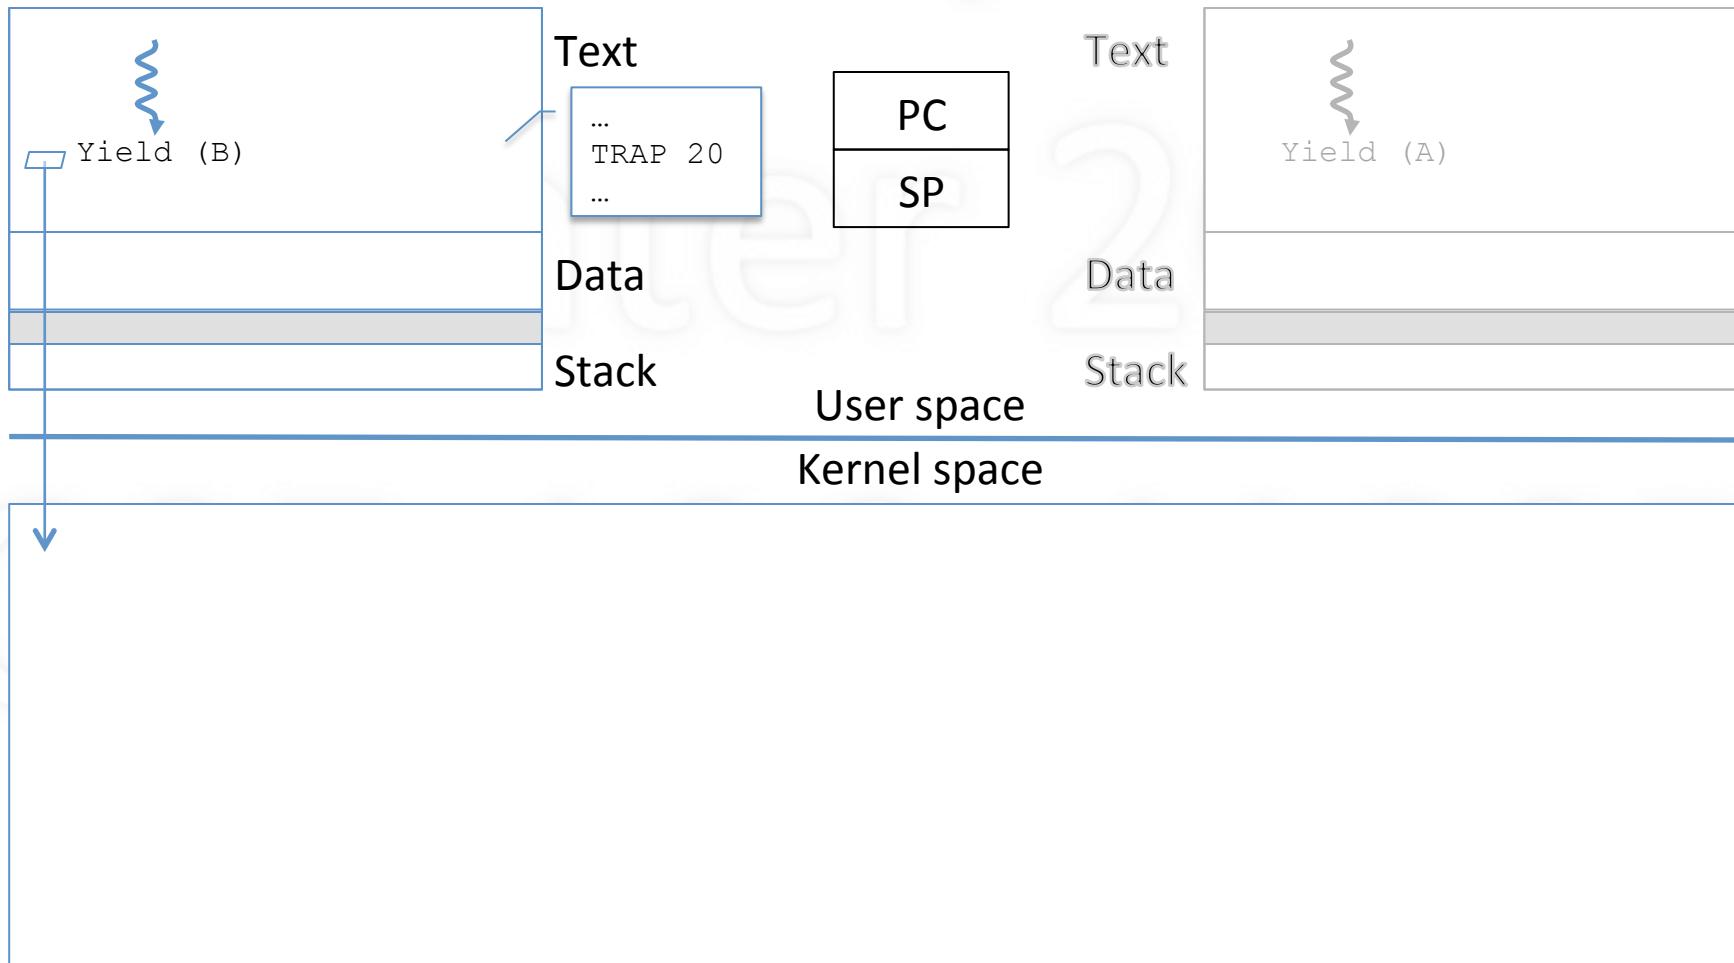


Process B

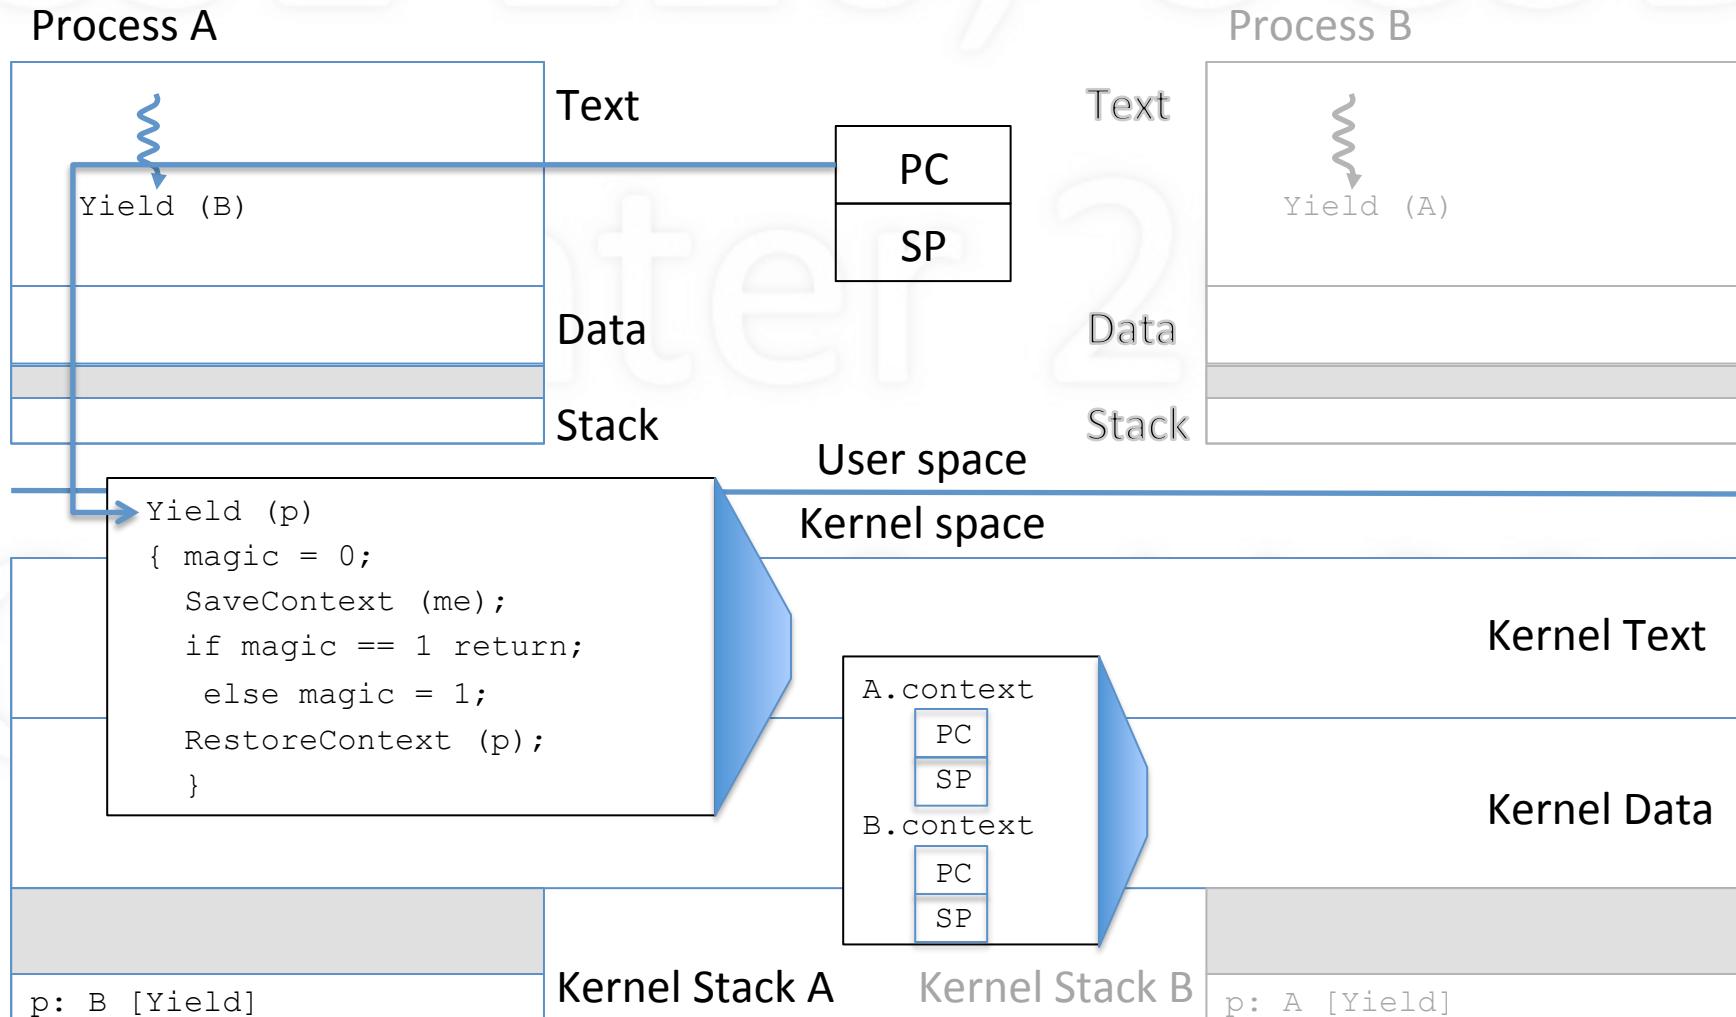


Trap Causes Kernel Entry

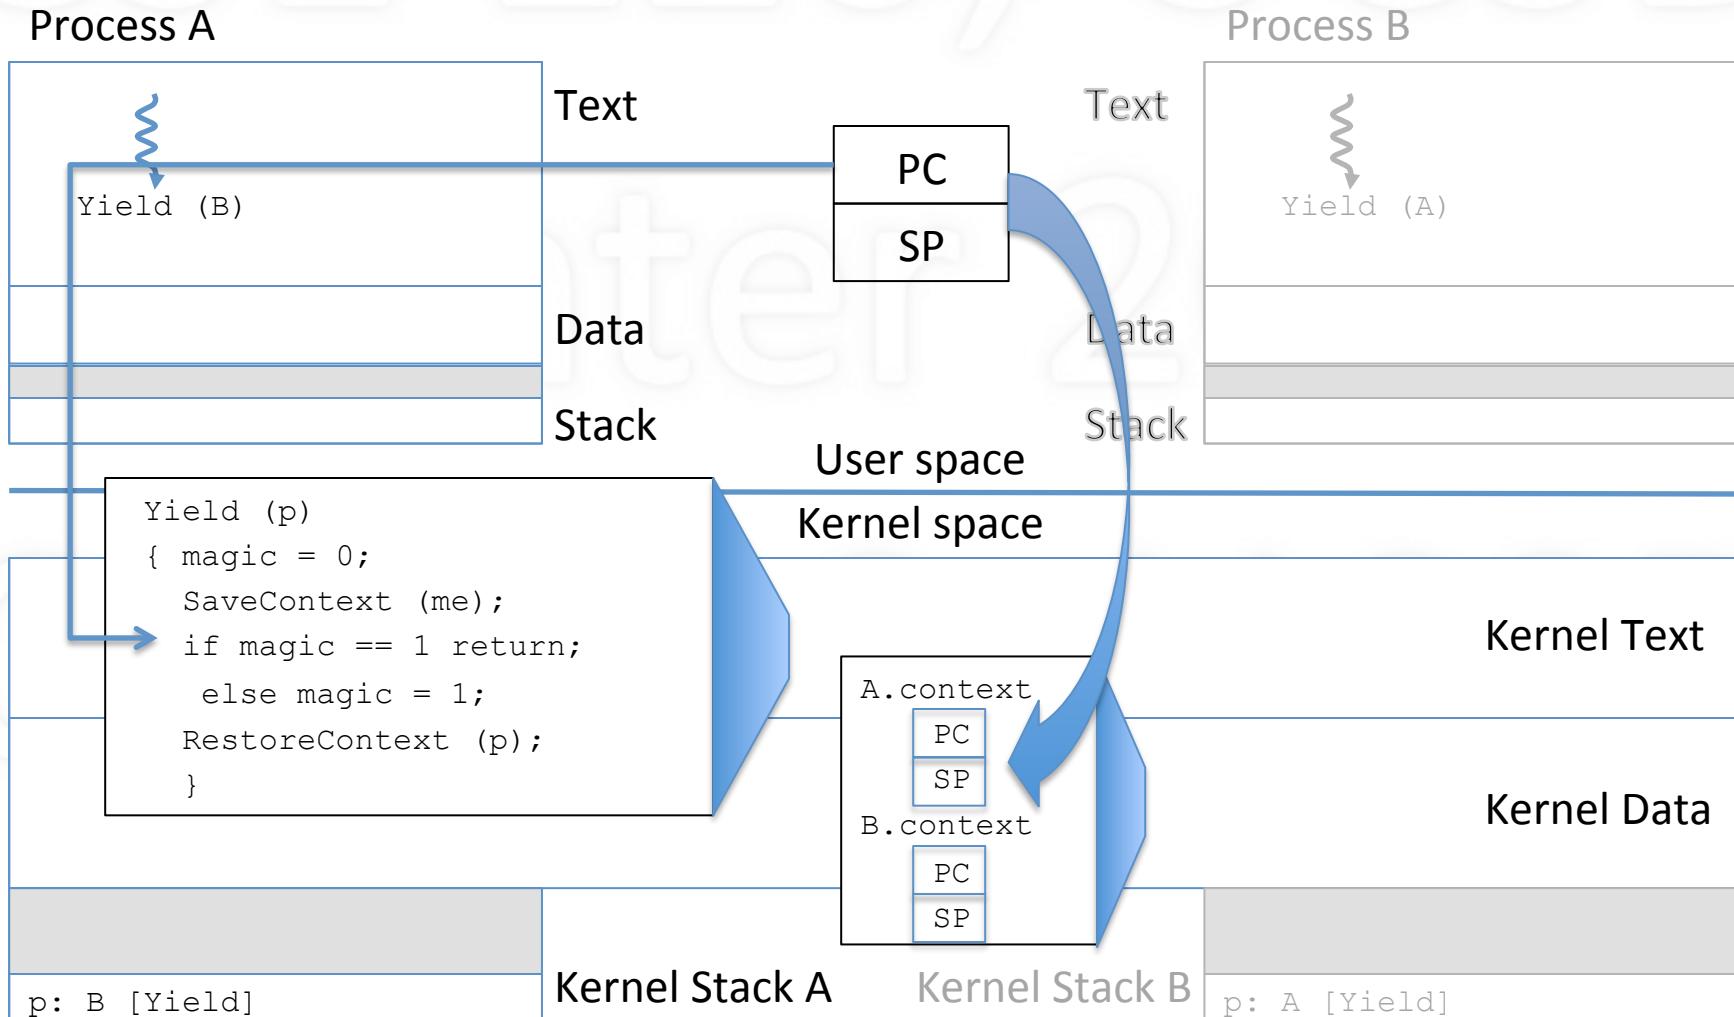
Process A



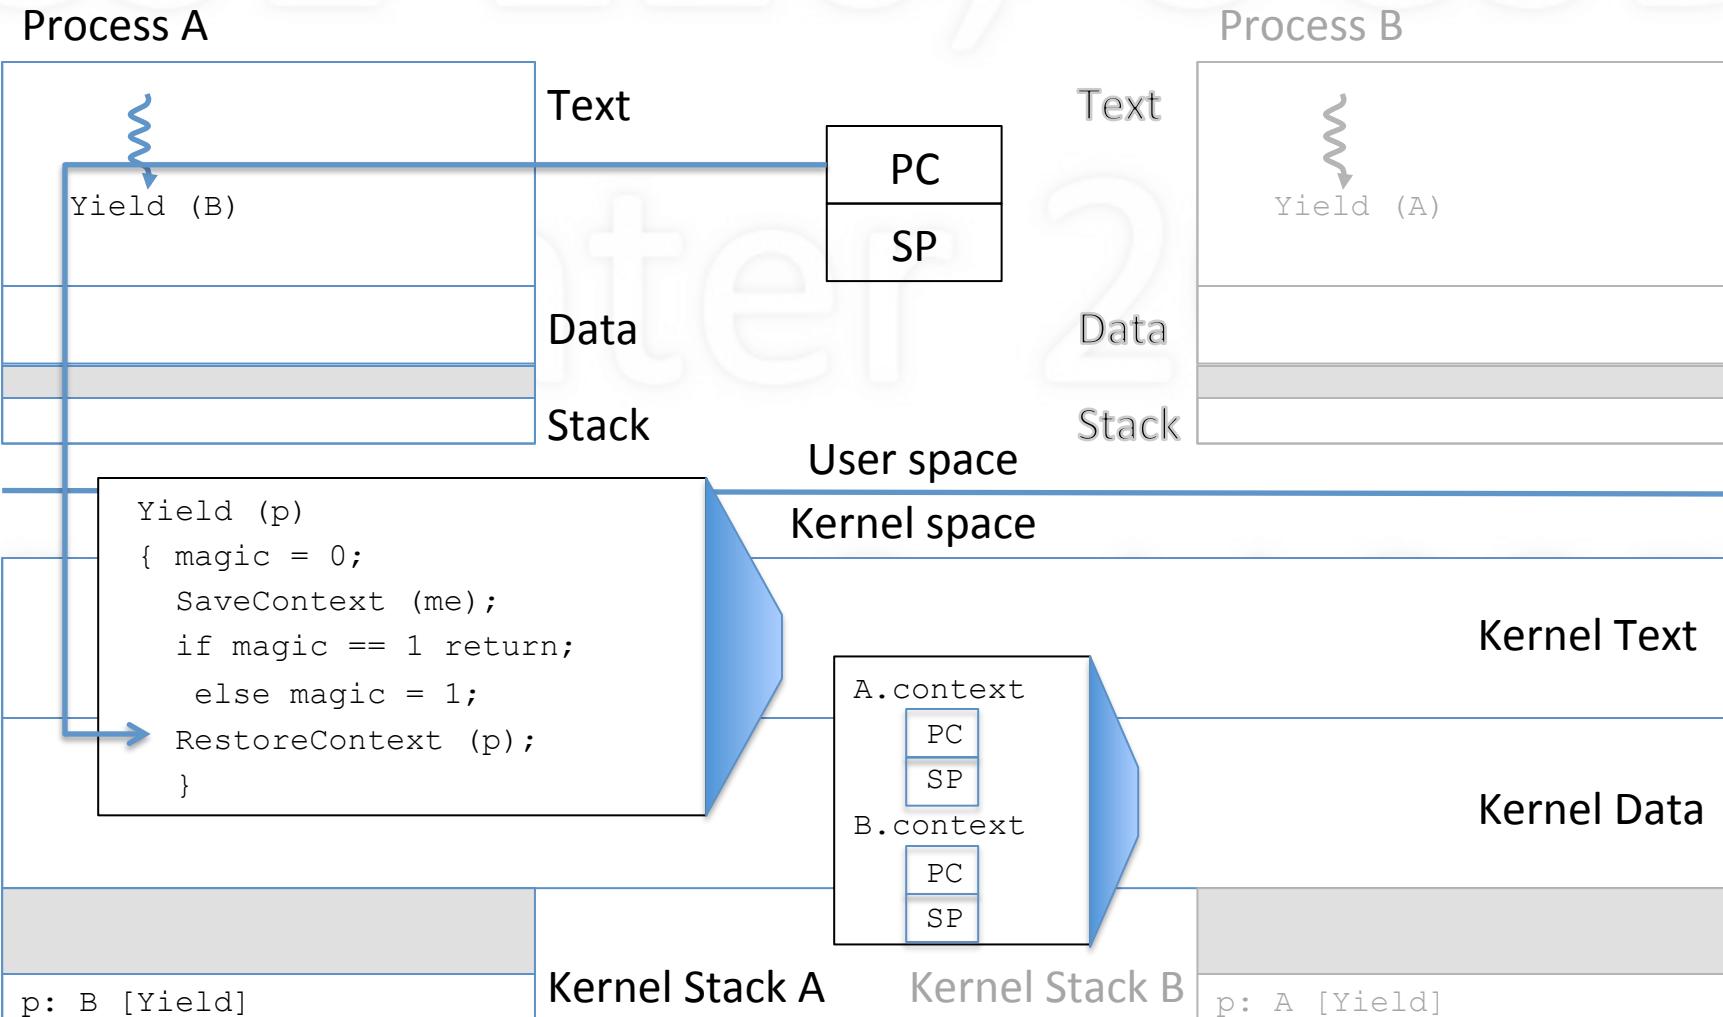
Routine for Yield Executes



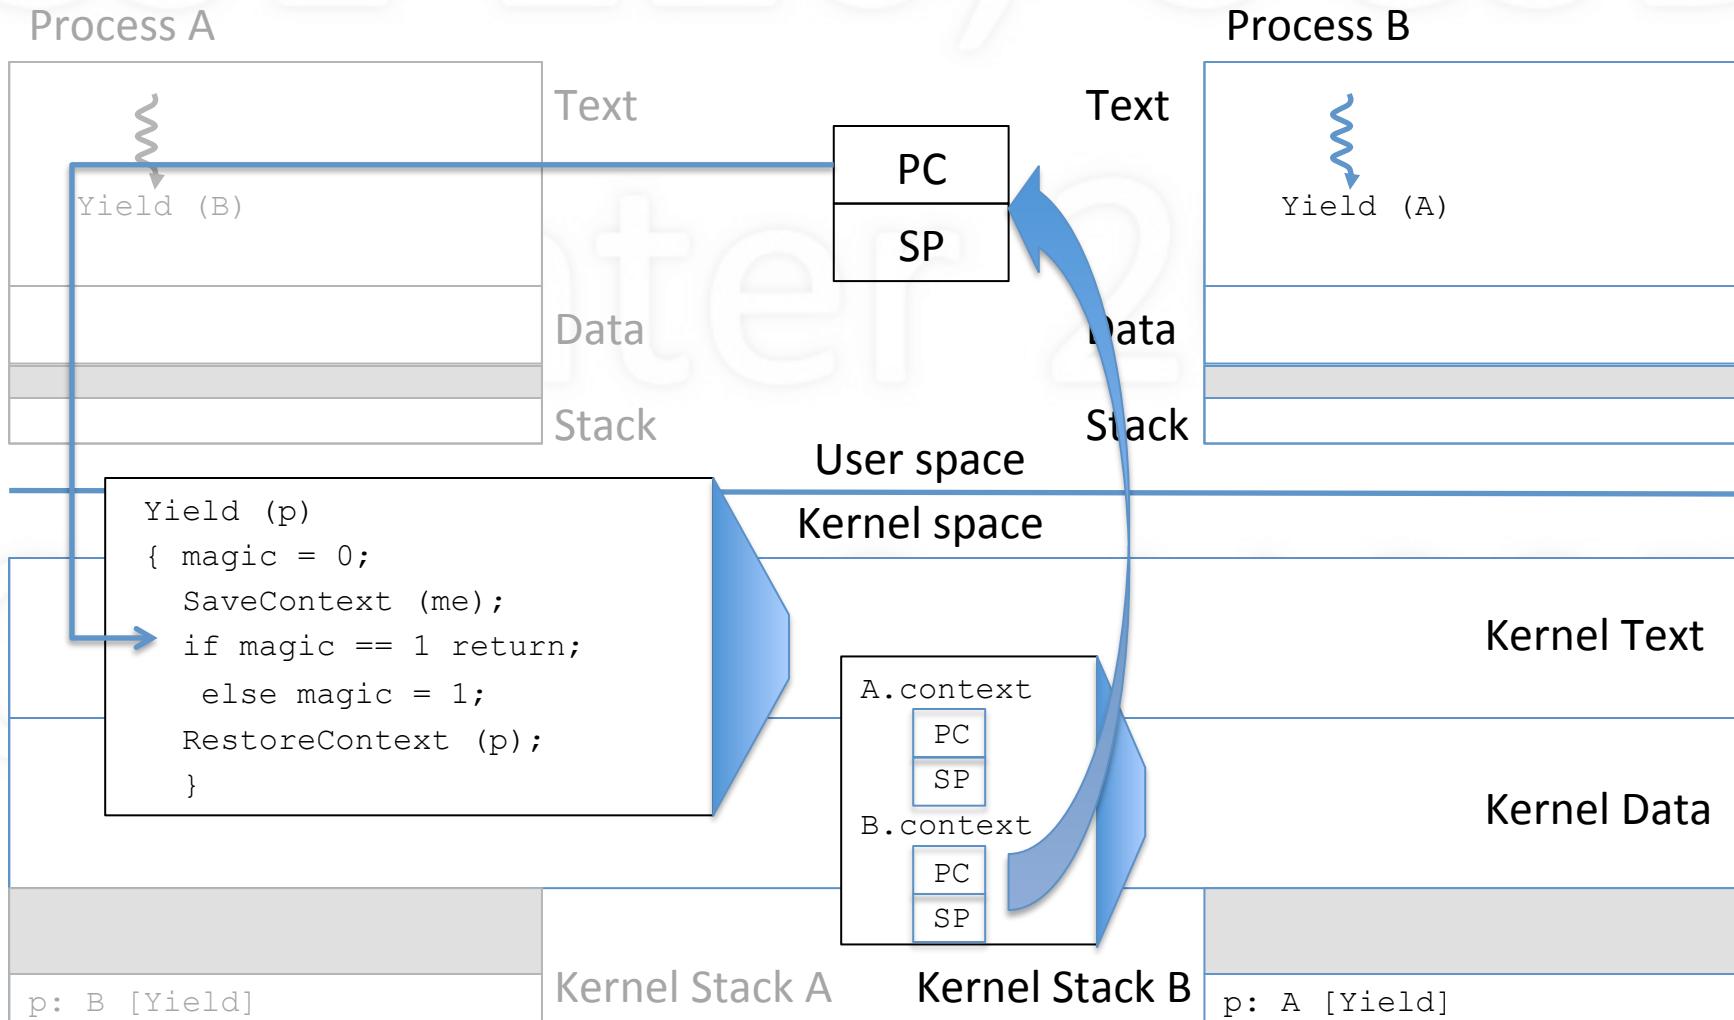
After Saving Context of A



About to Restore Context of B



After Restoring Context of B

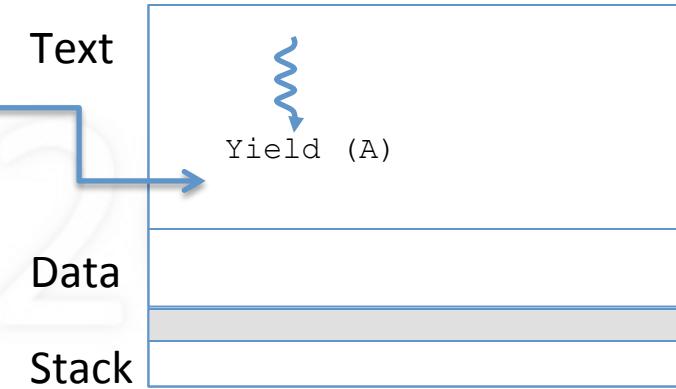


Return from Yield in B

Process A



Process B



User space

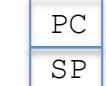
Kernel space

Kernel Text

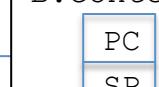
Kernel Data

```
Yield (p)
{ magic = 0;
  SaveContext (me);
  if magic == 1 return;
  else magic = 1;
  RestoreContext (p);
}
```

A.context



B.context



p: B [Yield]

Kernel Stack A

Kernel Stack B

Summary

- Process
 - abstraction of a running program
- Multiprogramming
 - allow for multiple processes (despite single CPU)
- Yield
 - when one process gives up CPU to another
- Context switching
 - mechanism that reassigns CPU between processes

Textbook

- OSP: Chapter 2
- OSC: Chapters 3, 4
 - Lecture-related: 3.1-3.3, 3.9, 4.1, 4.3, 4.8
 - Recommended: 4.2, 4.4-4.7