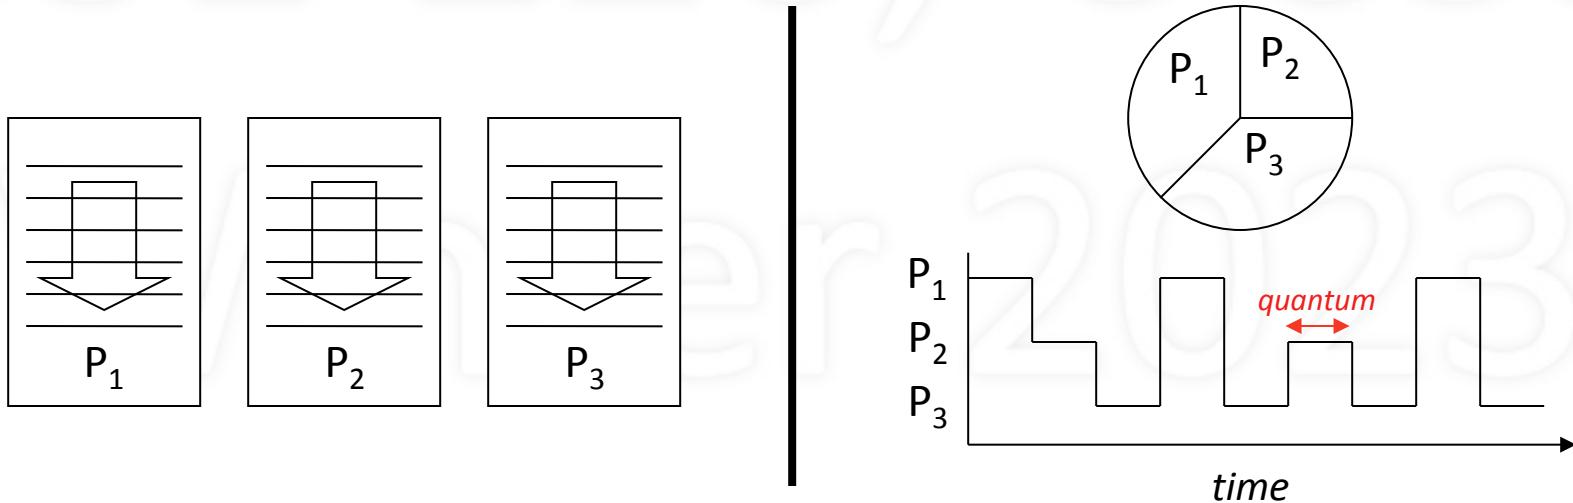


# CSE 120: Principles of Operating Systems

## Lecture 3: Timesharing

Prof. Joseph Pasquale  
University of California, San Diego  
January 18, 2023

# Timesharing

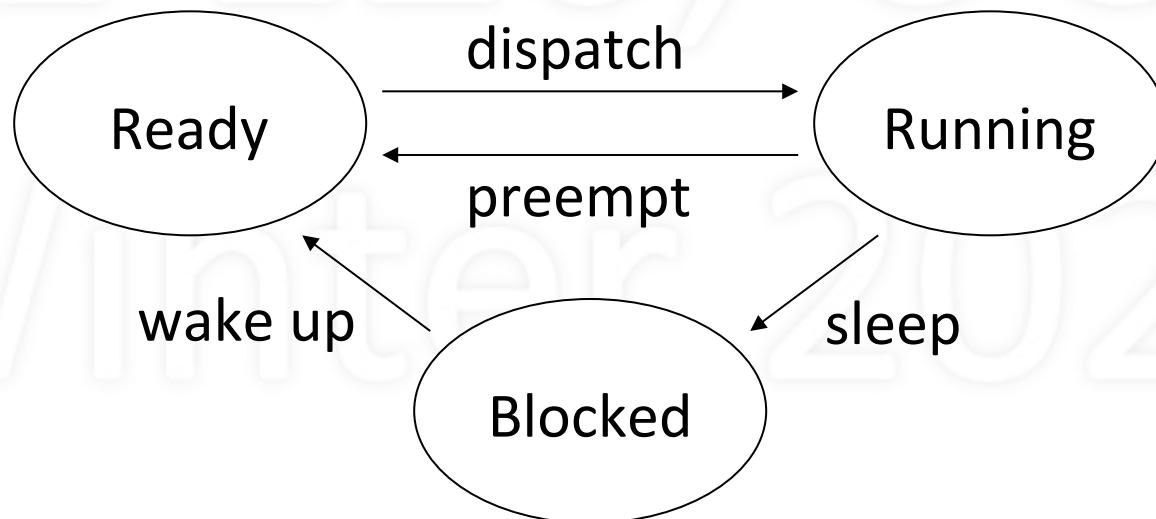


- Timesharing: multiplexing use of CPU over time
- Multiple processes, single CPU (uniprocessor)
- Conceptually, each process makes progress over time
- In reality, each periodically gets quantum of CPU time
- Illusion of parallel progress by rapidly switching CPU

# How is Timesharing Implemented?

- Kernel keeps track of progress of each process
- Characterizes state of process's progress
  - Running: actually making progress, using CPU
  - Ready: able to make progress, but not using CPU
  - Blocked: not able to make progress, can't use CPU
- Kernel selects a ready process, lets it run
  - Eventually, the kernel gets back control
  - Selects another ready process to run, ...

# Process State Diagram



- State transitions
  - Dispatch: allocate the CPU to a process
  - Preempt: take away CPU from process
  - Sleep: process gives up CPU to wait for event
  - Wakeup: event occurred, make process ready

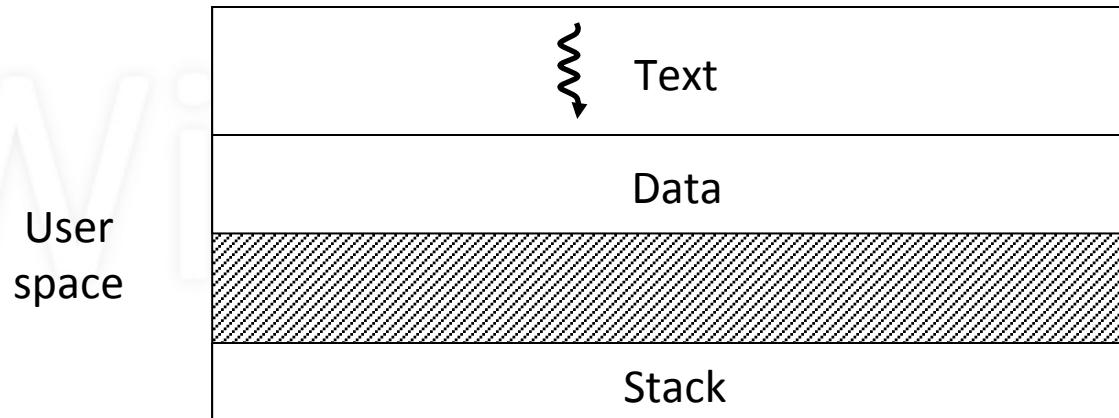
# Logical vs. Physical Execution

		<u>Logical Execution</u>	
		Able to execute	Not able to execute
<u>Physical Execution</u>	Actually executing	Run	X
	Not actually executing	Ready	Blocked

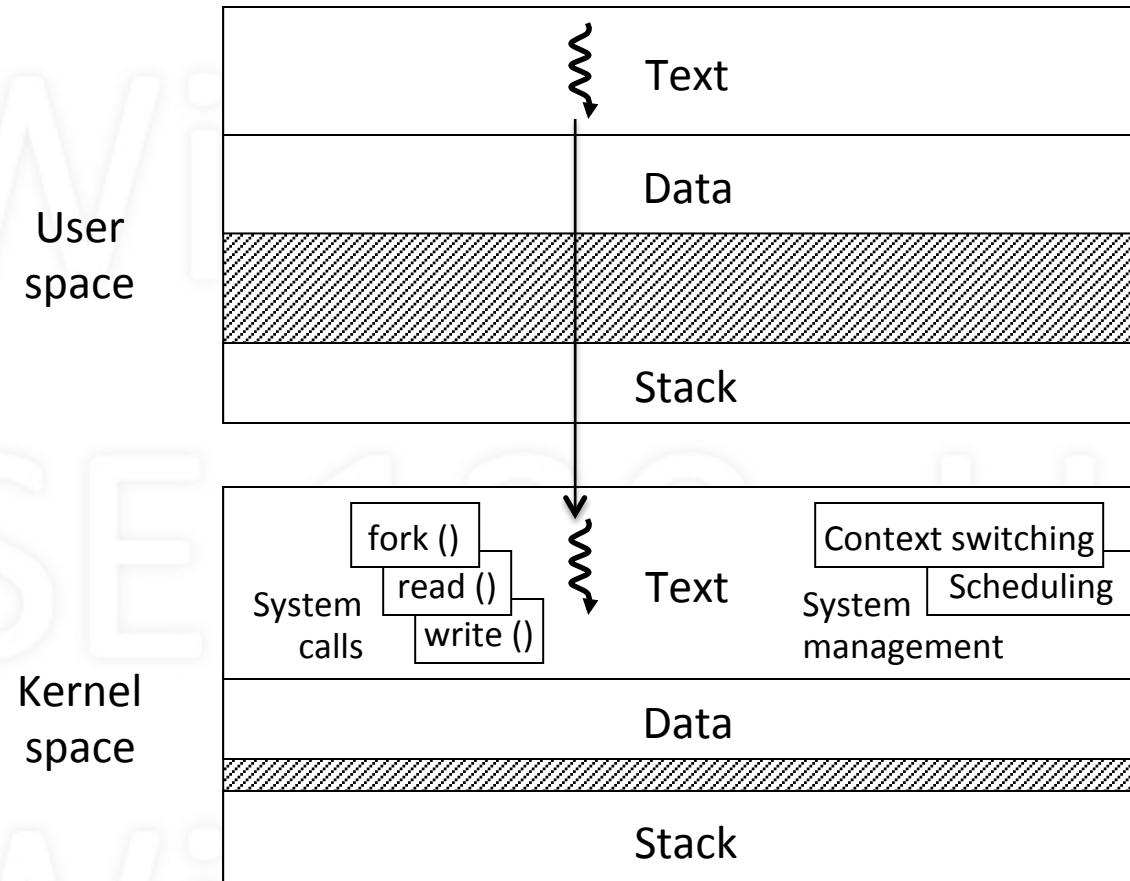
# Process vs. Kernel

- Kernel: code that supports processes
  - system calls: fork ( ), exit ( ), read ( ), write ( ), ...
  - management: context switching, scheduling, ...
- When does the kernel run?
  - when system call or hardware interrupt occurs
- The kernel runs as part of the running process
  - due to that process having made a system call
  - in response to device issuing interrupt

# Process Running in User Space



# Process Running in Kernel Space



# Kernel Maintains List of Processes

Process ID	State	Other info
1534	Ready	Saved context, ...
34	Running	Memory areas used, ...
487	Ready	Saved context, ...
9	Blocked	Condition to unblock, ...

- All processes: unique names (IDs) and states
- Other info kernel needs for managing system
  - contents of CPU contexts
  - areas of memory being used
  - reasons for being blocked

# How Does Kernel Get Control

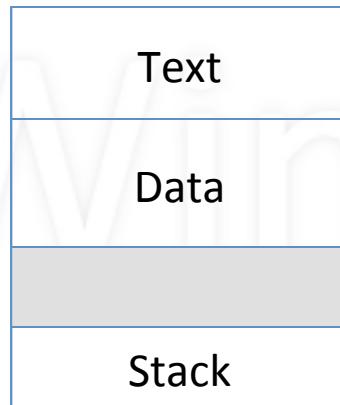
- Process can give up control voluntarily
  - Makes system call that blocks, e.g., `read()`
  - System-call function calls `yield()` to give up CPU
  - Kernel selects a ready process, dispatches it
- Or, CPU is forcibly taken away: *preemption*
  - Interrupt generated when hardware timer expires
  - Interrupt forces control to go to kernel
  - While kernel running, resets timer for next time

# How a Context Switch Occurs

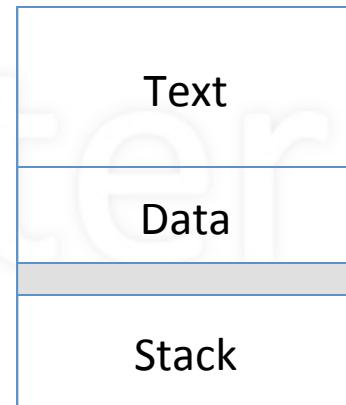
- Process makes system call or interrupt occurs
- What's done by hardware
  - Switch from user to kernel mode: amplifies power
  - Go to fixed kernel location: trap/interrupt handler
- What's done in software (in the kernel)
  - Save context of current process
  - Select a process that is ready; restore its context
  - RTI: return from interrupt/trap

# Example

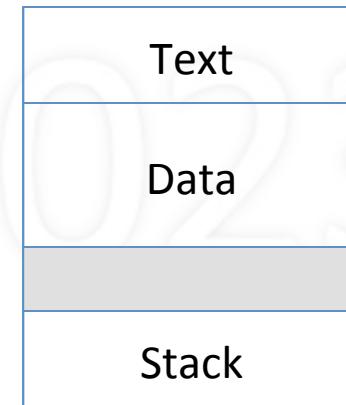
Process 1



Process 2



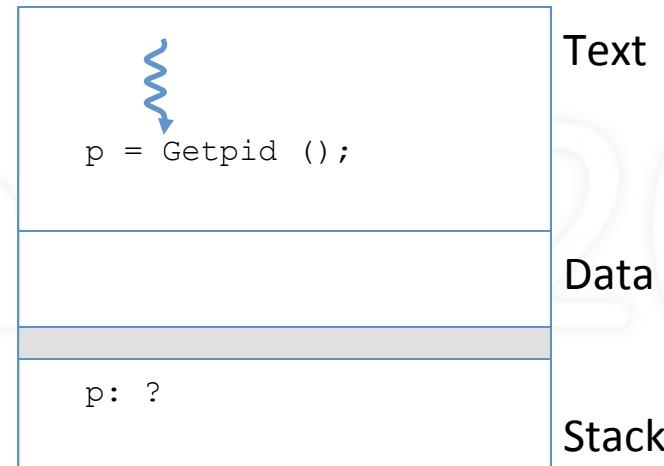
Process n



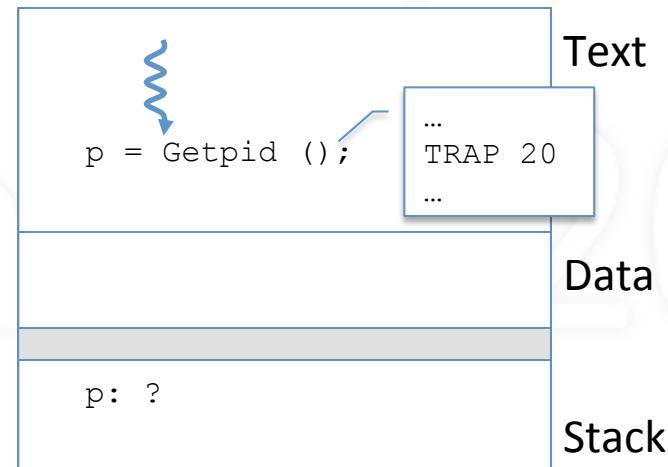
...

Kernel

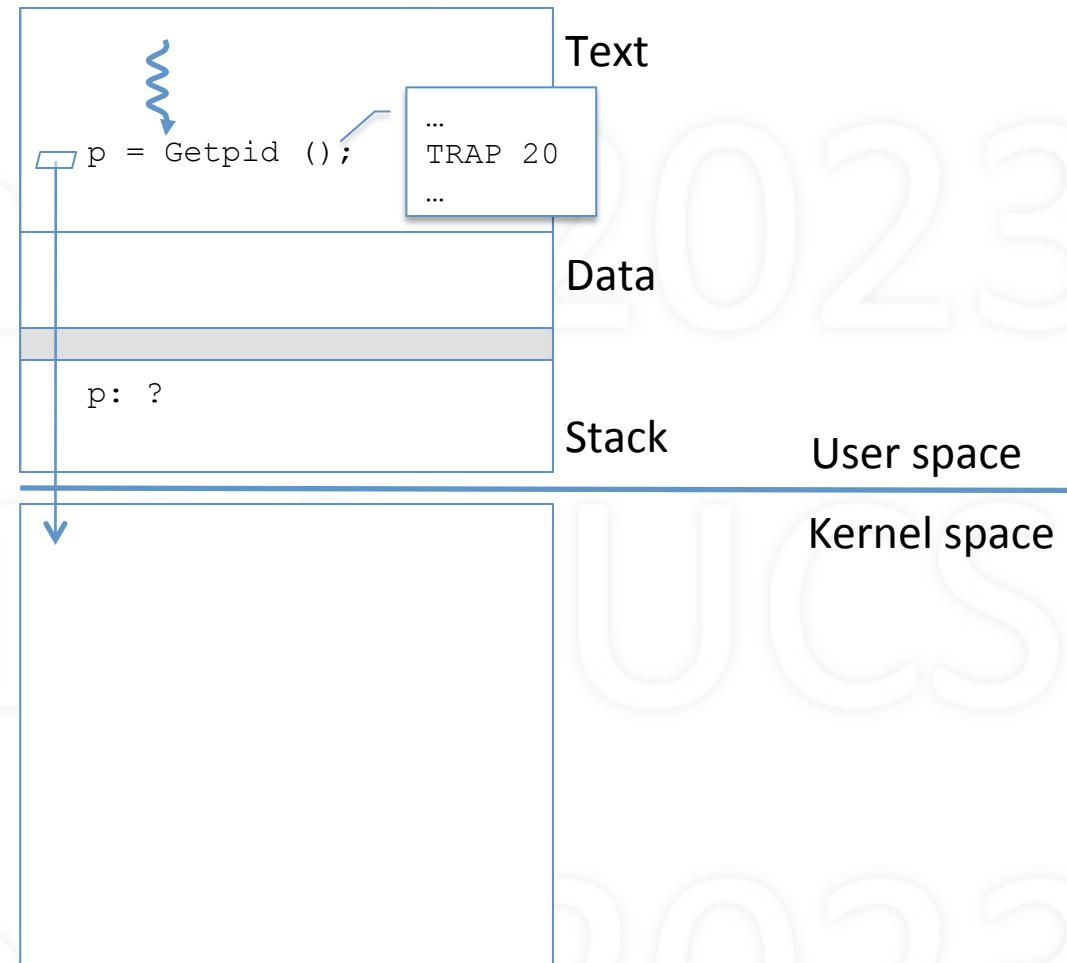
# Process makes system call



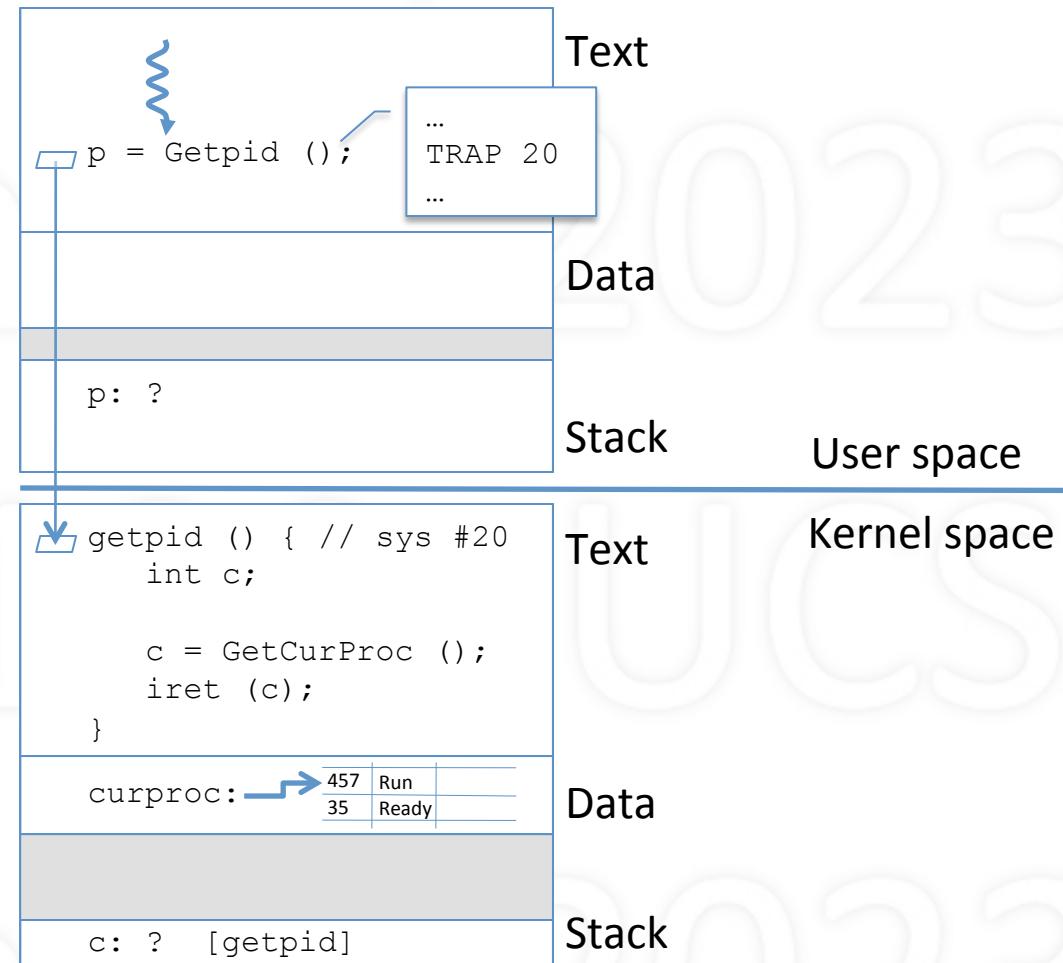
# Translates into a TRAP



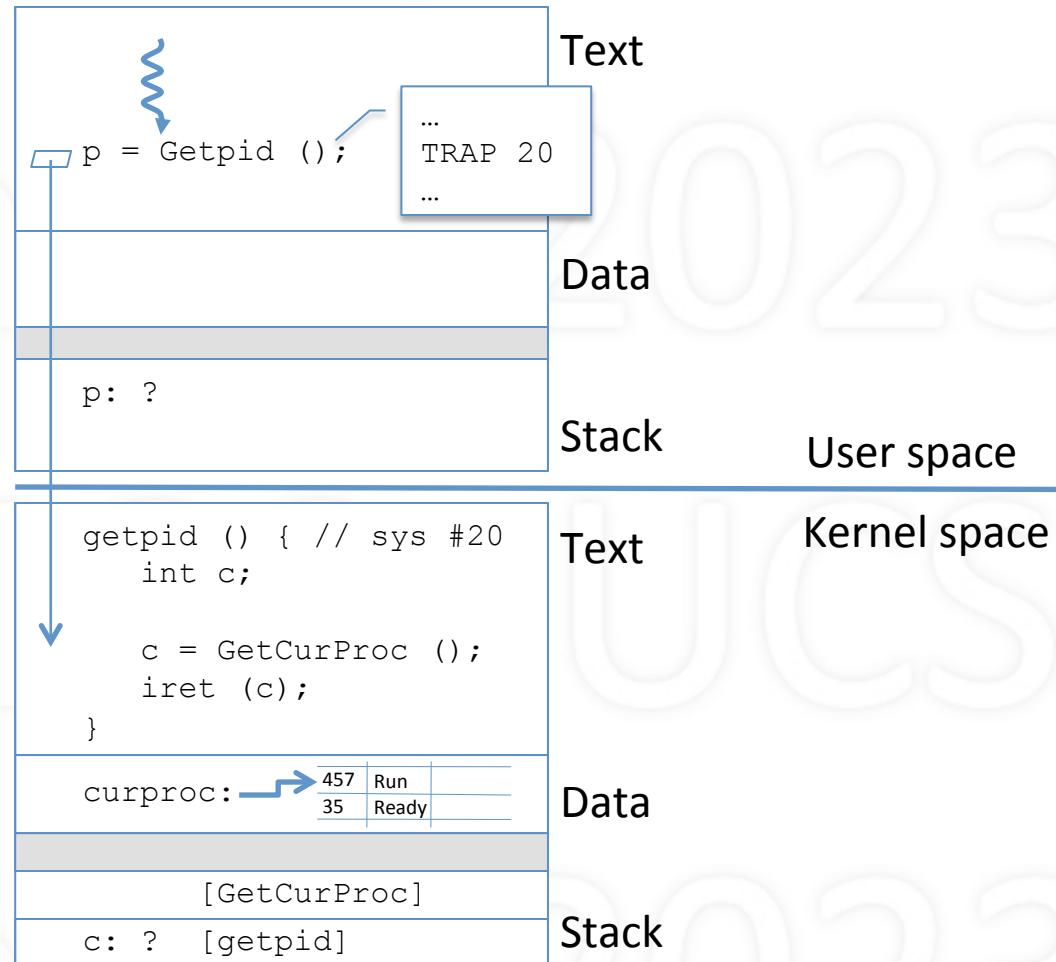
# TRAP causes entry into kernel



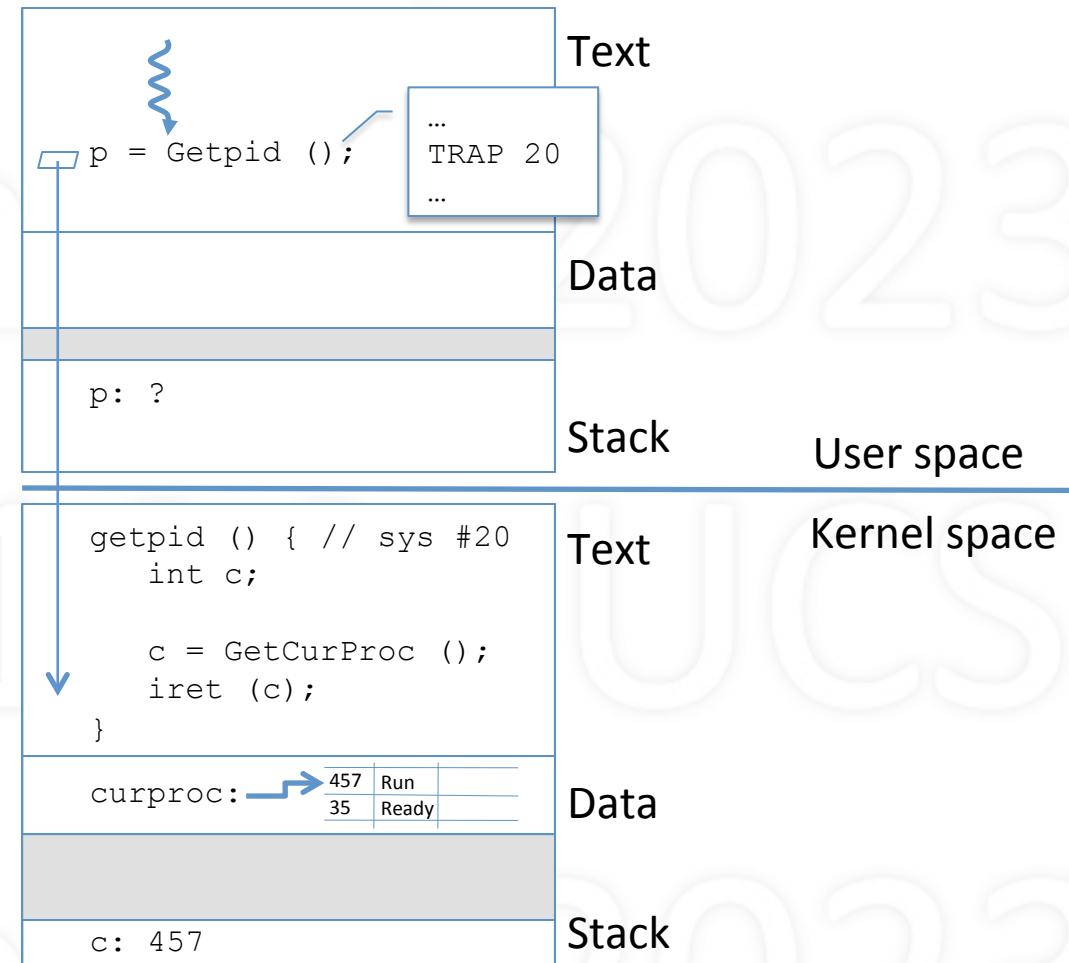
# Jump to proper system function



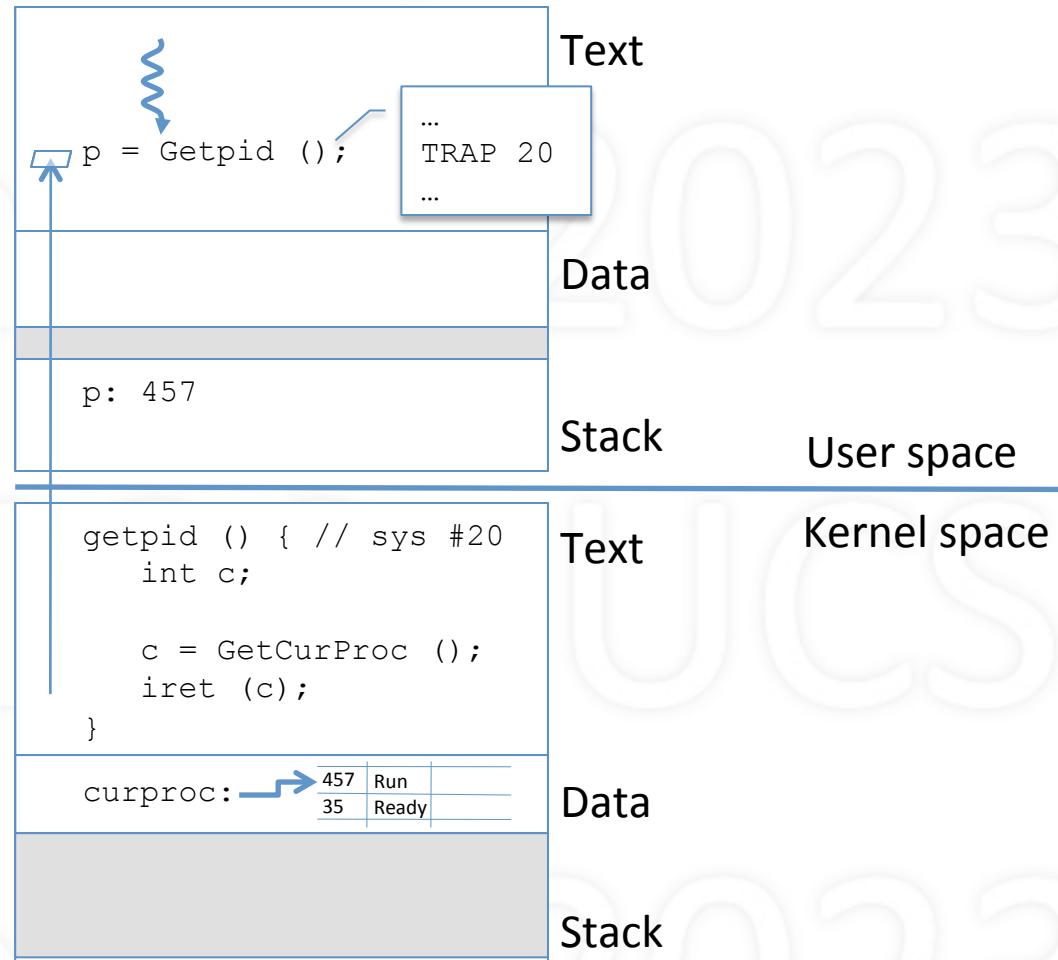
# System function starts executing



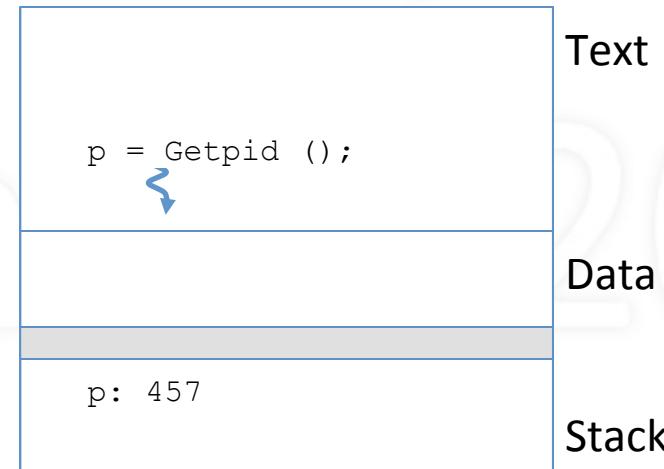
# System function completes



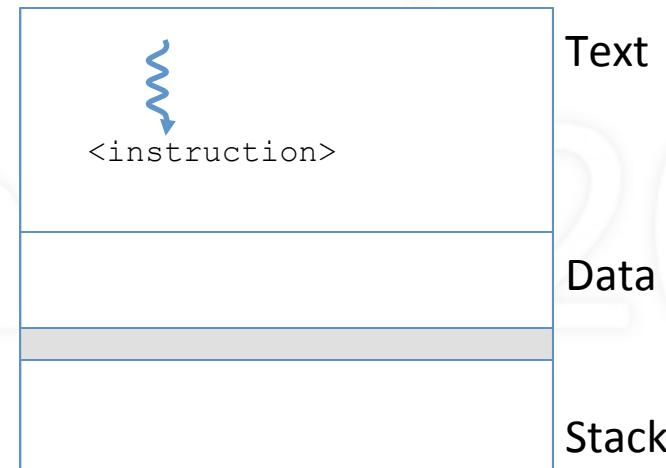
# Return from kernel to process



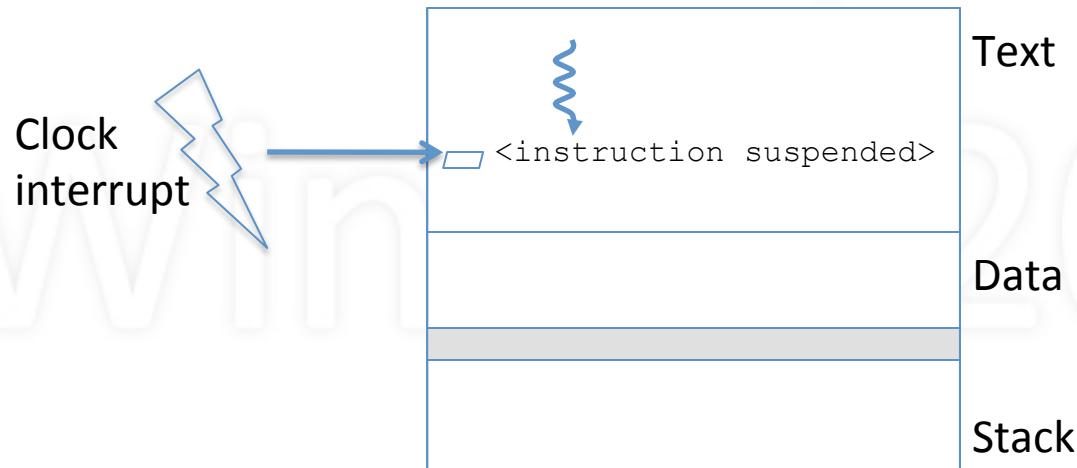
# Process continues after system call



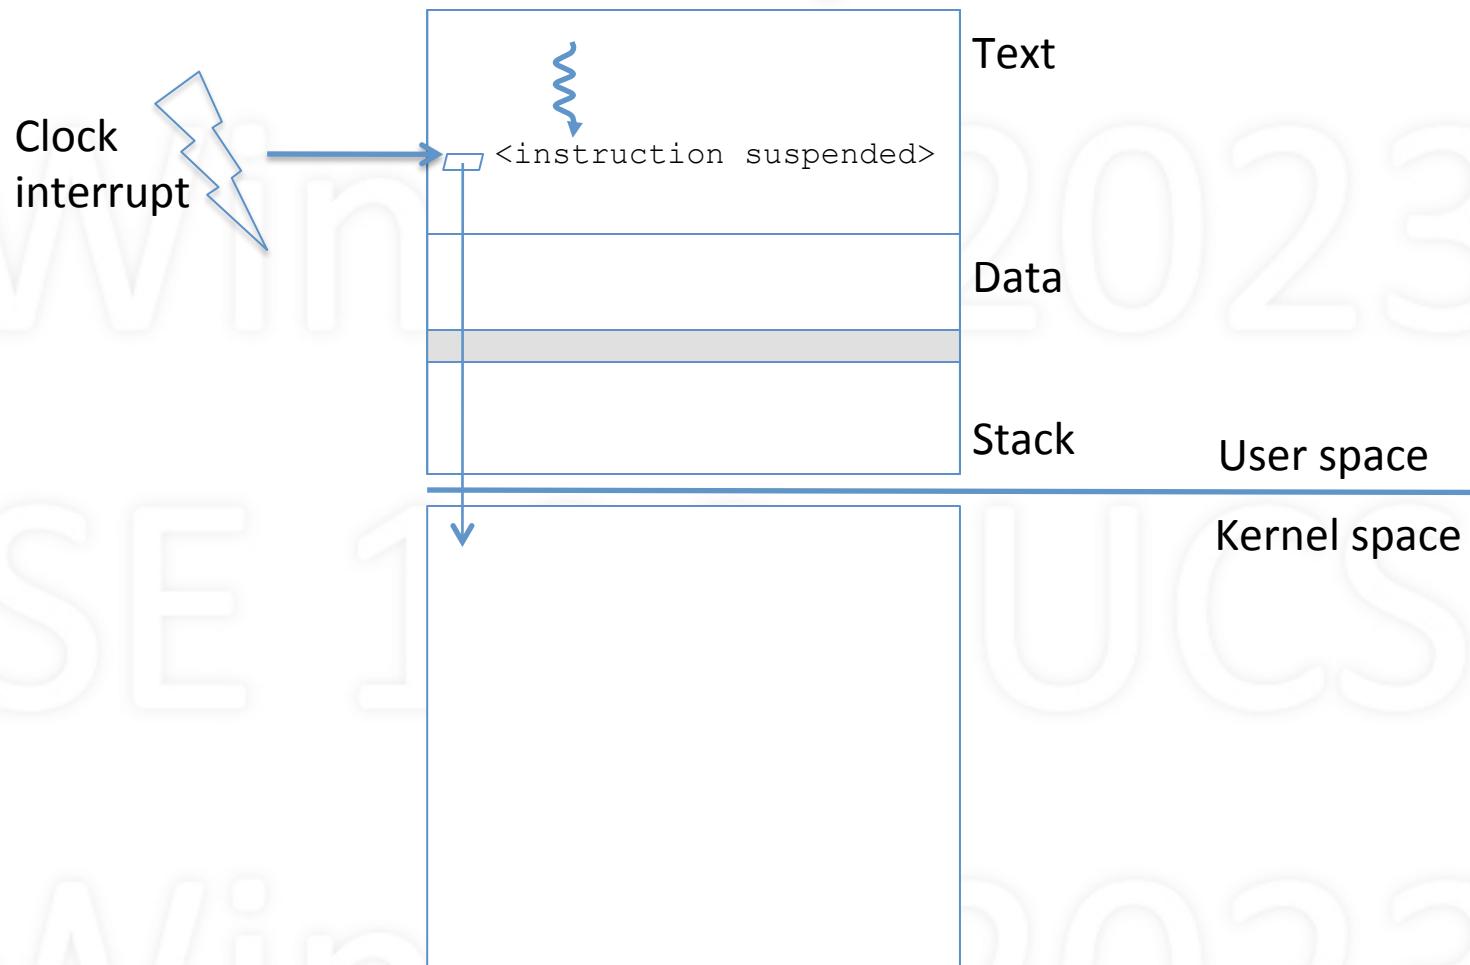
# Kernel entry via clock interrupt



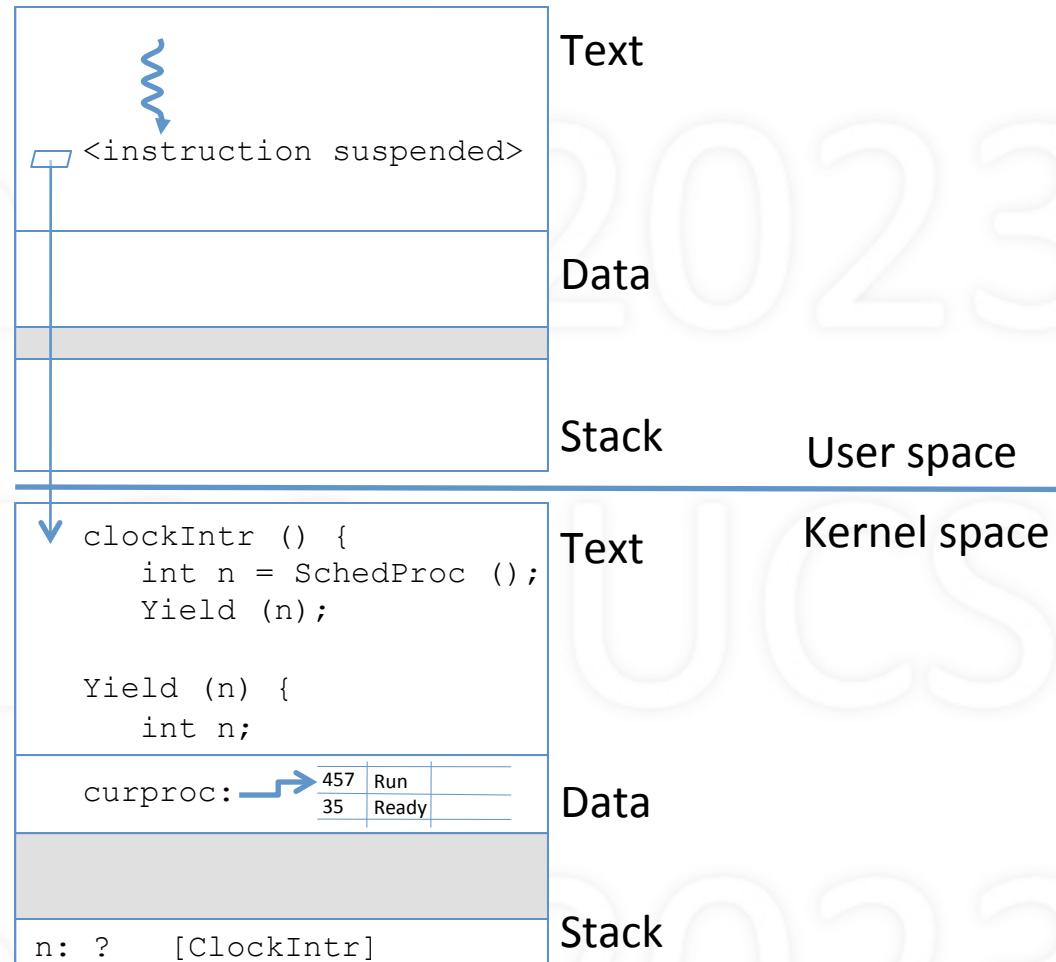
# Asynchronous clock interrupt



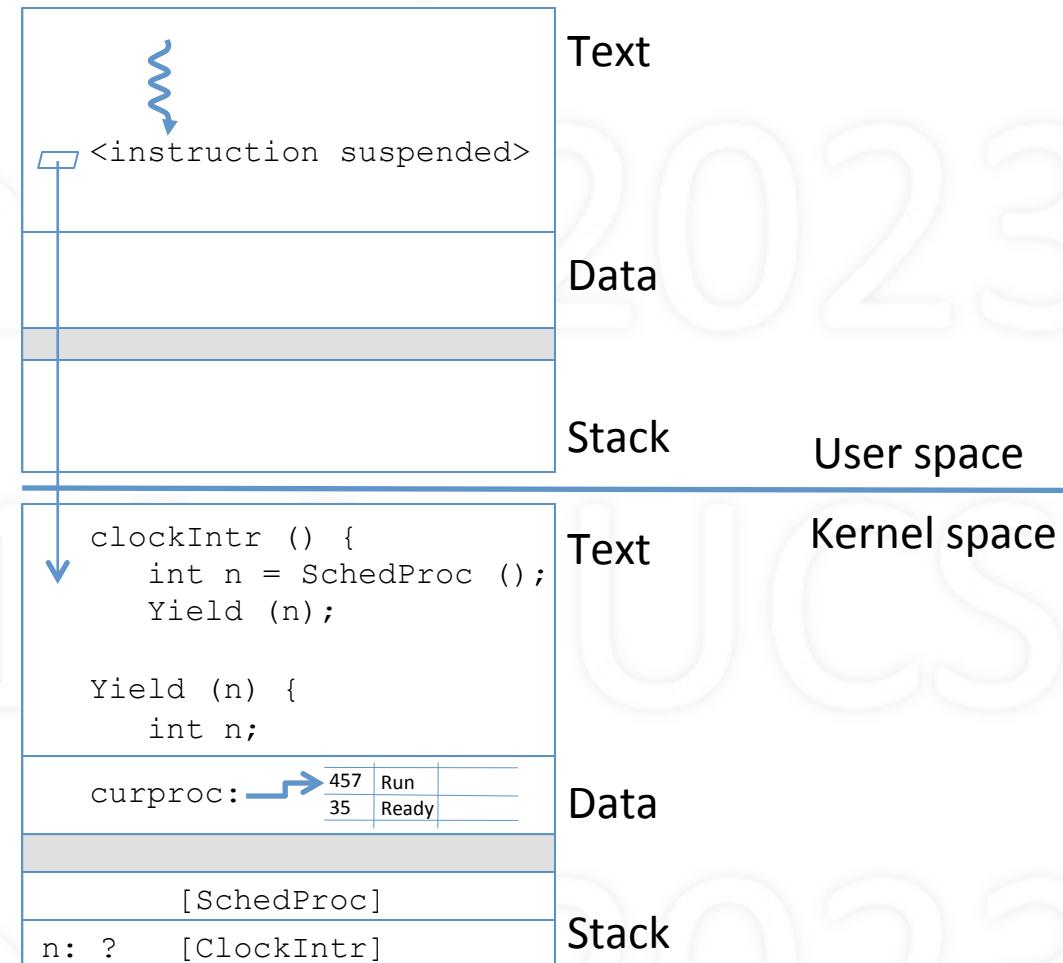
# Causes forced kernel entry



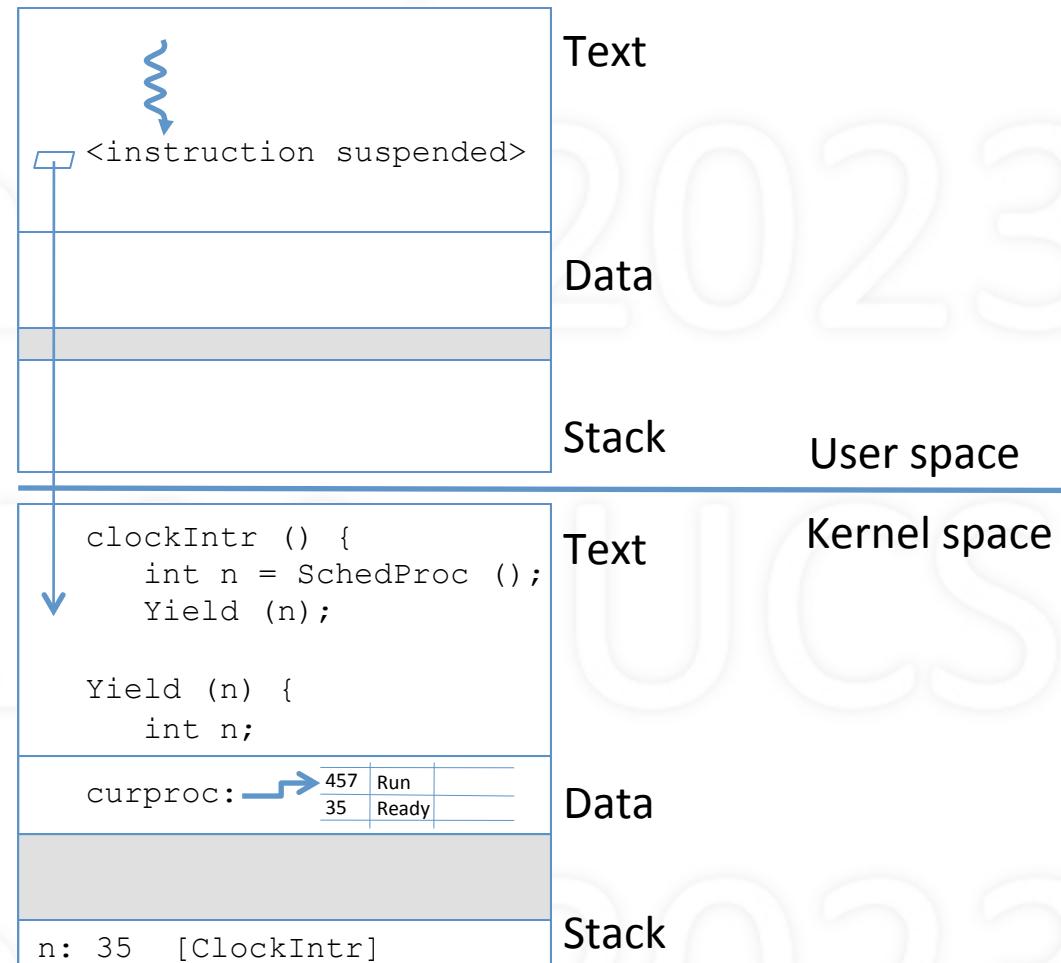
# Clock interrupt handler executes



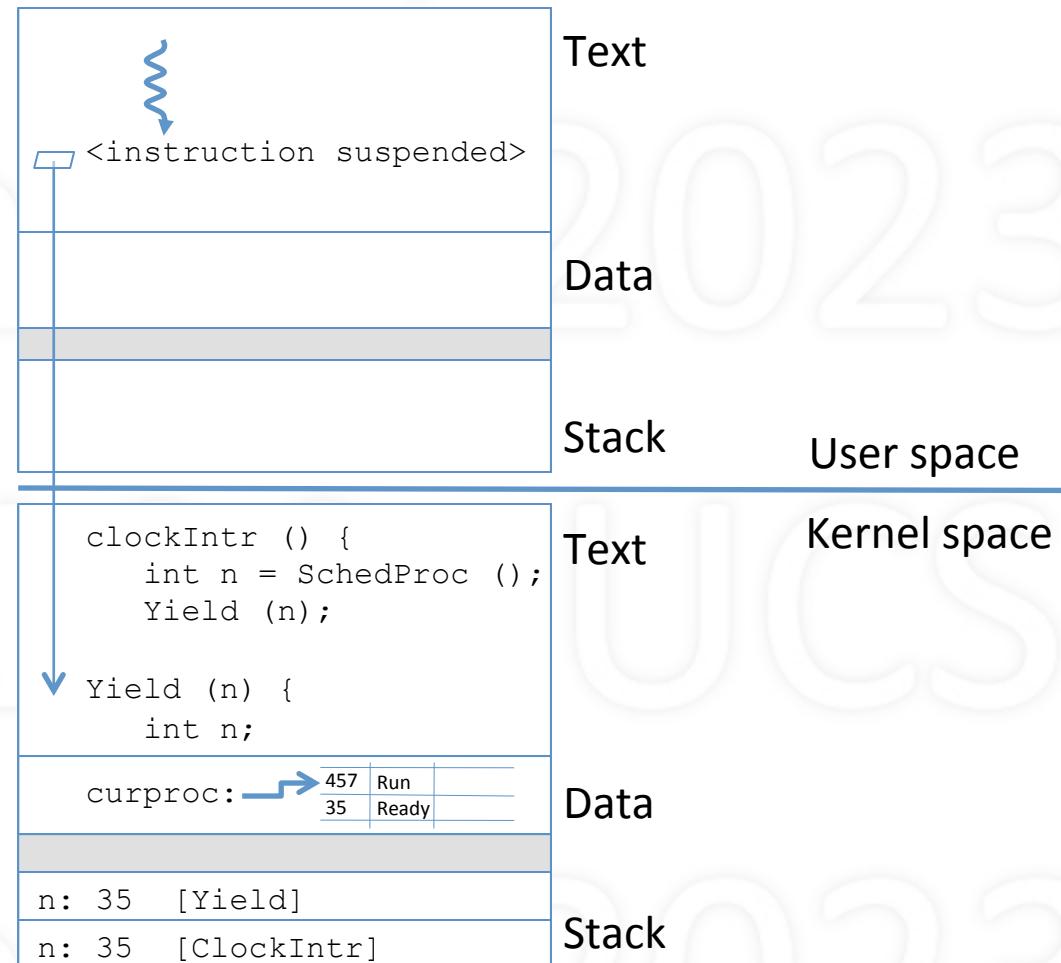
# Decide which process to run next



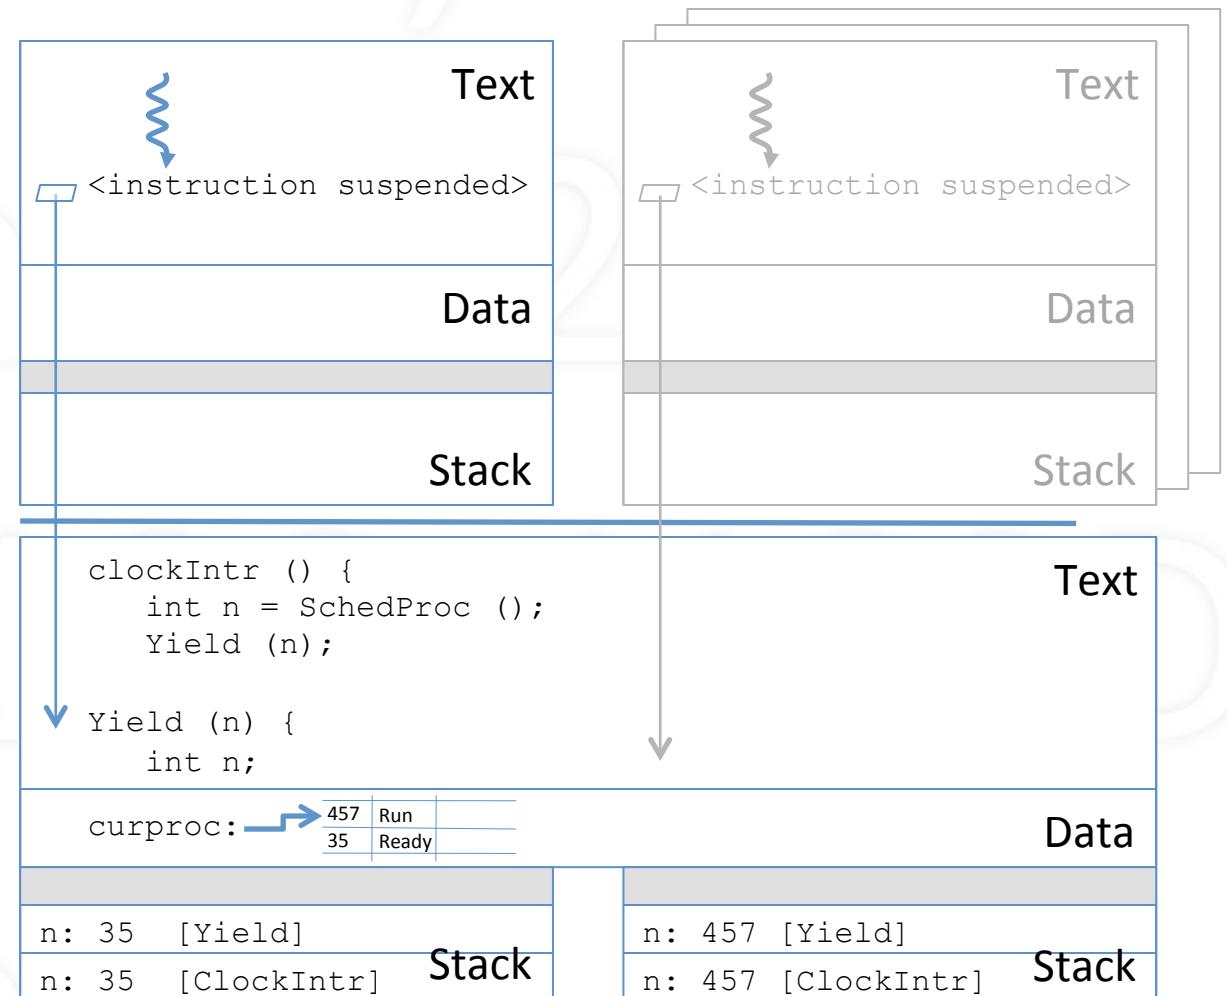
# Yield to that process



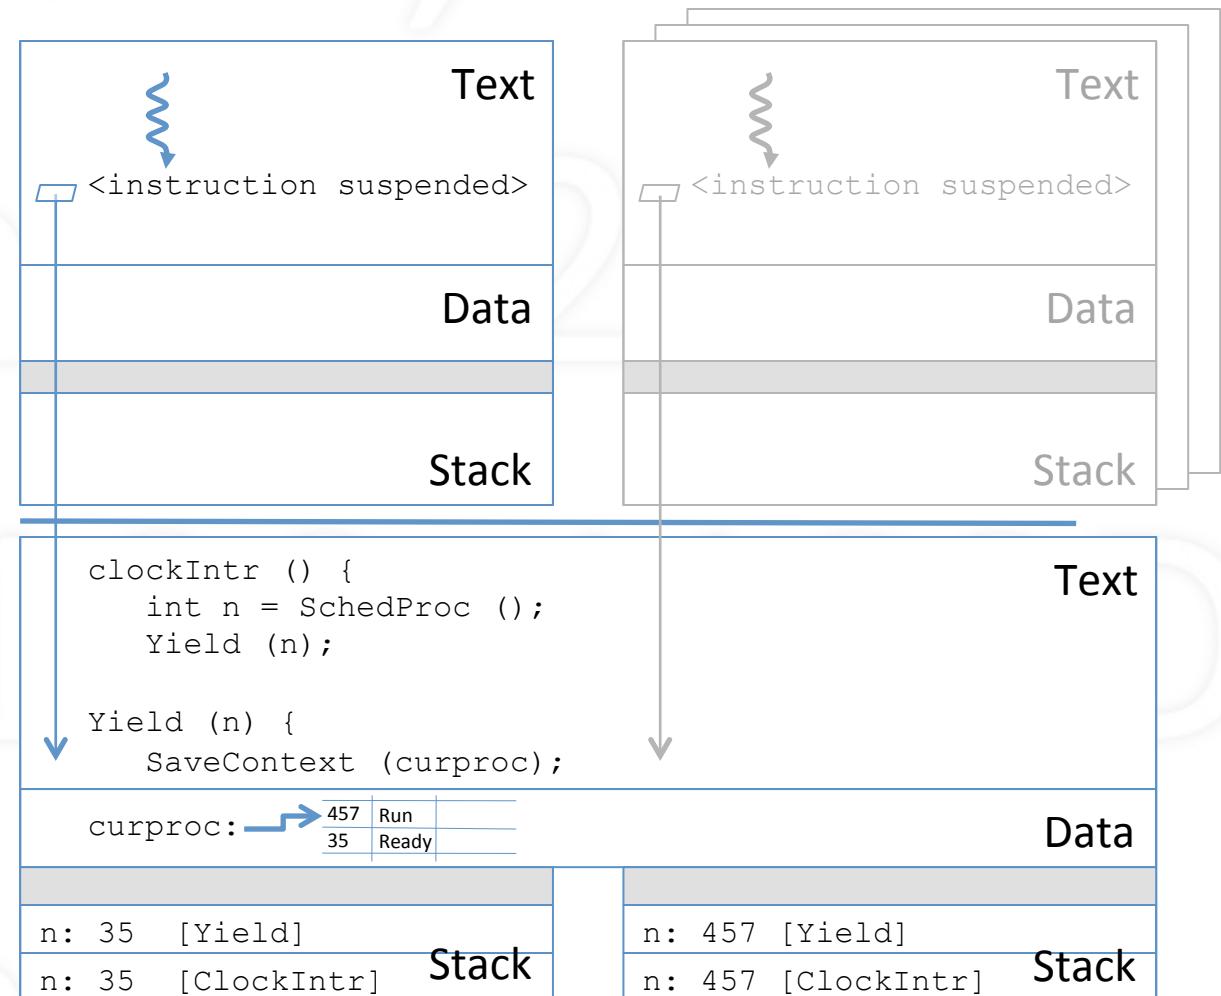
# Call to yield



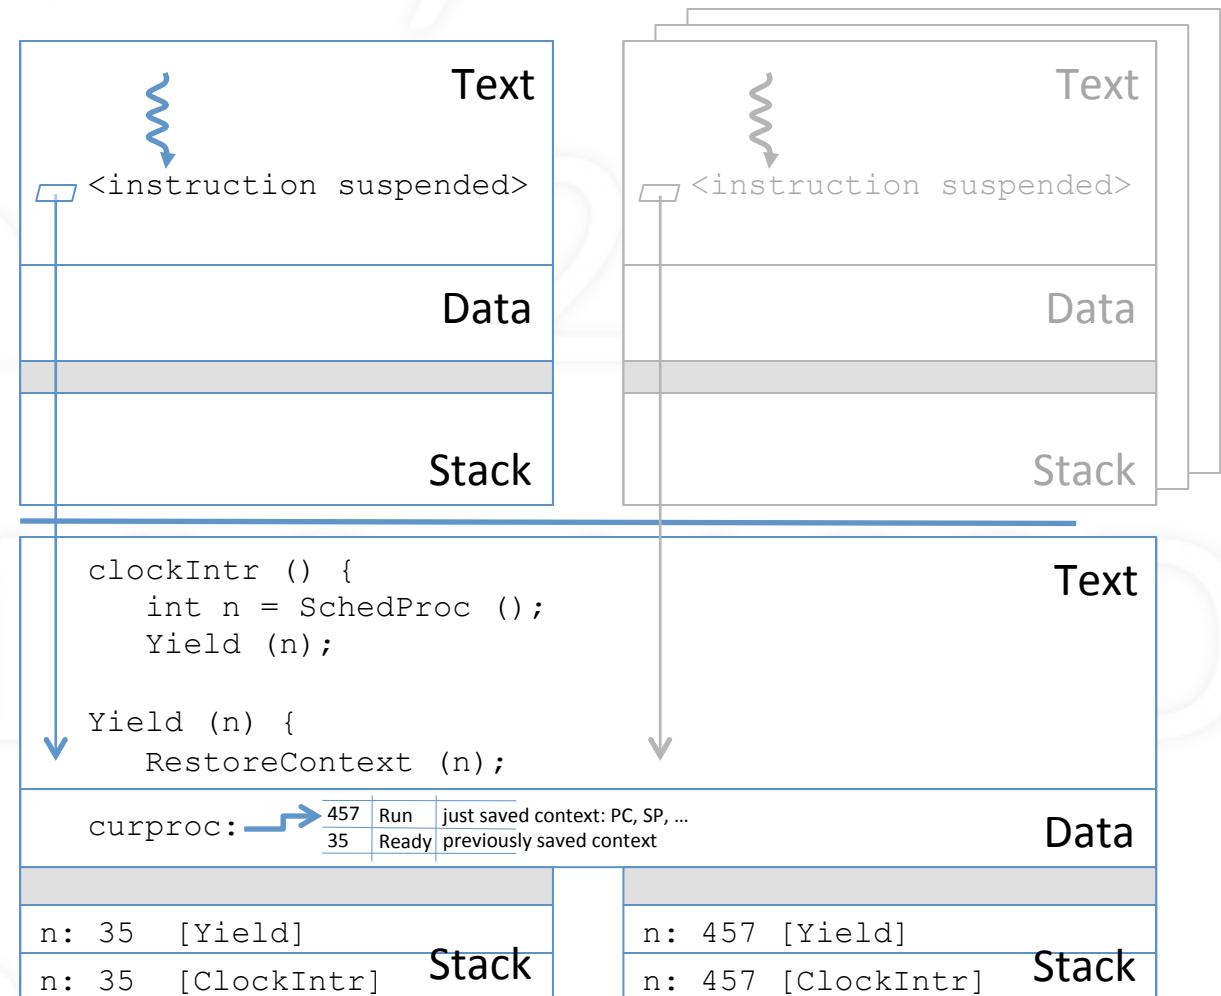
# Yield magic



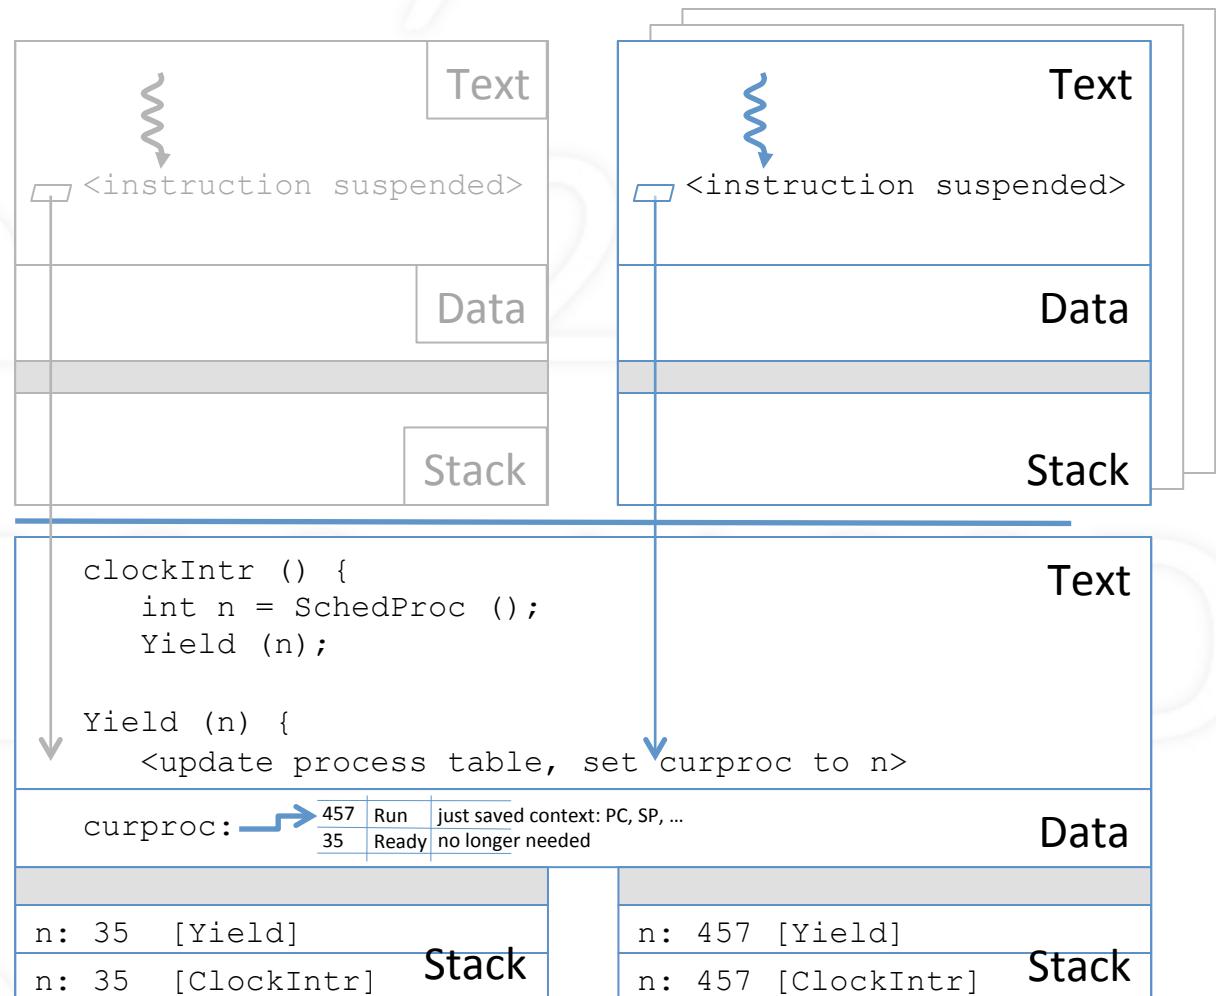
# Save context



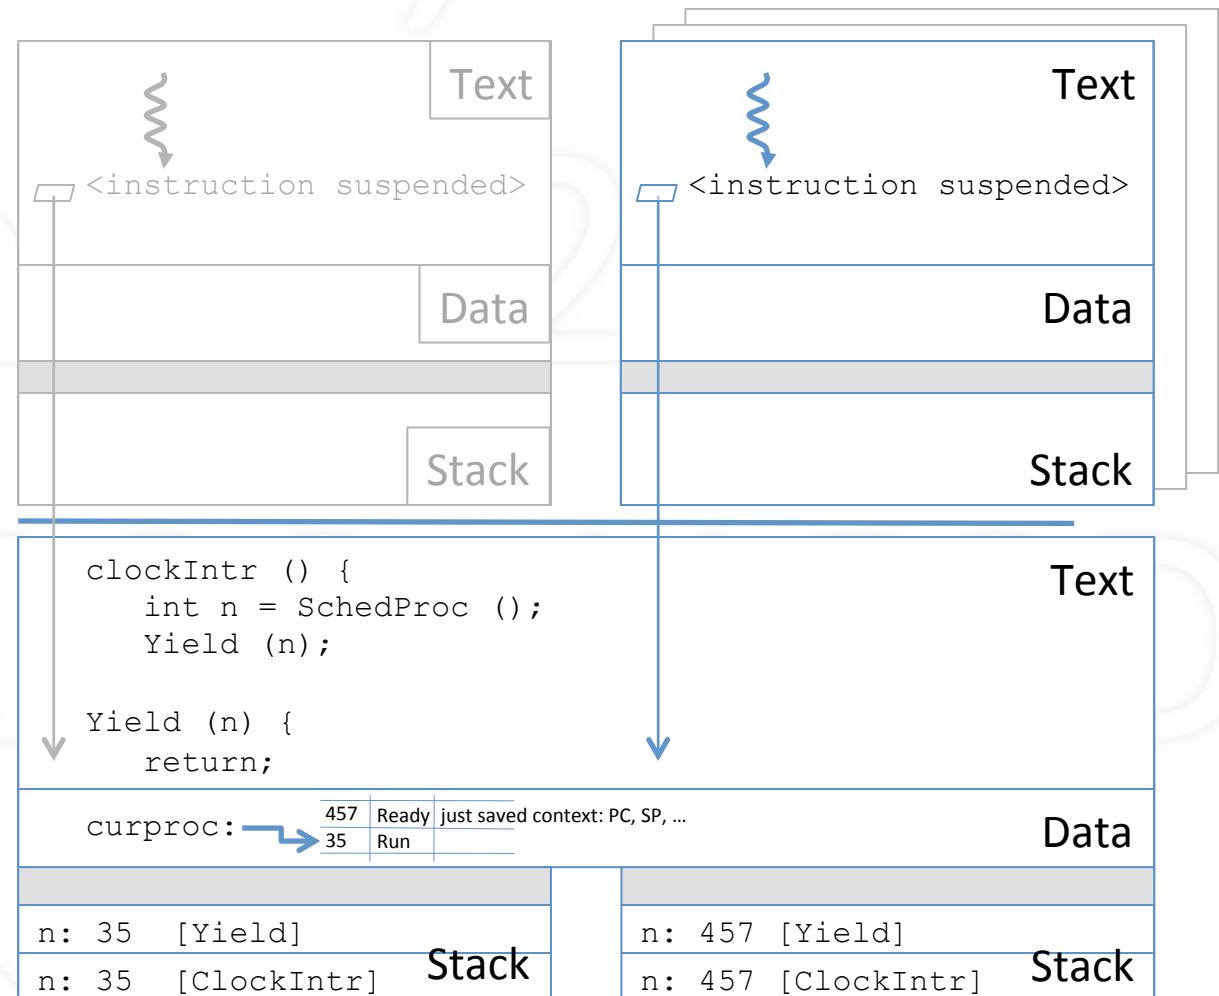
# Restore context of next process



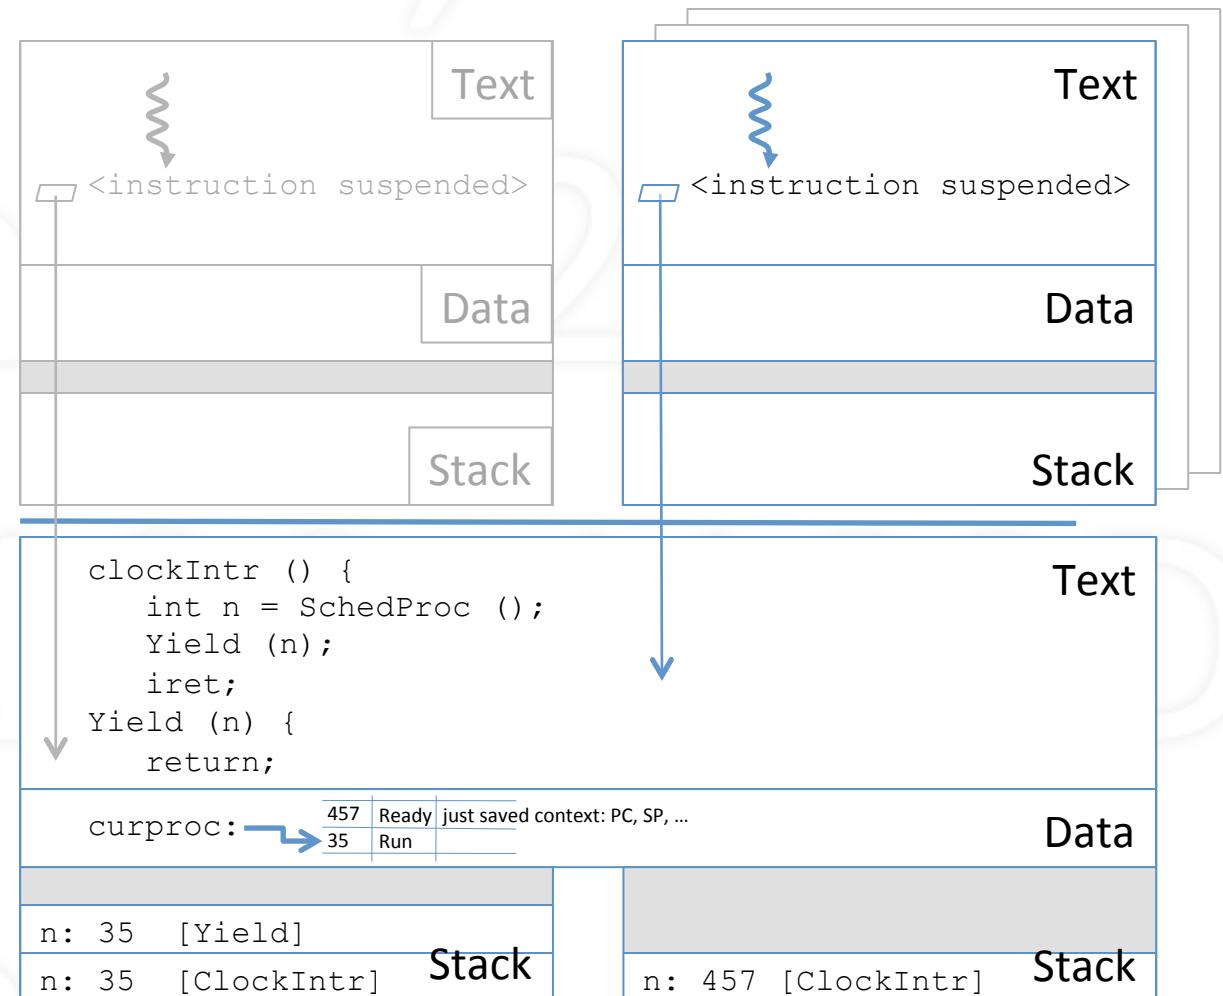
# Next process almost running



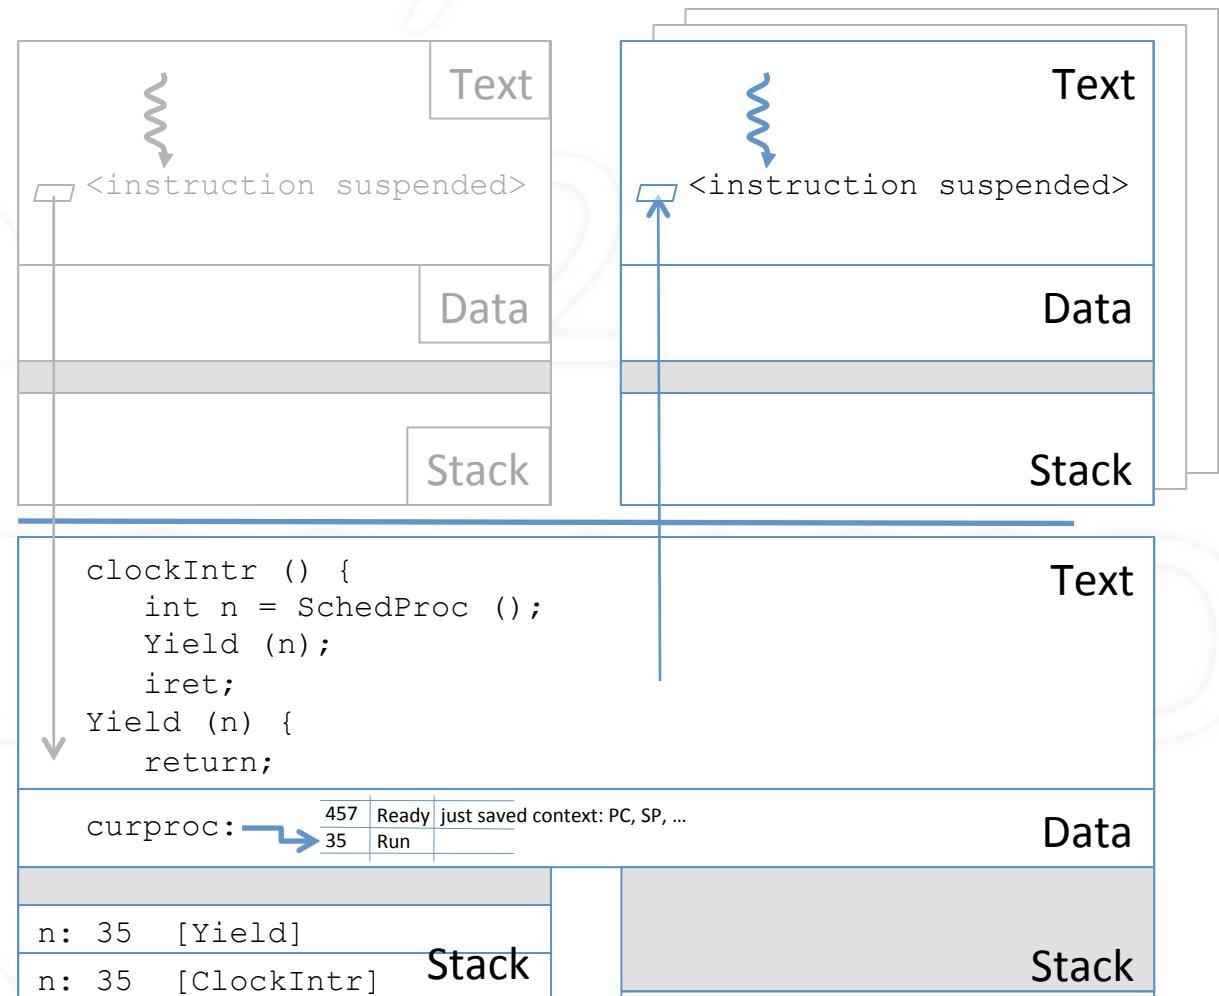
# Update process table



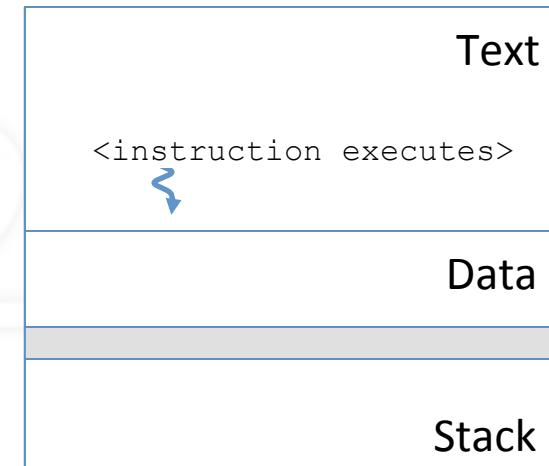
# About to return from Yield



# Return to process code

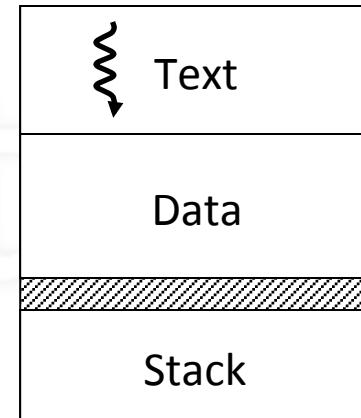


# Resumed process now running

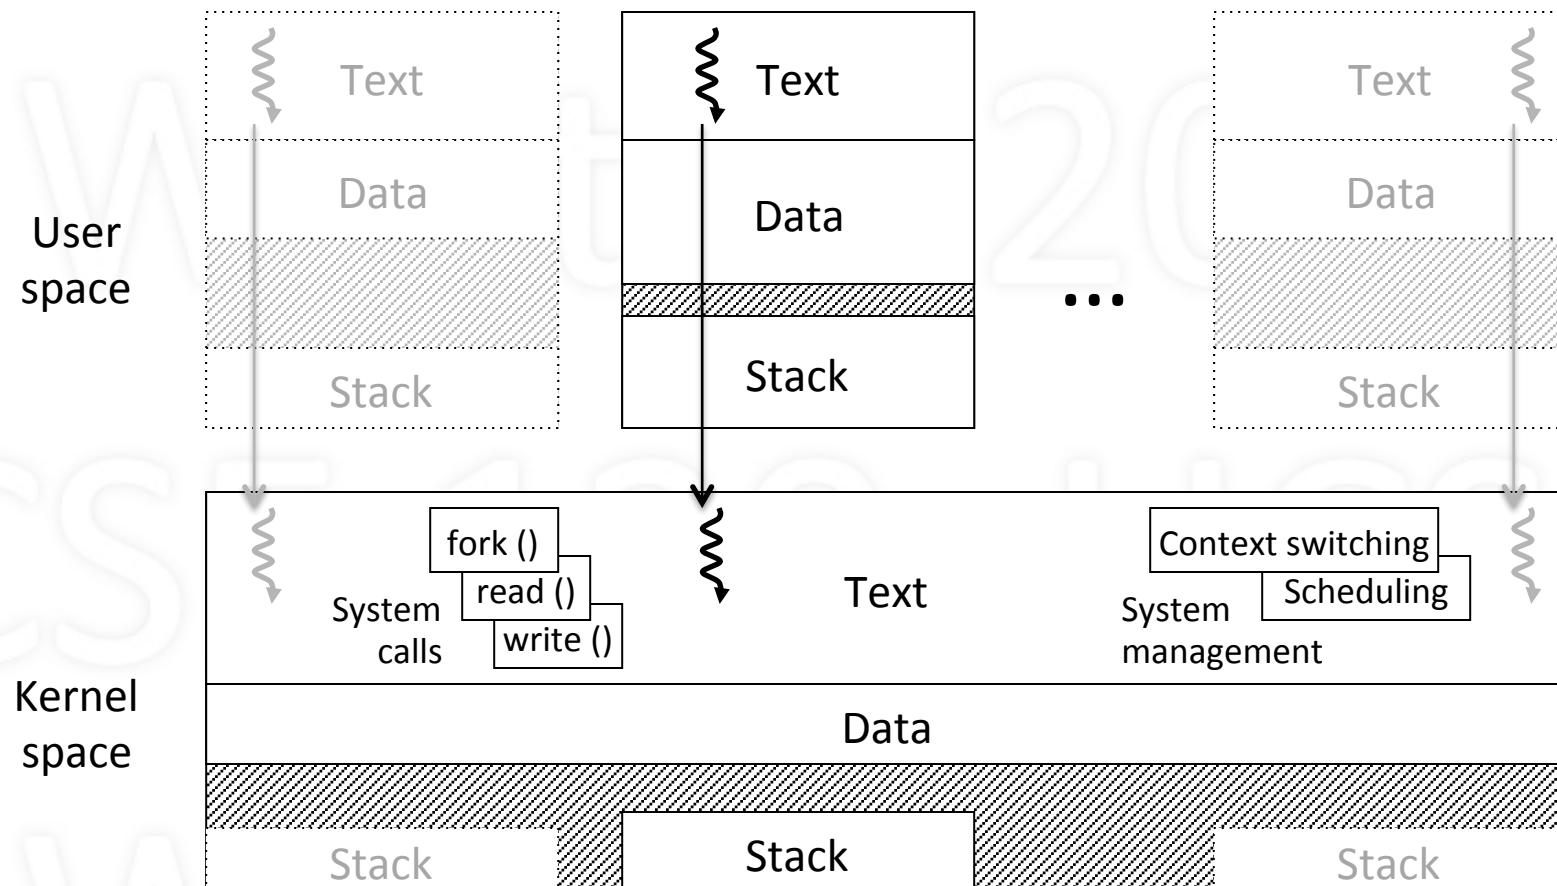


# Process Running in User Space

User  
space



# Process Running in Kernel Space



# How to Get Parallelism in Process

- Process is a “program in execution”
  - assumed (so far) a single path of execution
  - in a memory composed of text, data, stack
- What if we want multiple paths of execution?
  - Single text, but multiple executions in parallel
  - Single data, any execution can see others’ updates
  - Need separate stacks: one per ongoing execution
- Multiple processes? No (separate memories)

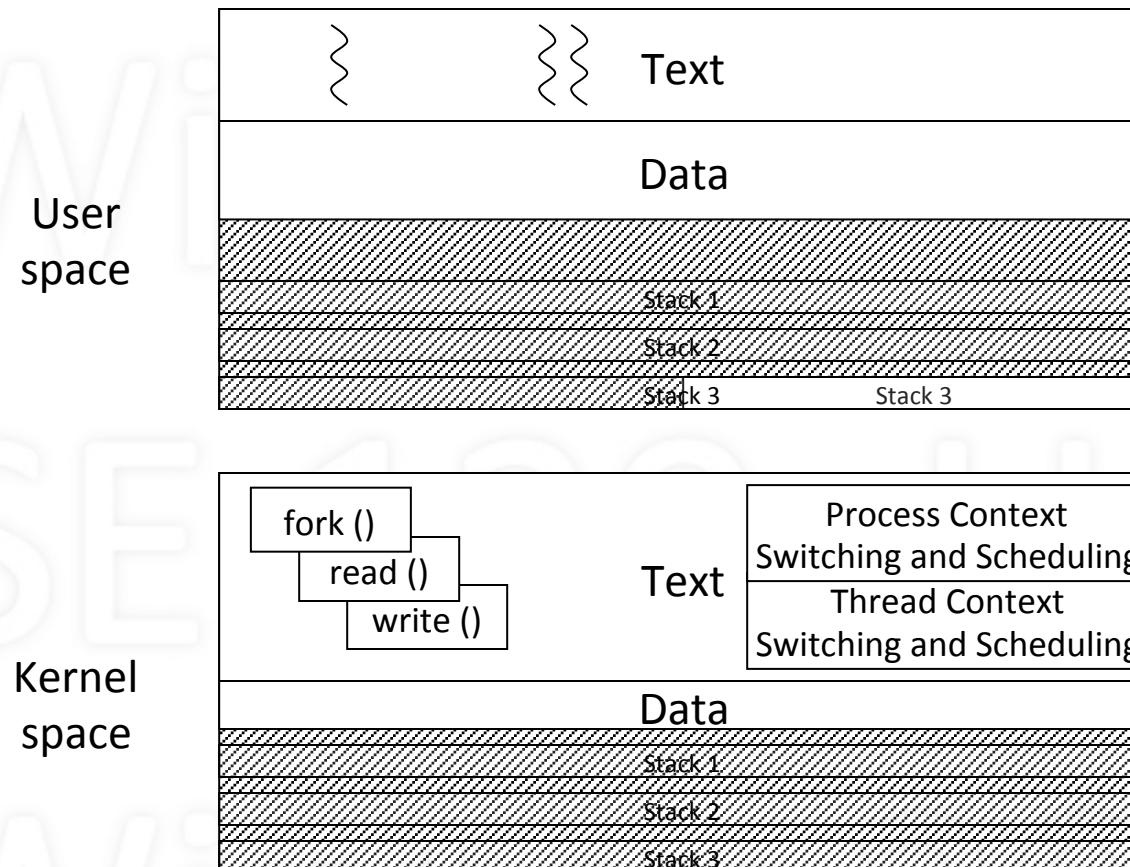
# Threads

- Thread: single sequential path of execution
- Abstraction is independent of memory
  - Contrast to process: path of execution + memory
- A thread is part of a process
  - Lives in the memory of a process
  - Distinction allows multiple threads in a process
- To the user: unit of parallelism
- To the kernel: unit of schedulability

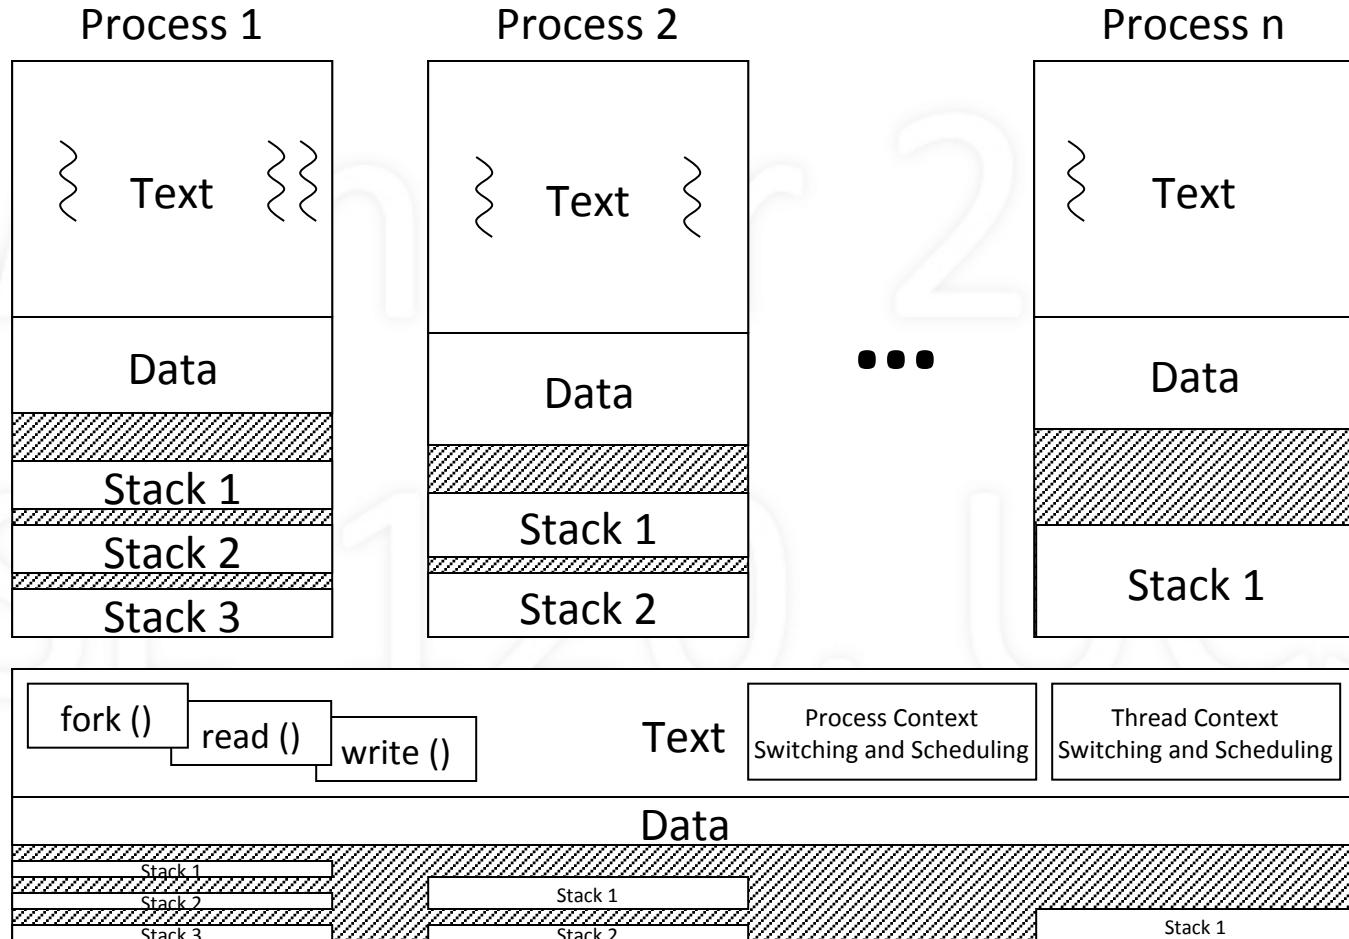
# Implementing Threads

- Thread calls are system calls
  - ForkThread (): like process Fork () but for threads
  - Thread system call functions are in kernel
- Thread management functions are in kernel
  - Thread context switching
  - Thread scheduling
- Each thread requires user and kernel stacks
- Kernel can schedule threads on separate CPUs

# Single Process, Multiple Threads



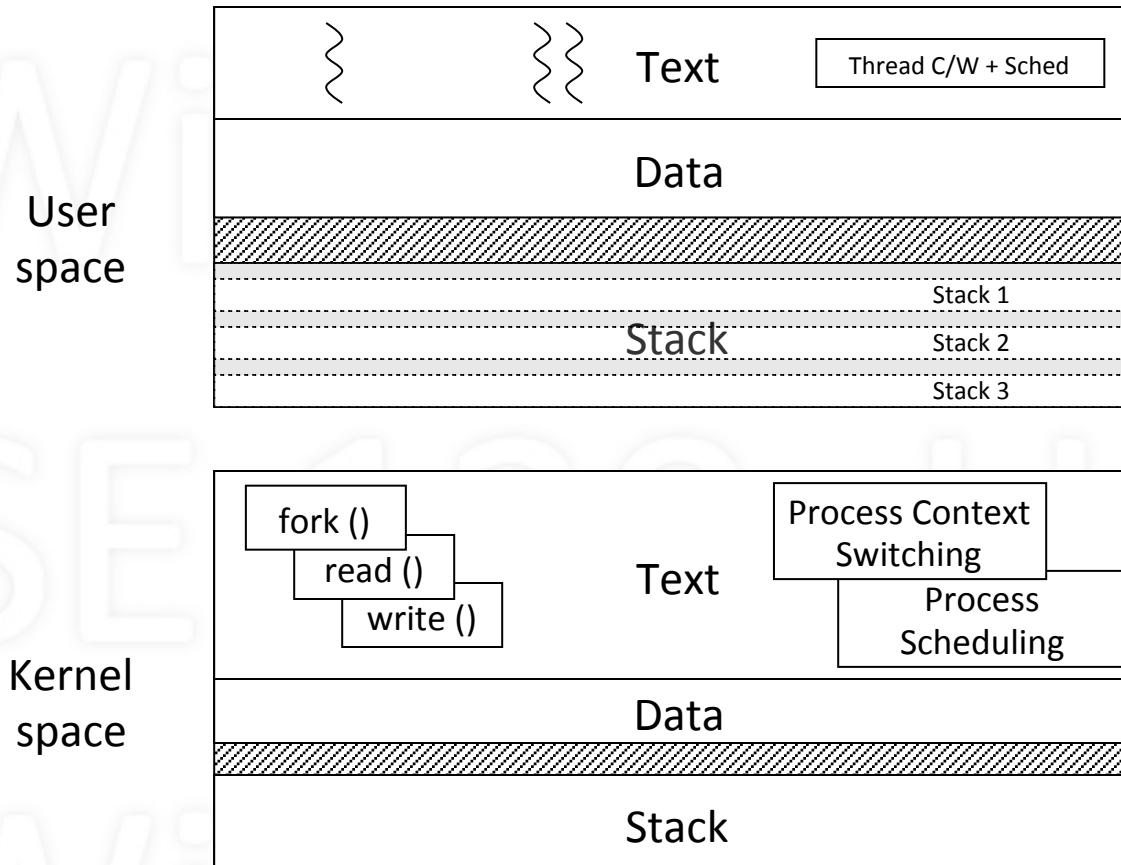
# Many Processes with Threads



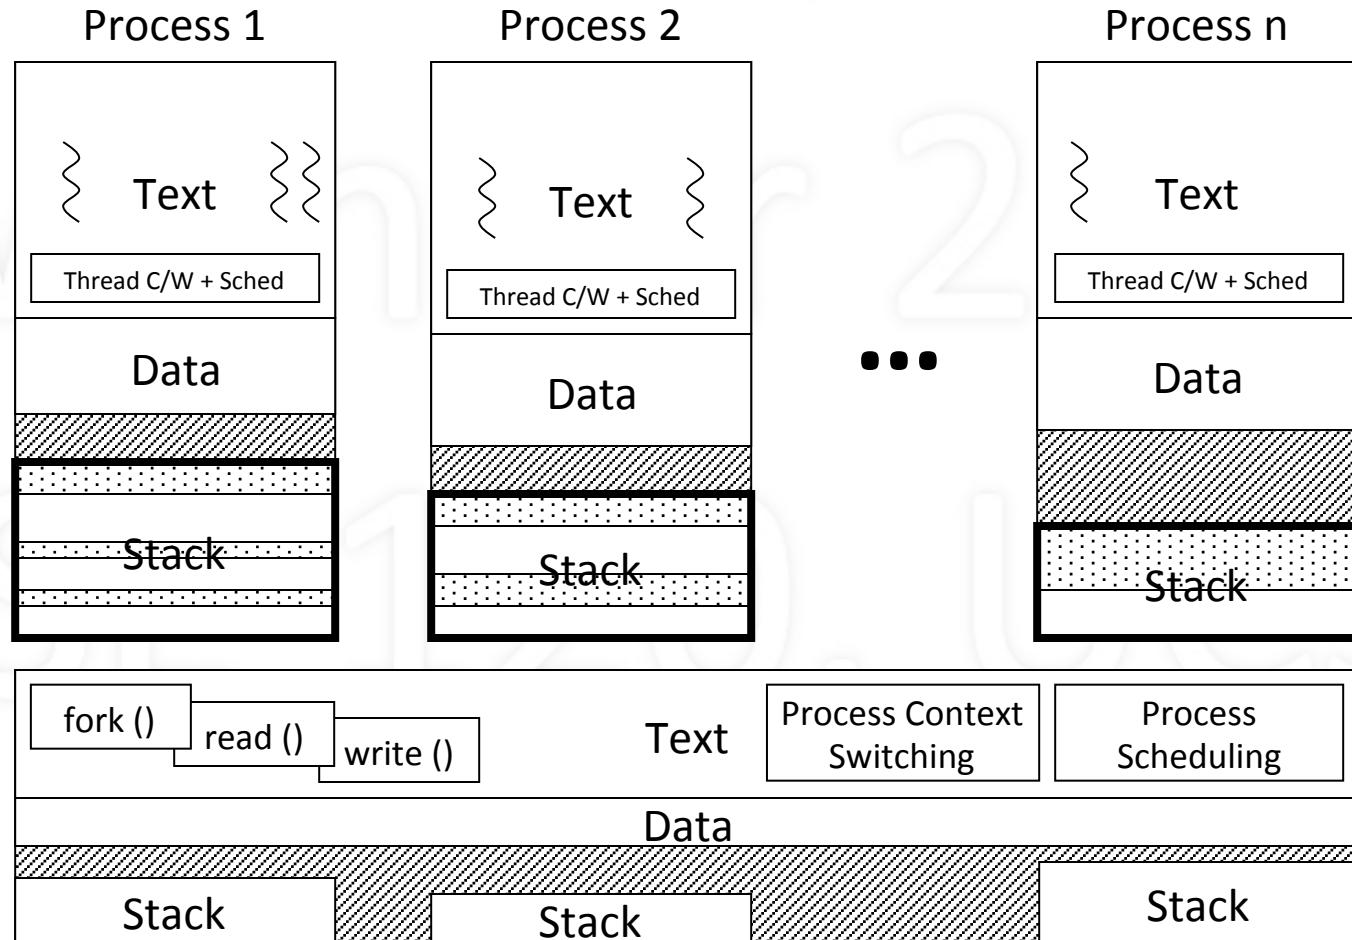
# User-Level Threads

- Can support threads at user level
- Included via thread library
- Thread calls at user level
  - ForkThread (), YieldThread (), ...
- Thread Management at user level
- Supports threads regardless of kernel support
- However, no true parallelism

# User-Level Threads



# User-Level Threads



# Pros and Cons

- User-level threads
  - Portability: works on any kernel
  - Efficient: thread-switching occurs in user space
  - User can decide on scheduling policy
  - But no true parallelism (without special support)
- Kernel-level threads
  - Can achieve true parallelism
  - Overhead: thread switch requires kernel call

# Thread Support vs. Execution

- Distinguish between
  - Where is thread abstraction supported?
  - Where is thread executing?
- User-level vs. kernel-level threads
  - Is thread *support* part of user or kernel code?
- Running in user space vs. kernel space
  - Is thread *running in* user or kernel space?
- Make sure you understand the distinction!

# Summary

- Timesharing: rapidly switching the CPU
  - creates illusion of parallel progress of processes
- Process states: running, ready, blocked
  - distinguishes logical vs. physical execution
- Kernel: extension of all processes
  - gets control by process yielding or preemption
- Thread: single sequential path of execution
  - kernel threads vs. user threads

# Textbook

- OSP: Chapter 3
- OSC: Chapters 3, 4, 12
  - Lecture-related: 3.1-3.3, 3.9, 4.1, 4.3, 4.8, 12.3.3
  - Recommended: 4.2, 4.4-4.7