

CSE 120: Principles of Operating Systems

Lecture 4: Scheduling

Prof. Joseph Pasquale
University of California, San Diego
January 23, 2023

The CPU Scheduling Problem

- Scheduling: which process gets CPU and when
- Given multiple processes, but only one CPU
- How much CPU time does each process get?
- Possibilities
 - Keep CPU till done
 - Each process uses CPU a bit and passes it on
 - Each process gets proportional to what they pay
- Which is the best policy?

There is No Single Best Policy

- Depends on the goals of the system
- Different for
 - your personal computer
 - large time-shared computer
 - computer controlling a nuclear power plant
- Might even have multiple (conflicting) goals

Scheduling: An Example

Process	Arrival Time	Service Time
A	0	5
B	0	3
C	0	1

- Arrival time: time that process is created
- Service time: CPU time needed to complete
- Turnaround time: from arrival to departure
 - Process arrives, waits for CPU, then uses (in bursts)
 - Departs after CPU usage equals service time
- What order minimizes average turnaround time?

Longest First vs. Shortest First

Longest First										
	0	1	2	3	4	5	6	7	8	9 TT
A										5
B										8
C										9
Average Turnaround Time										22/3

Shortest First										
	0	1	2	3	4	5	6	7	8	9 TT
A										9
B										4
C										1
Average Turnaround Time										14/3

Shortest First Is Provably Optimal

- Given n processes with service times S_1, \dots, S_n
 - Note: processes are numbered $1, 2, 3, \dots, n$
- Average turnaround time computed as follows

Shortest First Is Provably Optimal

- Given n processes with service times S_1, \dots, S_n
 - Note: processes are numbered $1, 2, 3, \dots, n$
- Average turnaround time computed as follows
 - $[S_1 + (S_1 + S_2) + (S_1 + S_2 + S_3) + \dots + (S_1 + \dots + S_n)] / n$

Shortest First Is Provably Optimal

- Given n processes with service times S_1, \dots, S_n
 - Note: processes are numbered $1, 2, 3, \dots, n$
- Average turnaround time computed as follows
 - $[S_1 + (S_1 + S_2) + (S_1 + S_2 + S_3) + \dots + (S_1 + \dots + S_n)] / n$

Shortest First Is Provably Optimal

- Given n processes with service times S_1, \dots, S_n
 - Note: processes are numbered $1, 2, 3, \dots, n$
- Average turnaround time computed as follows
 - $[S_1 + (S_1 + S_2) + (S_1 + S_2 + S_3) + \dots + (S_1 + \dots + S_n)] / n$
 - $[(n \times S_1) +$

Shortest First Is Provably Optimal

- Given n processes with service times S_1, \dots, S_n
 - Note: processes are numbered $1, 2, 3, \dots, n$
- Average turnaround time computed as follows
 - $[S_1 + (S_1 + S_2) + (S_1 + S_2 + S_3) + \dots + (S_1 + \dots + S_n)] / n$
 - $[(n \times S_1) +$

Shortest First Is Provably Optimal

- Given n processes with service times S_1, \dots, S_n
 - Note: processes are numbered $1, 2, 3, \dots, n$
- Average turnaround time computed as follows
 - $[S_1 + (S_1 + S_2) + (S_1 + S_2 + S_3) + \dots + (S_1 + \dots + S_n)] / n$
 - $[(n \times S_1) +$

Shortest First Is Provably Optimal

- Given n processes with service times S_1, \dots, S_n
 - Note: processes are numbered $1, 2, 3, \dots, n$
- Average turnaround time computed as follows
 - $[S_1 + (S_1 + S_2) + (S_1 + S_2 + S_3) + \dots + (S_1 + \dots + S_n)] / n$
 - $[(n \times S_1) + ((n-1) \times S_2) +$

Shortest First Is Provably Optimal

- Given n processes with service times S_1, \dots, S_n
 - Note: processes are numbered $1, 2, 3, \dots, n$
- Average turnaround time computed as follows
 - $[S_1 + (S_1 + S_2) + (S_1 + S_2 + S_3) + \dots + (S_1 + \dots + S_n)] / n$
 - $[(n \times S_1) + ((n-1) \times S_2) +$

Shortest First Is Provably Optimal

- Given n processes with service times S_1, \dots, S_n
 - Note: processes are numbered $1, 2, 3, \dots, n$
- Average turnaround time computed as follows
 - $[S_1 + (S_1 + S_2) + (S_1 + S_2 + S_3) + \dots + (S_1 + \dots + S_n)] / n$
 - $[(n \times S_1) + ((n-1) \times S_2) +$

Shortest First Is Provably Optimal

- Given n processes with service times S_1, \dots, S_n
 - Note: processes are numbered $1, 2, 3, \dots, n$
- Average turnaround time computed as follows
 - $[S_1 + (S_1 + S_2) + (S_1 + S_2 + S_3) + \dots + (S_1 + \dots + S_n)] / n$
 - $[(n \times S_1) + ((n-1) \times S_2) + ((n-2) \times S_3) + \dots]$

Shortest First Is Provably Optimal

- Given n processes with service times S_1, \dots, S_n
 - Note: processes are numbered $1, 2, 3, \dots, n$
- Average turnaround time computed as follows
 - $[S_1 + (S_1 + S_2) + (S_1 + S_2 + S_3) + \dots + (S_1 + \dots + S_n)] / n$
 - $[(n \times S_1) + ((n-1) \times S_2) + ((n-2) \times S_3) + \dots + S_n] / n$

Shortest First Is Provably Optimal

- Given n processes with service times S_1, \dots, S_n
 - Note: processes are numbered $1, 2, 3, \dots, n$
- Average turnaround time computed as follows
 - $[S_1 + (S_1 + S_2) + (S_1 + S_2 + S_3) + \dots + (S_1 + \dots + S_n)] / n$
 - $[(n \times S_1) + ((n-1) \times S_2) + ((n-2) \times S_3) + \dots + S_n] / n$
- In general: order by shortest to longest
 - S_1 has maximum weight (n), minimize it
 - S_2 has next-highest ($n-1$), minimize it after S_1 ; ...

Consider Different Arrival Times

Process	Arrival Time	Service Time
A	0	5
B	1	3
C	2	1



FCFS: First Come First Served

	0	1	2	3	4	5	6	7	8	9	TT
A											5
B											7
C											7
Average Turnaround Time											19/3

- Allocate CPU to processes in order of arrival
- A B]A B]A CB]A CB]A CB]A C]B C]B C]B C
 - A At 0, A arrives and runs at 0
 - B]A At 1-, B arrives, enters queue (A still running)
 - B]A At 1, B is in the queue, A is running

FCFS: First Come First Served

	0	1	2	3	4	5	6	7	8	9	TT
A											5
B											7
C											7
Average Turnaround Time											19/3

- Allocate CPU to processes in order of arrival
- A B]A B]A CB]A CB]A CB]A C]B C]B C]B C
- Average turnaround time = $(5 + 7 + 7)/3 = 6.3$
- Non-preemptive, simple, no starvation
- Poor for short processes

RR: Round Robin

	0	1	2	3	4	5	6	7	8	9	TT
A											9
B											6
C											2
Average Turnaround Time											17/3

- Time-slice: each process gets quantum in turn
- A B]A A]B CA]B BC]A AB]C A]B B]A A]B A A
- Average turnaround time = $(9 + 6 + 2)/3 = 5.7$
- Preemptive, simple, no starvation
- Process waits at most $(n - 1) \times \text{quantum}$

SPN: Shortest Process Next

	0	1	2	3	4	5	6	7	8	9	TT
A											5
B											8
C											4
Average Turnaround Time											17/3

- Select process with shortest service time
- A B]A B]A BC]A BC]A BC]A B]C B B B
- Average turnaround time = $(5 + 8 + 4)/3 = 5.7$
- Optimal for non-preemptive, allows starvation
- Assumes service times are known

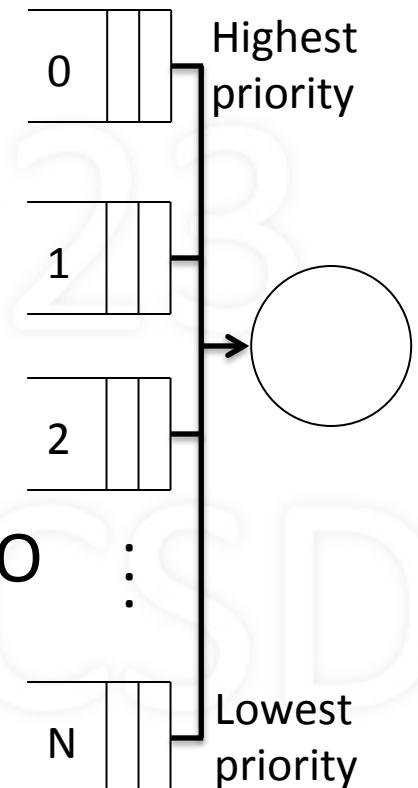
SRT: Shortest Remaining Time

	0	1	2	3	4	5	6	7	8	9	TT
A											9
B											4
C											1
Average Turnaround Time											14/3

- Select process with shortest remaining time
- A B]A A]B AC]B AB]C A]B A]B A A A A
- Average turnaround time = $(9 + 4 + 1)/3 = 4.7$
- Assumes service times are known
- Optimal for preemptive, but allows starvation

Multi-Level Feedback Queues

- Priority queues: 0 (high), ..., N (low)
- New processes enter queue 0
- Select from highest priority queue
- Run for $T = 2^k$ quantums
 - Used T : move to next lower queue, FIFO
 - Used $< T$: back to same queue, RR
 - Due to yield or higher priority arrival
- Periodically boost (e.g., all to highest queue)

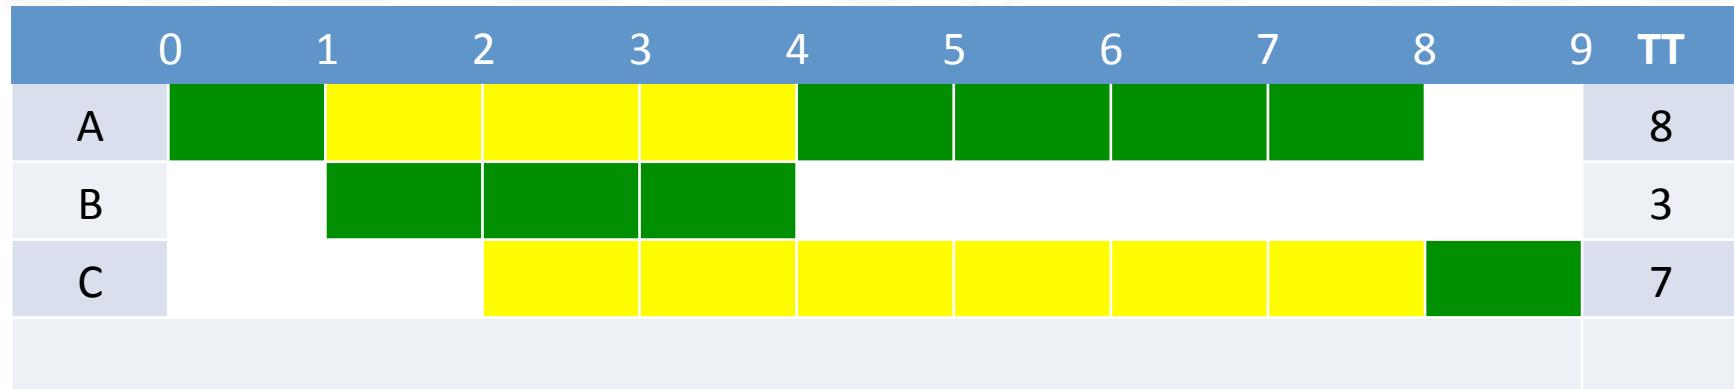


Multi-Level Feedback Queues

	0	1	2	3	4	5	6	7	8	9	TT
A											9
B											6
C											1
Average Turnaround Time											16/3

- Select from highest queue k , run 2^k quantums
- A B]₀A A]₁B A]₁C]₀B BA]₁C B]₁A(2) A]₂B(2) A
- Average Turnaround Time = $(9 + 6 + 1)/3 = 5.3$
- Complex, adaptive, highly responsive
- Favors shorter over longer, possible starvation

Priority Scheduling



- Select process with highest priority
 - Example: $P_A = \text{medium}$, $P_B = \text{high}$, $P_C = \text{low}$
- $A_M \rightarrow B_H \rightarrow A_M \rightarrow A_M \rightarrow B_H \rightarrow C_L \rightarrow A_M \rightarrow B_H \rightarrow C_L \rightarrow A_M \rightarrow (2) \rightarrow C_L \rightarrow A_M \rightarrow (4) \rightarrow C$
- Allows scheduling based on external criteria
 - E.g., priority = $1/\text{CPU_time_used}$, or = $f(\text{pay})$

Fair Share (Proportional Share)

	1	2	3	4	5	6	7	8	9	10
A	100%	50%	33%	50%	40%	50%	43%	50%	44%	50%
B	0%	50%	33%	25%	20%	17%	14%	13%	11%	10%
C	0%	0%	33%	25%	40%	33%	43%	38%	44%	40%

- Each process requests some CPU utilization
 - Utilization: what percentage of time resource is used
 - Example of requests: A: 50%; B: 10%; C: 40%
- Goal: utilization over long run, actual \approx request
- How do we determine who runs each quantum?
- Select process with minimum actual/request ratio

FS Example: A 50%, B 10%, C 40%

	1	2	3	4	5	6	7	8	9	10
A										
B										
C										

- At time 0 A: 0/50 B: 0/10 C: 0/40
- All 0, so arbitrary select A

FS Example: A 50%, B 10%, C 40%

	1	2	3	4	5	6	7	8	9	10
A	100%									
B	0%									
C	0%									

- After 1q A: 100/50 B: 0/10 C: 0/40
- A got 100%, but only requested 50% (100/50)
 - Got *more* than its fair share ($100/50 = 2 > 1$)!
- Both B, C got 0%, less than their fair shares
- Select process with min actual/request ratio: B

FS Example: A 50%, B 10%, C 40%

	1	2	3	4	5	6	7	8	9	10
A	100%	50%								
B	0%	50%								
C	0%	0%								

- After 2q A: 50/50 B: 50/10 C: 0/40
- A got 50%: used 1 out of 2q; same for B
 - A got *exactly* its fair share ($50/50 = 1$)
 - B got *more* than its fair share ($50/10 = 5 > 1$)
- Select process with min actual/request ratio: C

FS Example: A 50%, B 10%, C 40%

	1	2	3	4	5	6	7	8	9	10
A	100%	50%	33%							
B	0%	50%	33%							
C	0%	0%	33%							

- After 3q A: $33/50$ B: $33/10$ C: $33/40$
- A got 33% as it used 1 out of 3; same for B, C
 - A, C got *less* than their fair share ($33/50, 33/40 < 1$)
 - B got *more* than its fair share ($33/10 > 1$)
- Select process with min actual/request ratio: A

FS Example: A 50%, B 10%, C 40%

	1	2	3	4	5	6	7	8	9	10
A	100%	50%	33%	50%						
B	0%	50%	33%	25%						
C	0%	0%	33%	25%						

- After 4q A: 50/50 B: 25/10 C: 25/40
- A got 50% as it used 2 out of 4
- B and C got 25%, using 1 out of 4
- A got its fair share; B got more; C got less
- Select process with min actual/request ratio: C

FS Example: A 50%, B 10%, C 40%

	1	2	3	4	5	6	7	8	9	10
A	100%	50%	33%	50%	40%	50%	43%	50%	44%	50%
B	0%	50%	33%	25%	20%	17%	14%	13%	11%	10%
C	0%	0%	33%	25%	40%	33%	43%	38%	44%	40%

- After 10q A: 50/50 B: 10/10 C: 40/40
- Recall goal: over long run, actual \approx request
 - Process utilizations are what they requested!
- However, algorithm shown is inefficient
 - Each q, each p: calculate actual/request, select min

Stride Scheduling

- For processes A, B, C ... with requests $R_A, R_B, R_C \dots$
- Calculate ***strides***: $S_A = 1/R_A, S_B = 1/R_B, S_C = 1/R_C \dots$
- For each process x , maintain ***pass*** value P_x (init 0)
- Schedule: repeat every quantum
 - Select process x with minimum pass value P_x , run
 - Increment pass value by stride value: $P_x = P_x + S_x$
- Optimization: use only integers for R_x, S_x and P_x
 - Calculate $S_x = L/R_x$ using very large L, e.g., $L = 100000$

Stride Scheduling Example

Process x	Runs Now	A	B	C	To Run Next
Requested Utilization: R_x (in %)		50	10	40	
Stride: $S_x = L/R_x$ ($L = 100000$)		2000	10000	2500	
Pass: $P_x = P_x + S_x$; init = 0		0	0	0	A
quantum 1					
quantum 2					
quantum 3					
quantum 4					
quantum 5					
quantum 6					
quantum 7					
quantum 8					
quantum 9					
quantum 10					

- All pass values 0, select one arbitrarily: A

Stride Scheduling Example

Process x	Runs Now	A	B	C	To Run Next
Requested Utilization: R_x (in %)		50	10	40	
Stride: $S_x = L/R_x$ ($L = 100000$)		2000	10000	2500	
Pass: $P_x = P_x + S_x$; init = 0		0	0	0	A
quantum 1	A	2000	0	0	B
quantum 2					
quantum 3					
quantum 4					
quantum 5					
quantum 6					
quantum 7					
quantum 8					
quantum 9					
quantum 10					

- $P_A = 2000$, P_B and $P_C = 0$; select B (or C)

Stride Scheduling Example

Process x	Runs Now	A	B	C	To Run Next
Requested Utilization: R_x (in %)		50	10	40	
Stride: $S_x = L/R_x$ ($L = 100000$)		2000	10000	2500	
Pass: $P_x = P_x + S_x$; init = 0		0	0	0	A
quantum 1	A	2000	0	0	B
quantum 2	B	2000	10000	0	C
quantum 3					
quantum 4					
quantum 5					
quantum 6					
quantum 7					
quantum 8					
quantum 9					
quantum 10					

- $P_A = 2000, P_B = 10000, P_C = 0$; select C

Stride Scheduling Example

Process x	Runs Now	A	B	C	To Run Next
Requested Utilization: R_x (in %)		50	10	40	
Stride: $S_x = L/R_x$ ($L = 100000$)		2000	10000	2500	
Pass: $P_x = P_x + S_x$; init = 0		0	0	0	A
quantum 1	A	2000	0	0	B
quantum 2	B	2000	10000	0	C
quantum 3	C	2000	10000	2500	A
quantum 4					
quantum 5					
quantum 6					
quantum 7					
quantum 8					
quantum 9					
quantum 10					

- $P_A = 2000, P_B = 10000, P_C = 2500$; select A

Stride Scheduling Example

Process x	Runs Now	A	B	C	To Run Next
Requested Utilization: R_x (in %)		50	10	40	
Stride: $S_x = L/R_x$ ($L = 100000$)		2000	10000	2500	
Pass: $P_x = P_x + S_x$; init = 0		0	0	0	A
quantum 1	A	2000	0	0	B
quantum 2	B	2000	10000	0	C
quantum 3	C	2000	10000	2500	A
quantum 4	A	4000	10000	2500	C
quantum 5					
quantum 6					
quantum 7					
quantum 8					
quantum 9					
quantum 10					

- $P_A = 4000, P_B = 10000, P_C = 2500$; select C

Stride Scheduling Example

Process x	Runs Now	A	B	C	To Run Next
Requested Utilization: R_x (in %)		50	10	40	
Stride: $S_x = L/R_x$ ($L = 100000$)		2000	10000	2500	
Pass: $P_x = P_x + S_x$; init = 0		0	0	0	A
quantum 1	A	2000	0	0	B
quantum 2	B	2000	10000	0	C
quantum 3	C	2000	10000	2500	A
quantum 4	A	4000	10000	2500	C
quantum 5	C	4000	10000	5000	A
quantum 6					
quantum 7					
quantum 8					
quantum 9					
quantum 10					

- $P_A = 4000, P_B = 10000, P_C = 5000$; select A

Stride Scheduling Example

Process x	Runs Now	A	B	C	To Run Next
Requested Utilization: R_x (in %)		50	10	40	
Stride: $S_x = L/R_x$ ($L = 100000$)		2000	10000	2500	
Pass: $P_x = P_x + S_x$; init = 0		0	0	0	A
quantum 1	A	2000	0	0	B
quantum 2	B	2000	10000	0	C
quantum 3	C	2000	10000	2500	A
quantum 4	A	4000	10000	2500	C
quantum 5	C	4000	10000	5000	A
quantum 6	A	6000	10000	5000	C
quantum 7					
quantum 8					
quantum 9					
quantum 10					

- $P_A = 6000, P_B = 10000, P_C = 5000$; select C

Stride Scheduling Example

Process x	Runs Now	A	B	C	To Run Next
Requested Utilization: R_x (in %)		50	10	40	
Stride: $S_x = L/R_x$ ($L = 100000$)		2000	10000	2500	
Pass: $P_x = P_x + S_x$; init = 0		0	0	0	A
quantum 1	A	2000	0	0	B
quantum 2	B	2000	10000	0	C
quantum 3	C	2000	10000	2500	A
quantum 4	A	4000	10000	2500	C
quantum 5	C	4000	10000	5000	A
quantum 6	A	6000	10000	5000	C
quantum 7	C	6000	10000	7500	A
quantum 8					
quantum 9					
quantum 10					

- $P_A = 6000, P_B = 10000, P_C = 7500$; select A

Stride Scheduling Example

Process x	Runs Now	A	B	C	To Run Next
Requested Utilization: R_x (in %)		50	10	40	
Stride: $S_x = L/R_x$ ($L = 100000$)		2000	10000	2500	
Pass: $P_x = P_x + S_x$; init = 0		0	0	0	A
quantum 1	A	2000	0	0	B
quantum 2	B	2000	10000	0	C
quantum 3	C	2000	10000	2500	A
quantum 4	A	4000	10000	2500	C
quantum 5	C	4000	10000	5000	A
quantum 6	A	6000	10000	5000	C
quantum 7	C	6000	10000	7500	A
quantum 8	A	8000	10000	7500	C
quantum 9					
quantum 10					

- $P_A = 8000, P_B = 10000, P_C = 7500$; select C

Stride Scheduling Example

Process x	Runs Now	A	B	C	To Run Next
Requested Utilization: R_x (in %)		50	10	40	
Stride: $S_x = L/R_x$ ($L = 100000$)		2000	10000	2500	
Pass: $P_x = P_x + S_x$; init = 0		0	0	0	A
quantum 1	A	2000	0	0	B
quantum 2	B	2000	10000	0	C
quantum 3	C	2000	10000	2500	A
quantum 4	A	4000	10000	2500	C
quantum 5	C	4000	10000	5000	A
quantum 6	A	6000	10000	5000	C
quantum 7	C	6000	10000	7500	A
quantum 8	A	8000	10000	7500	C
quantum 9	C	8000	10000	10000	A
quantum 10					

- $P_A = 8000, P_B = 10000, P_C = 10000$; select A

Stride Scheduling Example

Process x	Runs Now	A	B	C	To Run Next
Requested Utilization: R_x (in %)		50	10	40	
Stride: $S_x = L/R_x$ ($L = 100000$)		2000	10000	2500	
Pass: $P_x = P_x + S_x$; init = 0		0	0	0	A
quantum 1	A	2000	0	0	B
quantum 2	B	2000	10000	0	C
quantum 3	C	2000	10000	2500	A
quantum 4	A	4000	10000	2500	C
quantum 5	C	4000	10000	5000	A
quantum 6	A	6000	10000	5000	C
quantum 7	C	6000	10000	7500	A
quantum 8	A	8000	10000	7500	C
quantum 9	C	8000	10000	10000	A
quantum 10	A	10000	10000	10000	... will repeat

- $P_A = 10000, P_B = 10000, P_C = 10000$

Stride Scheduling Example

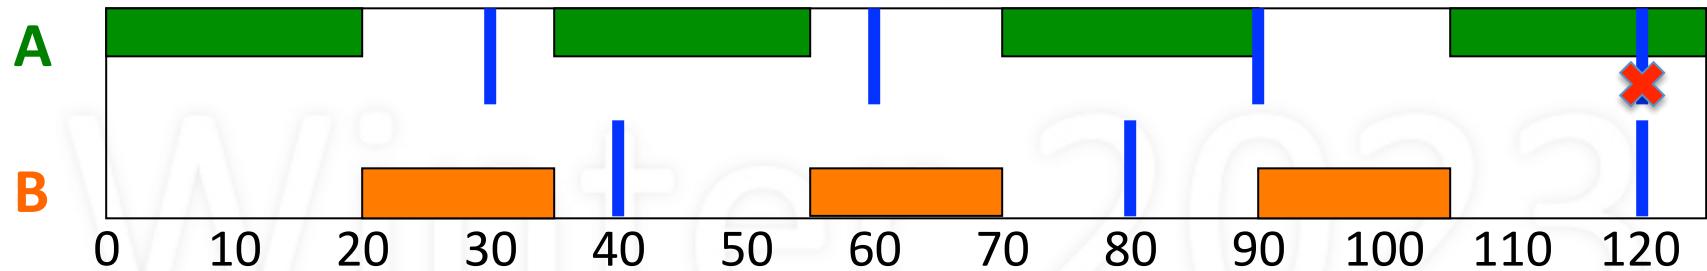
Process x	Runs Now	A	B	C	To Run Next
Requested Utilization: R_x (in %)		50	10	40	
Stride: $S_x = L/R_x$ ($L = 100000$)		2000	10000	2500	
Pass: $P_x = P_x + S_x$; init = 0		0	0	0	A
quantum 1	A	2000	0	0	B
quantum 2	B	2000	10000	0	C
quantum 3	C	2000	10000	2500	A
quantum 4	A	4000	10000	2500	C
quantum 5	C	4000	10000	5000	A
quantum 6	A	6000	10000	5000	C
quantum 7	C	6000	10000	7500	A
quantum 8	A	8000	10000	7500	C
quantum 9	C	8000	10000	10000	A
quantum 10	A	10000	10000	10000	... will repeat

- After 10 quantums, A ran 5; B ran 1; C ran 4

Real Time Scheduling

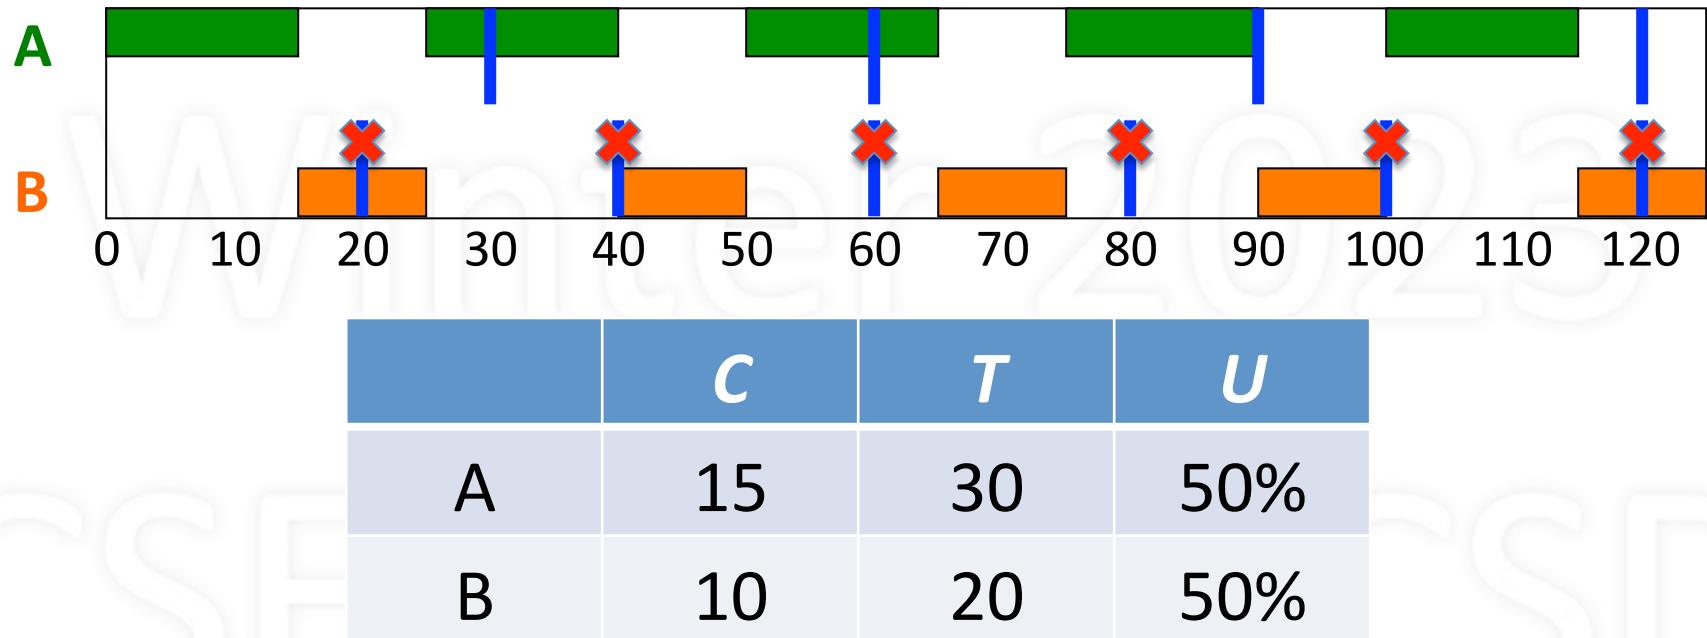
- Correctness of real-time systems depend on
 - logical result of computations
 - *and* the timing of these results
- Type of real-time systems
 - Hard vs. soft real-time
 - Periodic vs. aperiodic
- Scheduling
 - Earliest Deadline First (EDF)
 - Rate Monotonic Scheduling (RMS)

Periodic Processes (or Tasks)



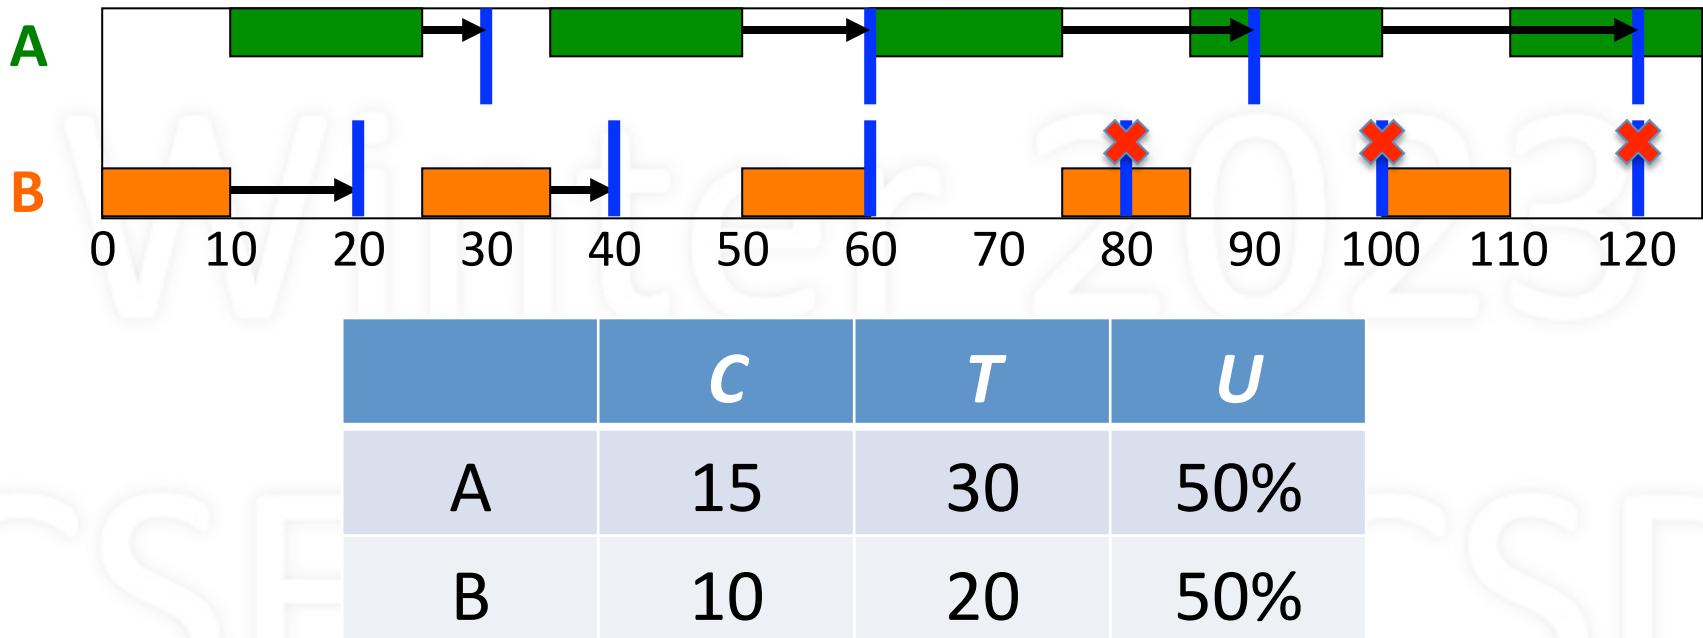
- Periodic processes: computation is cyclic
- For each process, given
 C = CPU burst, T = period, $U = C/T$ = utilization
- Can processes be ordered so deadlines are met?
- Consider orders: ABABAB..., vs. BABABA...

Periodic Processes



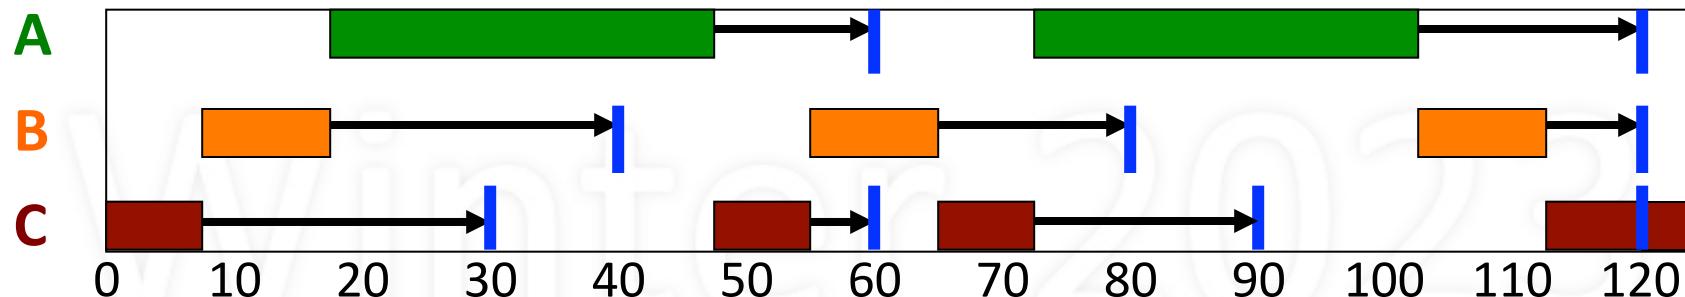
- Sum of utilizations does not exceed 100%
- For order ABABAB..., B misses all deadlines!

Periodic Processes



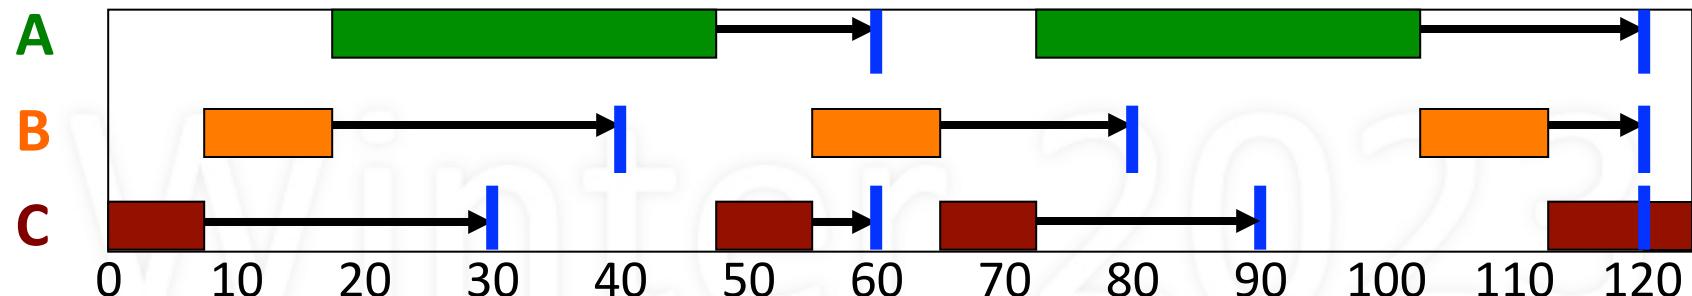
- Sum of utilizations does not exceed 100%
- For order BABABA..., B misses some deadlines

EDF: Earliest Deadline First



- Schedule process with earliest deadline
- If earlier deadline process appears, preempt
- Works for periodic and aperiodic processes
- Achieves 100% utilization (ignoring overhead!)
- Expensive: requires ordering by deadlines

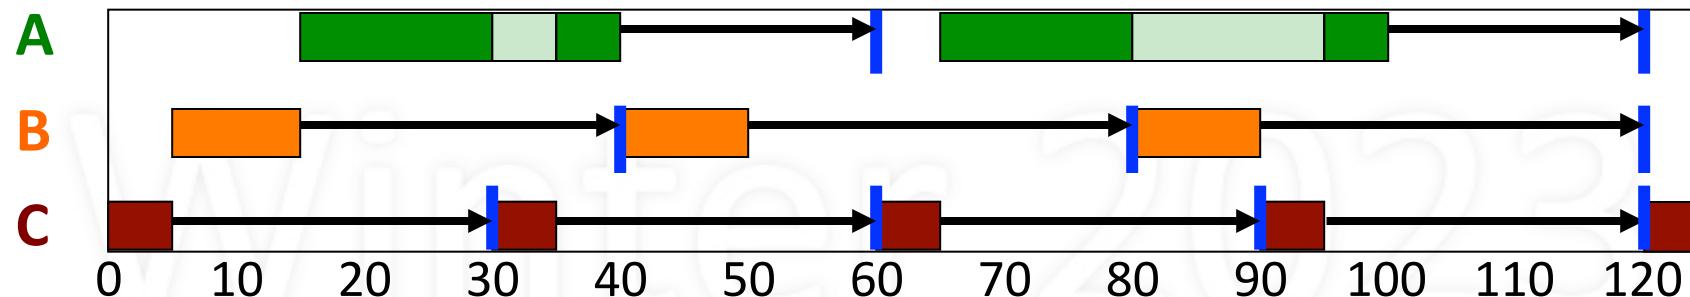
EDF: Earliest Deadline First



	<i>C</i>	<i>T</i>	<i>U</i>
A	30	60	50%
B	10	40	25%
C	7.5	30	25%

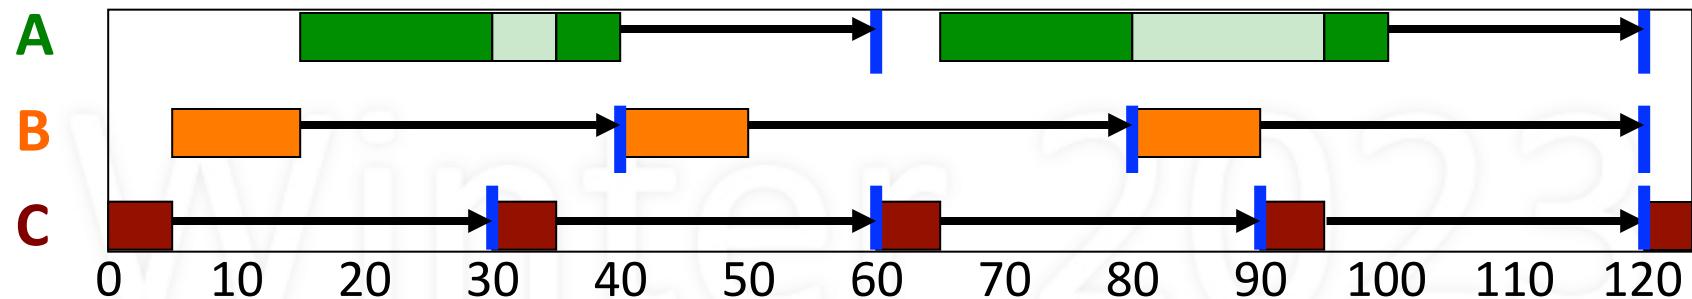
- Meets all deadlines, sum of utilizations = 100%

RMS: Rate Monotonic Scheduling



- If periodic processes, prioritize based on rates
- At start of period, select highest priority
- Preempt if necessary
- When burst done, wait till next period
- If $U_1 + \dots + U_n \leq n (2^{1/n} - 1)$, all deadlines met

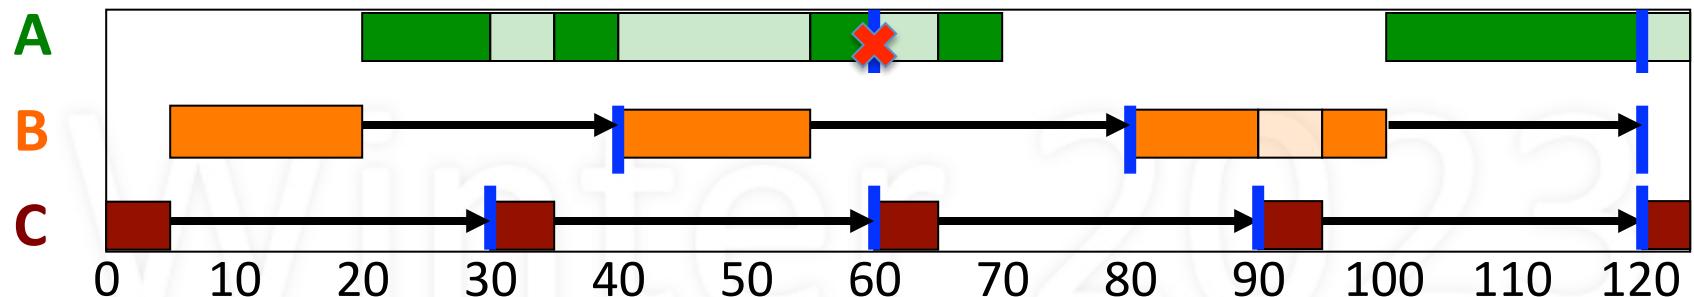
RMS Test Passes, All Deadlines Met



	<i>C</i>	<i>T</i>	<i>U</i>	Rate	Prio
A	20	60	33%	$1/60 = 0.017$	Low
B	10	40	25%	$1/40 = 0.025$	Med
C	5	30	17%	$1/30 = 0.033$	High

- Passes: $U_A + U_B + U_C = 75\% \leq 3(2^{1/3} - 1) = 78\%$

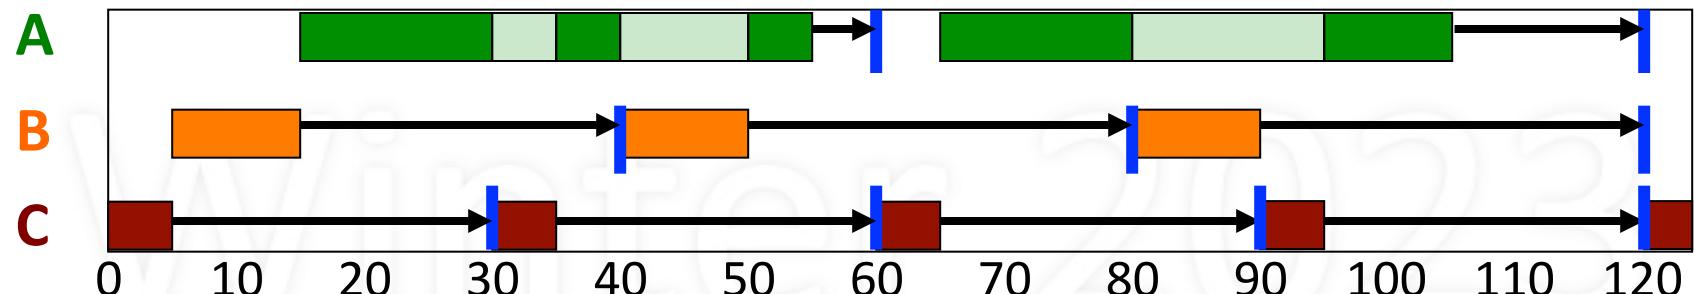
RMS Test Fails, Deadline Missed



	<i>C</i>	<i>T</i>	<i>U</i>	Rate	Prio
A	25	60	42%	$1/60 = 0.017$	Low
B	15	40	38%	$1/40 = 0.025$	Med
C	5	30	17%	$1/30 = 0.033$	High

- Fails: $U_A + U_B + U_C = 97\% > 3(2^{1/3} - 1) = 78\%$

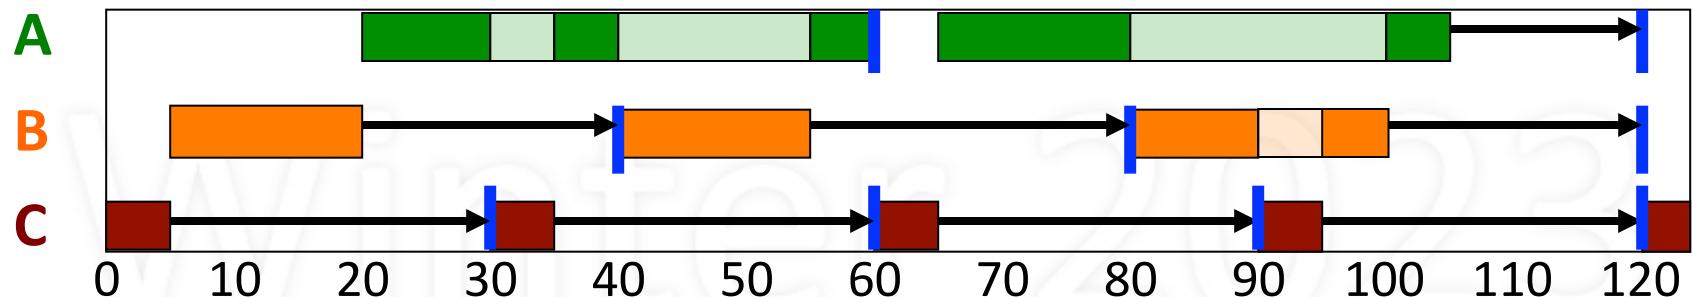
RMS Test Fails, But Deadlines Met



	<i>C</i>	<i>T</i>	<i>U</i>	Rate	Prio
A	25	60	42%	$1/60 = 0.017$	Low
B	10	40	25%	$1/40 = 0.025$	Med
C	5	30	17%	$1/30 = 0.033$	High

- Fails: $U_A + U_B + U_C = 84\% > 3(2^{1/3} - 1) = 78\%$

RMS Test Fails, But Deadlines Met



	<i>C</i>	<i>T</i>	<i>U</i>	Rate	Prio
A	20	60	33%	$1/60 = 0.017$	Low
B	15	40	38%	$1/40 = 0.025$	Med
C	5	30	17%	$1/30 = 0.033$	High

- Fails: $U_A + U_B + U_C = 88\% > 3(2^{1/3} - 1) = 78\%$

RMS Optimal But Limited

- RMS is simple and efficient
 - Static priority scheduling based on rates
- RMS is optimal for static priority algorithms
 - If RMS can't schedule, no other static priority can
- RMS is limited in what it guarantees
 - Utilization bounded by $n (2^{1/n} - 1) > \ln 2 \approx 69\%$
 - Deadline guarantee applies only if test passes
- RMS is limited to periodic processes

Summary

- CPU scheduling is policy: depends on goals
 - First come first served
 - Round robin
 - Shortest process next
 - Shortest remaining time
 - Multi-level feedback
 - Priority
 - Fair share
 - Earliest deadline first
 - Rate monotonic sched

Textbook

- OSP: Chapter 4
- OSC: Chapter 6
 - Lecture-related: 5.1-5.3, 5.6, 5.9
 - Recommended: 5.7-5.8