

CSE 120: Principles of Operating Systems

Lecture 5: Synchronization

Prof. Joseph Pasquale
University of California, San Diego
January 30, 2023

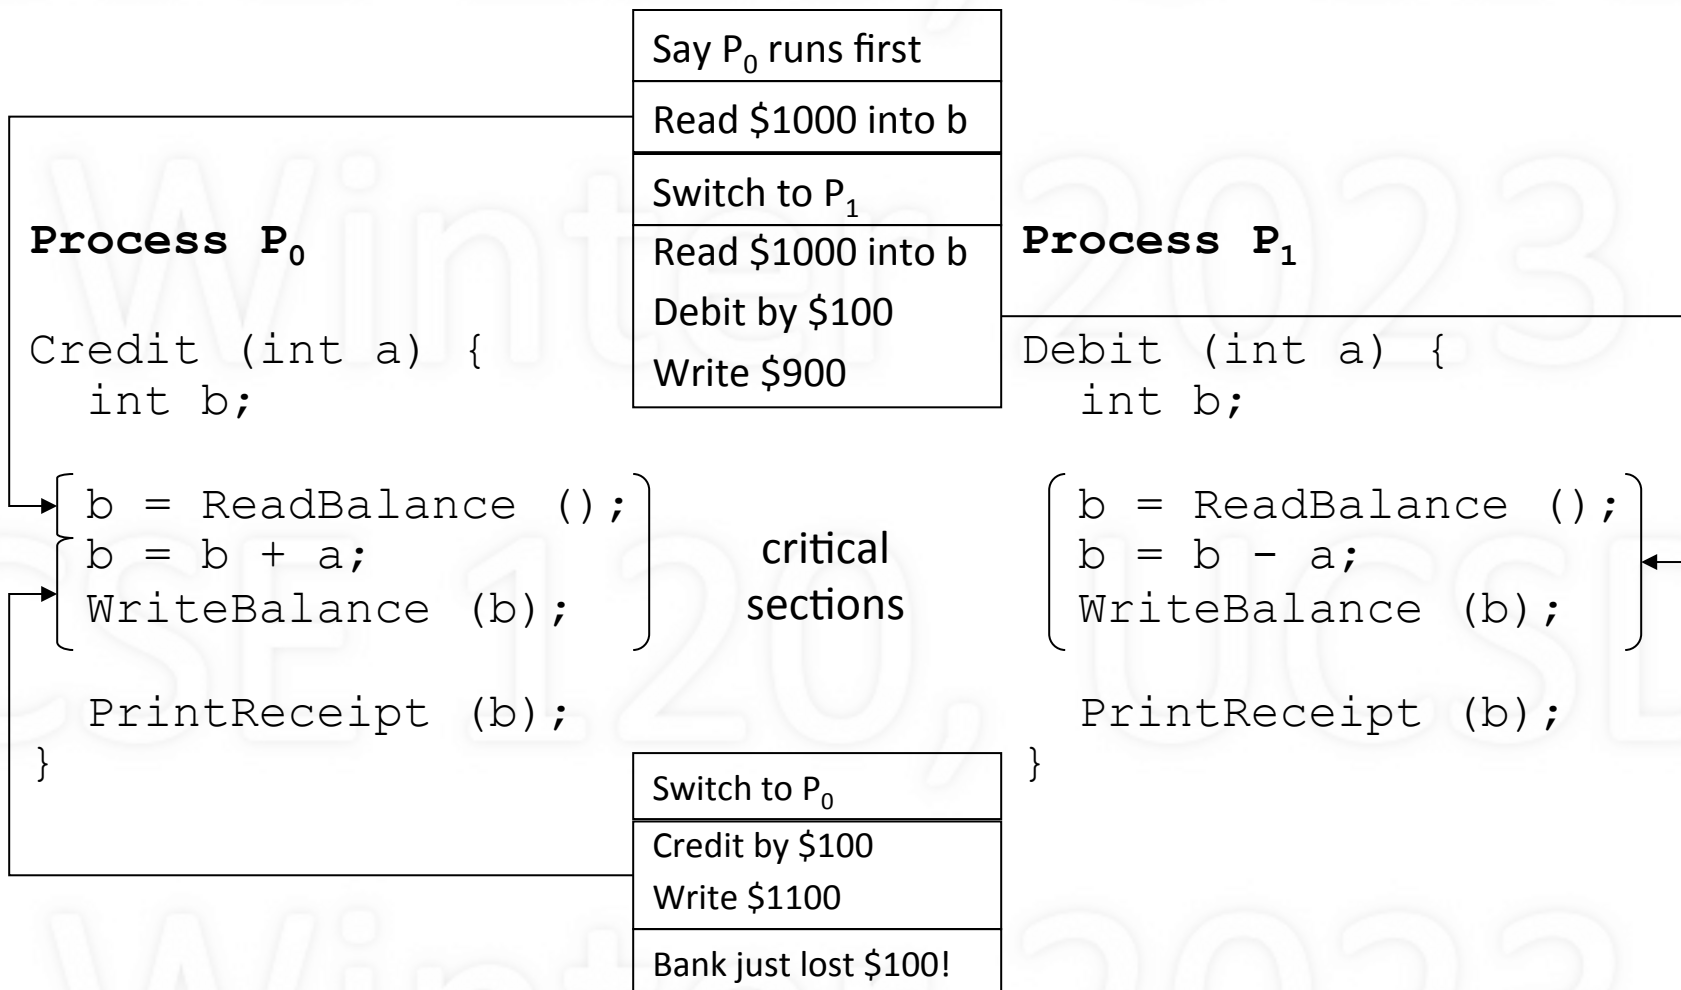
Synchronization

- Synchronize: events happen at the same time
- Process synchronization
 - Events in processes that occur “at the same time”
 - Actually, when one process waits for another
- Uses of synchronization
 - Prevent race conditions
 - Wait for resources to become available


The Credit/Debit Problem

- Say you have \$1000 in your bank account
- You deposit \$100
- You also withdraw \$100
- How much should be in your account?
- What if deposit/withdraw occur at same time?

Credit/Debit Problem: Race Condition



To Avoid Race Conditions



- Identify related *critical sections*
 - Section(s) of code executed by different processes
 - Must run *atomically, with respect to each other*
- Enforce mutual exclusion
 - Only one process active in a critical section

What Does Atomic Really Mean?

- Atomic means “indivisible”
- We seek *effective* atomicity
 - can interrupt, as long as interruption has no effect
- It *is* OK to interrupt process in critical section
 - as long as other processes have no effect
- How to determine
 - Consider effect of critical section in isolation
 - Next consider interruptions: if same result, OK

How to Achieve Mutual Exclusion?

< entry code >

< critical section >

< exit code >

< entry code >

< critical section >

< exit code >

- Surround critical section with entry/exit code
- Entry code should act as a barrier
 - If another process is in critical section, block
 - Otherwise, allow process to proceed
- Exit code should release other entry barriers

Requirements for Good Solution

- Given multiple cooperating processes
 - Each process has a critical section
 - All critical sections are to be mutually exclusive
1. At most one in a critical section at a time
 2. Can't prevent entry if no others are in theirs
 3. Should eventually be able to enter
 4. No assumptions about CPU speed or number

Software Lock?

```
shared int lock = OPEN;
```

P₀

```
while (lock == CLOSED);  
lock = CLOSED;
```

```
< critical section >
```

```
lock = OPEN;
```

P₁

```
while (lock == CLOSED);  
lock = CLOSED;
```

```
< critical section >
```

```
lock = OPEN;
```

- Lock indicates if any process in critical section
- Is there a problem?

Take Turns?

```
shared int turn = 0;           // arbitrary set to  $P_0$ 
```

P_0

```
while (turn != 0);  
< critical section >  
turn = 1;
```

P_1

```
while (turn != 1);  
< critical section >  
turn = 0;
```

- Alternate which process enters critical section
- Is there a problem?

State Intention?

```
shared boolean intent[2] = {FALSE, FALSE};
```

P₀

```
intent[0] = TRUE;  
while (intent[1]);  
< critical section >  
intent[0] = FALSE;
```

P₁

```
intent[1] = TRUE;  
while (intent[0]);  
< critical section >  
intent[1] = FALSE;
```

- Process states intent to enter critical section
- Is there a problem?

Peterson's Solution

```
shared int turn;  
shared boolean intent[2] = {FALSE, FALSE};
```

P₀

```
intent[0] = TRUE;  
turn = 1;  
while (intent[1] && turn==1);
```

< critical section >

```
intent[0] = FALSE;
```

P₁

```
intent[1] = TRUE;  
turn = 0;  
while (intent[0] && turn==0);
```

< critical section >

```
intent[1] = FALSE;
```

- If competition, take turns; otherwise, enter
- Is there a problem?
- There is a version for $n \geq 2$; more complex

What about Disabling Interrupts?

- Reasoning
 - No interrupts \Rightarrow no uncontrolled context switches
 - No uncontrolled context switches \Rightarrow no races
 - No races \Rightarrow mutual exclusion
- Is there a problem?

Test-and-Set Lock Instruction: TSL

- TSL mem (test-and-set lock: contents of mem)
do atomically (i.e., locking the memory bus)
[**test** if mem == 0 AND **set** mem = 1]
- Operations occur without interruption
 - Memory bus is locked
 - Not affected by hardware interrupts

What TSL Does, Expressed in C

- Assume C function, TSL(int *), that is *atomic*

```
TSL(int *lockptr)
{
    int oldval;

    oldval = *lockptr
    *lockptr = 1;
    return ((oldval == 0) ? 1 : 0);
}
```

Effectively, this is what the
TSL instruction does in one
atomic instruction

Mutual Exclusion Using TSL

```
shared int lock = 0;
```

P₀

```
while (! TSL(&lock));  
< critical section >  
lock = 0;
```

P₁

```
while (! TSL(&lock));  
< critical section >  
lock = 0;
```

- Shared variable solution using TSL(int *)
 - tests if lock == 0 (if so, will return 1; else 0)
 - before returning, sets lock to 1
- Simple, works for any number of processes
- Still “suffers” from busy waiting

Semaphores

- Synchronization variable
 - Takes on integer values
 - Can cause a process to block/unblock
- wait and signal operations
 - wait (s) decrement; block if < 0
 - signal (s) increment; if any blocked, unblock one
- No other operations allowed
 - In particular, cannot test value of semaphore!

Examples and Interpretation

- wait (s) decrement; block if < 0
- signal (s) increment; if any blocked, unblock

- wait (1) $s \rightarrow 0$ GO
- wait (0) $s \rightarrow -1$ STOP (i.e., block)
- signal (-1) $s \rightarrow 0$ GO *and* allow one to GO
- signal (0) $s \rightarrow 1$ GO

Mutual Exclusion

```
sem mutex = 1;           // declare and initialize
```

P₀

```
wait (mutex);  
< critical section >  
signal (mutex);
```

P₁

```
wait (mutex);  
< critical section >  
signal (mutex);
```

- Use “mutex” semaphore, initialized to 1
- Only one process can enter critical section
- Simple, works for n processes
- Is there any busy-waiting?

Order How Processes Execute

```
sem cond = 0;
```

P₀

< to be done before P₁ >

```
signal (cond);
```

P₁

```
wait (cond);
```

< to be done after P₀ >

- Cause a process to wait for another
- Use semaphore indicating condition; initially 0
 - the condition in this case: “P₀ has completed”
- Used for ordering processes
 - In contrast to mutual exclusion

Semaphores: Only Synchronization

- Semaphores only provide synchronization
 - Synchronization: when a process blocks for event
- But, no information transfer
 - No way for a process to tell it blocked

Semaphore Implementation

- Semaphore $s = [n, L]$
 - n : takes on integer values
 - L : list of processes blocked on s

- Operations

```
wait (sem s) {  
    s.n = s.n - 1;  
    if (s.n < 0) add calling process to S.L and block; }  
  
signal (sem s) {  
    s.n = s.n + 1;  
    if (s.L !empty) remove/unblock a process from s.L; }
```

Alternative Implementation

- Semaphore $s = [n, L]$
 - n : takes on integer values, non-negative
 - L : list of processes blocked on s

- Operations

```
wait (sem s) {  
    if (s.n == 0) add calling process to s.L and block;  
    else s.n = s.n - 1; }  
  
signal (sem s) {  
    if (s.L !empty) remove/unblock a process from s.L;  
    else s.n = s.n + 1; }
```

Wait and Signal Must Be Atomic

- Bodies of wait and signal are critical sections
- So, still need mechanism for mutual exclusion!
- Use a lower-level (more basic) mechanism
 - Test-and-set lock
 - Peterson's solution
- So, busy-waiting still exists (can never remove)
 - But at lower-level (within semaphore operations)
 - Occurrence limited to brief/known periods of time

Analysis: Lower-Level Busy Waiting

- A calls wait (s), switch to B, B calls wait (s)
 - Switch occurs while A executing body of wait
- Body of wait is critical section, so B must block
 - Use test-set lock or Peterson's: busy waiting
- How long will B be blocked?
 - For time it takes to execute body of wait
- Small/known amount of time!
 - Compare to user critical section: unknown time

Are These Equivalent?

Implementation 1

```
wait (sem s) {  
    s.n = s.n - 1;  
    if s.n < 0 {  
        addProc (me, s.L);  
        block (me);  
    }  
}  
  
signal (sem s) {  
    s.n = s.n + 1;  
    if (! empty (s.L)) {  
        p = removeProc (s.L);  
        unblock (p);  
    }  
}
```

Implementation 2

```
wait (sem s) {  
    if s.n ≤ 0 {  
        addProc (me, s.L);  
        block (me);  
    }  
    s.n = s.n - 1;  
}  
  
signal (sem s) { // same  
    s.n = s.n + 1;  
    if (! empty (s.L)) {  
        p = removeProc (s.L);  
        unblock (p);  
    }  
}
```

Summary

- Synchronization: process waiting for another
- Critical section: code allowing race condition
- Mutual exclusion: one process excludes others
- Mutual exclusion mechanism: obey four rules
- Peterson's solution: all software, but complex
- Semaphores: simple flexible synchronization
 - wait and signal must be atomic, thus requiring lower-level mutual exclusion (Peterson's, TSL)

Textbook

- OSP: Chapter 5
- OSC: Chapter 5 (Process Synchronization)
 - Lecture-related: 6.1-6.6

Supplementary

- For those who wish to understand more subtle and advanced issues
- Will not be on exams

Mutual Exclusion Using TSL

- Critical section entry code

```
                                ; assume lock initially 0
loop: TSL REG, lock            ; atomically {load REG with lock
                                ;           and store 1 into lock}
                                ;
                                CMP REG, #0          ; is REG (was lock) equal to 0?
                                JNE loop             ; if not equal to 0, check again
                                                ; also known as a "spin lock"
```

- Critical section exit code

```
MOV lock, #0                    ; reset lock to 0
```

Test and Test-and-Set

```
shared int lock = 0;
```

```
do {  
    while (lock == 1);           // cheap  
} while (! TSL(&lock));         // expensive  
< critical section >  
lock = 0;
```

- Busy-wait using simple reads of lock
 - Low overhead
- When lock opens, use test-and-set
 - Higher-overhead atomic operation less frequent

Efficient No-Spin Locking Code

- Entry code

```
while (! TSL(&lock)) {           // lock closed
    yield ();                     // give up CPU
}
```

- Exit code

```
lock = 0;                        // open lock
```


When is Busy-Waiting OK?

- Expected wait time $<$ scheduling overhead
- Lots of processors (i.e., if waste is OK)
- Blocking is not an option (e.g., inside kernel)

How Costly is Busy Waiting?

- Consider time spent in critical section
 - Chance of context switch increases with length
 - If switch to process seeking entry, it will busy wait
 - Wastes an entire quantum – this is cost
- So, try to minimize time in critical section
- Compare critical sections that are
 - user code (e.g., application code)
 - system code (e.g., semaphore operations)