

CSE 120: Principles of Operating Systems

Lecture 6: InterProcess Communication (IPC)

Prof. Joseph Pasquale

University of California, San Diego


February 8, 2023 (updated 3/11/23)

Cooperating Processes


- Why structure a computation as a set of cooperating (communicating) processes?
- Performance: speed
 - Exploit inherent parallelism of computation
 - Allow some parts to proceed while others do I/O
- Modularity: reusable self-contained programs
 - Each may do a useful task on its own
 - May also be useful as a sub-task for others

Examples of Cooperating Processes


Pipeline



Client/Server



Parent/Child




Inter-Process Communication

- To cooperate, need ability to communicate
- IPC: inter-process communication
 - Communication between processes
- IPC requires
 - *data transfer*
 - *synchronization*
- Need mechanisms for both

Three Abstractions for IPC

- Shared memory + semaphores
- Monitors
- Message passing

The Producer/Consumer Problem



- Producer produces data, inserts in shared buffer
- Consumer removes data from buffer, consumes it
- Cooperation: Producer feeds Consumer
 - How does data get from Producer to Consumer?
 - How does Consumer wait for Producer?

Producer/Consumer: Shared Memory

```
shared int buf[N], in = 0, out = 0;
```

Producer

```
while (TRUE) {  
    buf[in] = Produce ();  
    in = (in + 1)%N;  
}
```

Consumer

```
while (TRUE) {  
    Consume (buf[out]);  
    out = (out + 1)%N;  
}
```

- No synchronization
 - Consumer must wait for something to be produced
 - What about Producer?
- No mutual exclusion for critical sections
 - Relevant if multiple producers or multiple consumers

Recall Semaphores

- Semaphore: synchronization variable
 - Takes on integer values
 - Has an associated list of waiting processes
- Operations
 - wait (s) { $s = s - 1$; block if $s < 0$ }
 - signal (s) { $s = s + 1$; unblock a process if any }
- No other operations allowed (e.g., can't test s)

Semaphores for Synchronization

```
shared int buf[N], in = 0, out = 0;  
sem filledslots = 0, emptyslots = N;
```

Producer

```
while (TRUE) {  
    wait (emptyslots);  
    buf[in] = Produce ();  
    in = (in + 1)%N;  
    signal (filledslots);  
}
```

Consumer

```
while (TRUE) {  
    wait (filledslots);  
    Consume (buf[out]);  
    out = (out + 1)%N;  
    signal (emptyslots);  
}
```

- Buffer empty, Consumer waits
- Buffer full, Producer waits
- General synchronization vs. mutual exclusion

Multiple Producers

```
shared int buf[N], in = 0, out = 0;  
sem filledslots = 0, emptyslots = N;
```

Producer1

```
while (TRUE) {  
    wait (emptyslots);  
    buf[in] = Produce ();  
    in = (in + 1)%N;  
    signal (filledslots);  
}
```

Producer2

```
while (TRUE) {  
    wait (emptyslots);  
    buf[in] = Produce ();  
    in = (in + 1)%N;  
    signal (filledslots);  
}
```

Consumer

```
while (TRUE) {  
    wait (filledslots);  
    Consume (buf[out]);  
    out = (out + 1)%N;  
    signal (emptyslots);  
}
```

- There is a race condition in the Producer code
- Inconsistent updating of variables buf and in
- Need mutual exclusion

Semaphore for Mutual Exclusion

```
shared int buf[N], in = 0, out = 0;  
sem filledslots = 0, emptyslots = N, mutex = 1;
```

Producer1, 2, ...

```
while (TRUE) {  
    wait (emptyslots);  
    wait (mutex);  
    buf[in] = Produce ();  
    in = (in + 1)%N;  
    signal (mutex);  
    signal (filledslots);  
}
```

Consumer1, 2, ...

```
while (TRUE) {  
    wait (filledslots);  
    wait (mutex);  
    Consume (buf[out]);  
    out = (out + 1)%N;  
    signal (mutex);  
    signal (emptyslots);  
}
```

- Works for multiple producers and consumers
- But not easy to understand: can lead to bugs
 - example: what if wait statements interchanged?

Allowing For More Parallelism

```
shared int buf[N], in = 0, out = 0;  
sem filledslots = 0, emptyslots = N, pmutex=1, cmutex=1;
```

Producer1, 2, ...

```
while (TRUE) {  
    wait (emptyslots);  
    wait (pmutex);  
    buf[in] = Produce ();  
    in = (in + 1)%N;  
    signal (pmutex);  
    signal (filledslots);  
}
```

Consumer1, 2, ...

```
while (TRUE) {  
    wait (filledslots);  
    wait (cmutex);  
    Consume (buf[out]);  
    out = (out + 1)%N;  
    signal (cmutex);  
    signal (emptyslots);  
}
```

- Separate producer and consumer mutexes
- Decreases consumer/producer dependencies
 - increases parallelism

Monitors

- Programming language construct for IPC
 - Variables (shared) requiring controlled access
 - Accessed via procedures (mutual exclusion)
 - Condition variables (general synchronization)
 - wait (cond): block until another process signals cond
 - signal (cond): unblock a process waiting on cond
- Only one process can be active inside monitor
 - Active = running or able to run; others must wait

Producer/Consumer using a Monitor

```
monitor ProducerConsumer {
    int buf[N], in = 0, out = 0, count = 0;
    cond slotavail, itemavail;

    PutItem (int item) {
        if (count == N)
            wait (slotavail);
        buf[in] = item;
        in = (in + 1)%N;
        count++;
        signal (itemavail);
    }

    GetItem () {
        int item;
        if (count == 0)
            wait (itemavail);
        item = buf[out];
        out = (out + 1)%N;
        count--;
        signal (slotavail);
        return (item);
    }
}
```

Producer

```
while (TRUE) {
    PutItem (Produce ());
}
```

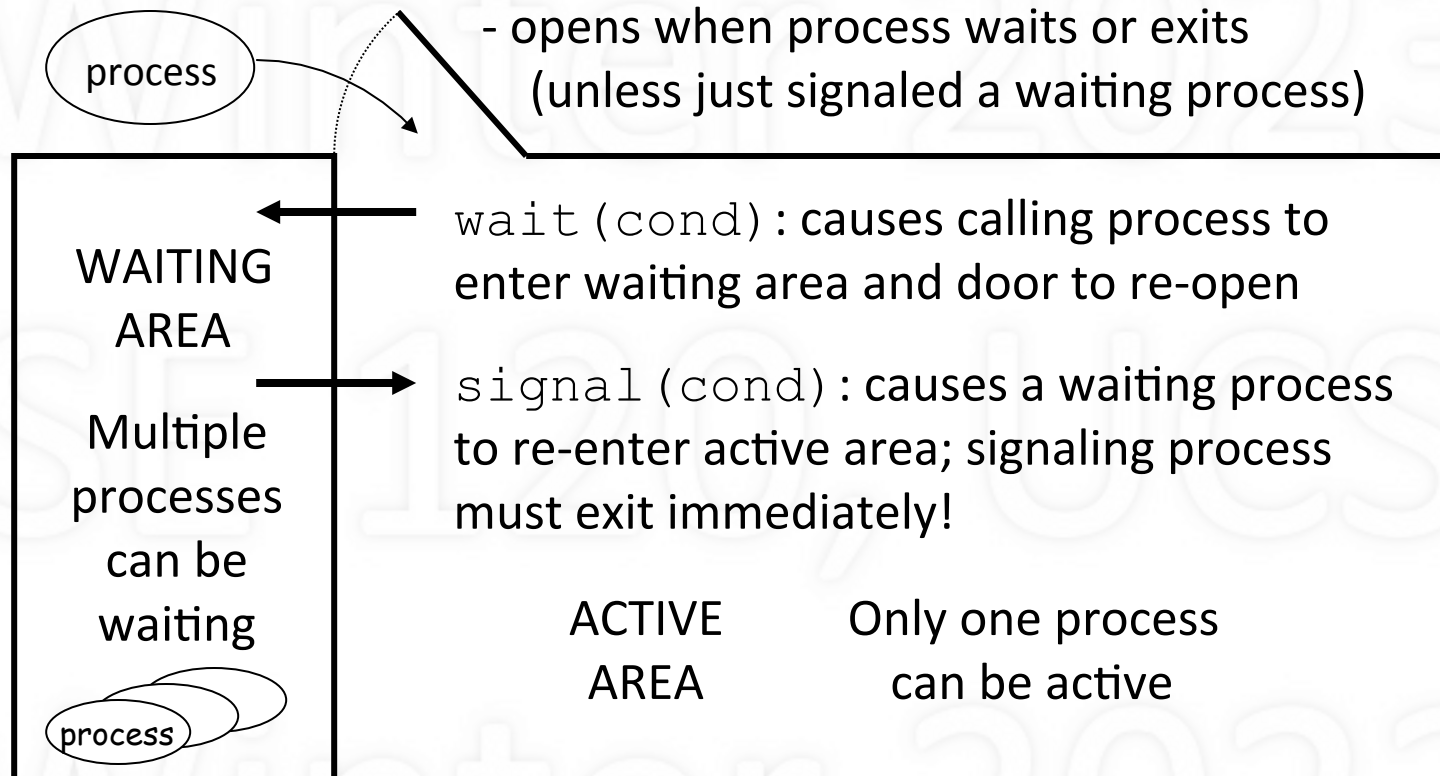
Consumer

```
while (TRUE) {
    Consume (GetItem ());
}
```

How Synchronization Works

Door with “monitor lock” enforces mutual exclusion:


- open if no process *active* in monitor
- closes when process enters
- opens when process waits or exits
(unless just signaled a waiting process)



Issues with Monitors

- Given P_1 waiting on condition c , P_2 signals c
 - P_1 and P_2 able to run: breaks mutual exclusion
 - One solution: signal just before returning
- Condition variables have no memory
 - Signal without someone waiting does nothing
 - Signal is “lost” (no memory, no future effect)
- Monitors bring structure to IPC
 - Localizes critical sections and synchronization

Message Passing



- Two methods
 - `send (destination, &message)`
 - `receive (source, &message)`
- Data transfer: in to and out of kernel message buffers
- Synchronization: receive blocks to wait for message

Producer/Consumer: Message-Passing

```
/* NO SHARED MEMORY */
```

Producer

```
int item;
```

```
while (TRUE) {  
    item = Produce ();  
    send (Consumer, &item);  
}
```

Consumer

```
int item;
```

```
while (TRUE) {  
    receive (Producer, &item);  
    Consume (item);  
}
```

With Flow Control

Producer

```
int item;
```

```
while (TRUE) {  
    receive (Consumer, &item);  
    item = Produce ();  
    send (Consumer, &item);  
}
```

Consumer

```
int item;
```

```
do N times {  
    send (Producer, &item);  
}
```

```
while (TRUE) {  
    receive (Producer, &item);  
    Consume (item);  
    send (Producer, &item);  
}
```

An Optimization

Producer

```
int item, empty;
```

```
while (TRUE) {  
    item = Produce ();  
    receive (Consumer, &empty);  
    send (Consumer, &item);  
}
```

Consumer

```
int item, empty;
```

```
do N times {  
    send (Producer, &empty);  
}
```

```
while (TRUE) {  
    receive (Producer, &item);  
    send (Producer, &empty);  
    Consume (item);  
}
```

Issues with Message Passing

- Who should messages be addressed to?
 - ports (“mailboxes”) rather than processes
- How to make process receive from anyone?
 - `pid = receive (*, &message)`
- Kernel buffering: outstanding messages
 - messages sent that haven’t been received yet
- Good paradigm for IPC over networks
- Safer than shared memory paradigms

Textbook

- OSP: Chapter 6
- OSC: Chapters 3, 6 (Process Synchronization)
 - Focus on semaphores, monitors, message-passing
 - Lecture-related: 3.4, 6.6, 6.7
 - Recommended: 3.5, 3.6