

CSE 120: Principles of Operating Systems

Lecture 9: Logical Memory

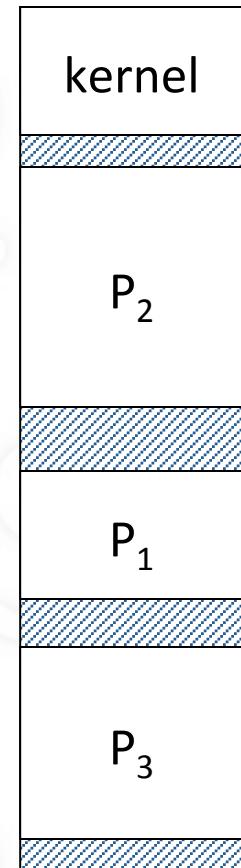
Prof. Joseph Pasquale
University of California, San Diego
February 27, 2023

What is Logical Memory

- Logical memory = a process's memory
- As viewed (referenced) by a process
- Allocated without regard to physical memory

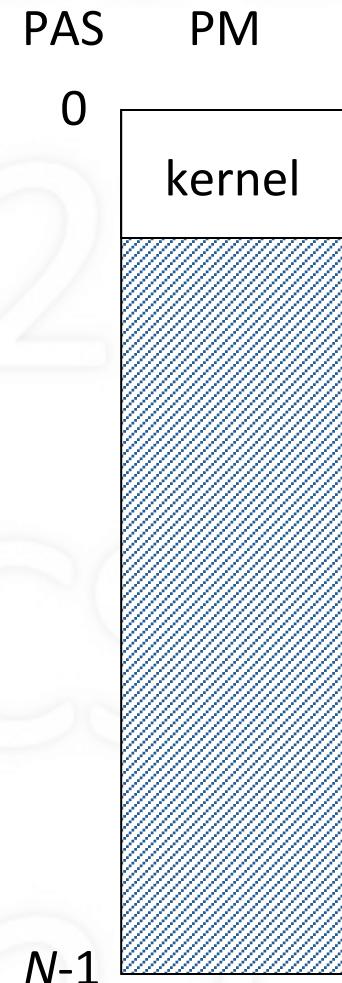
Problems with Sharing Memory

- The Addressing Problem
 - Compiler generates memory references
 - Unknown where process will be located
- The Protection Problem
 - Modifying another process's memory
- The Space Problem
 - The more processes there are, the less memory each individually can have



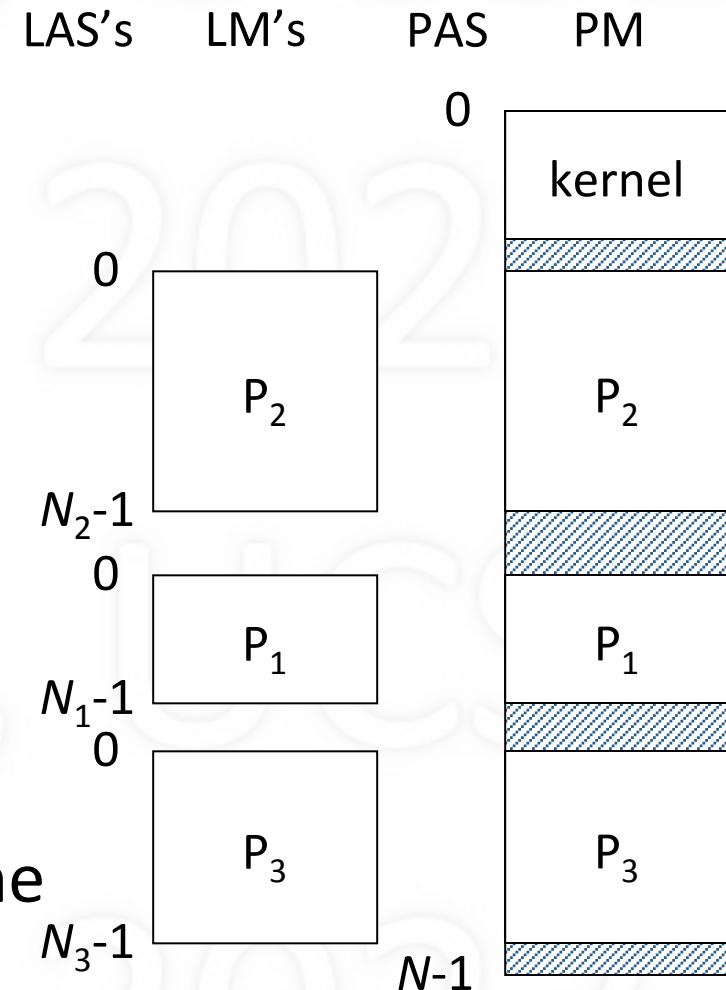
Address Spaces

- Address space
 - Set of addresses for memory
- Usually linear: 0 to $N-1$ (size N)
- Physical Address Space
 - 0 to $N-1$, $N = \text{size}$
 - Kernel occupies lowest addresses (typically)



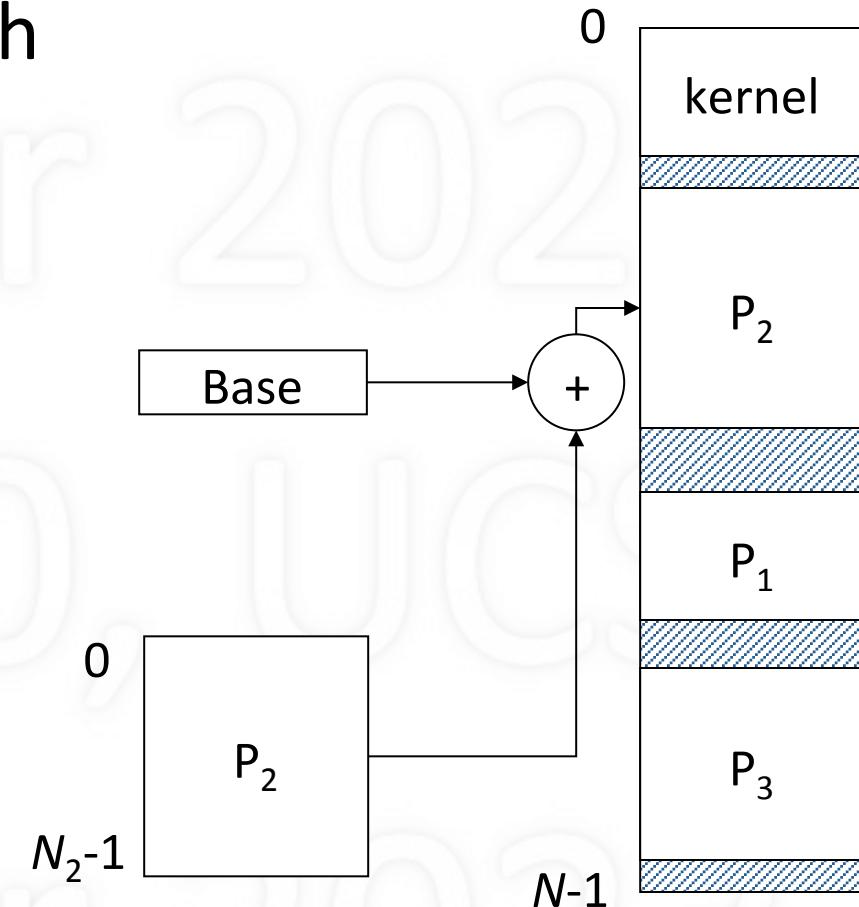
Logical vs. Physical Addressing

- Logical addresses
 - Assumes separate memory starting at 0
 - Compiler generated
 - Independent of location in physical memory
- Convert logical to physical
 - Via software: at load time
 - Via hardware: at access time



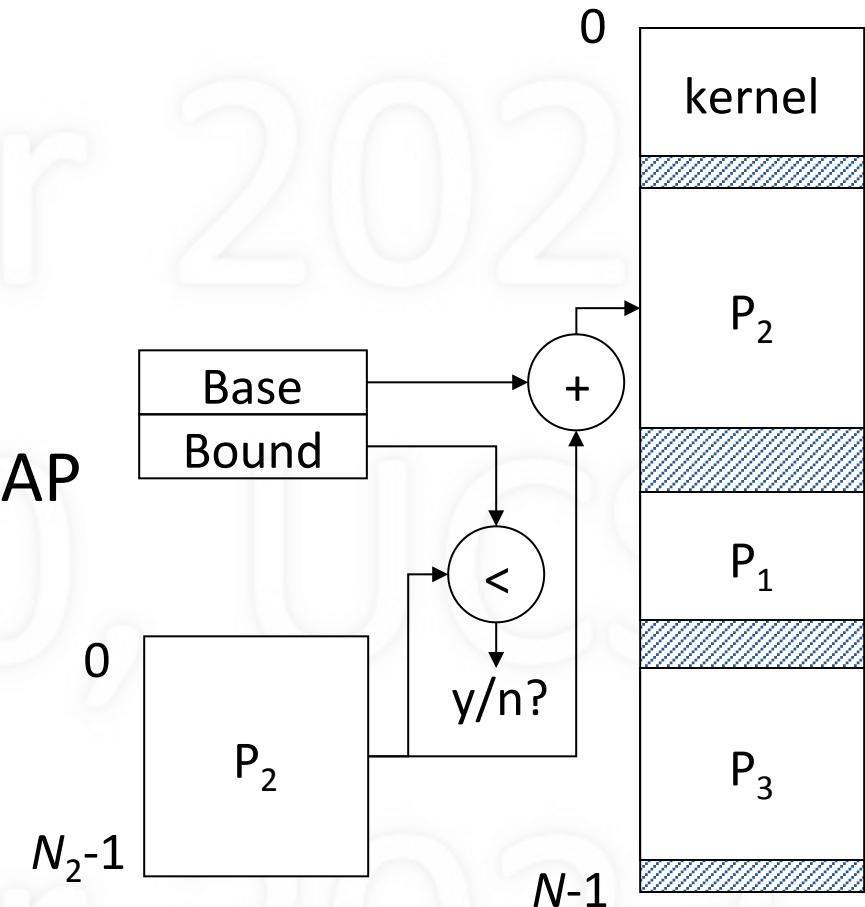
Hardware for Logical Addressing

- Base register filled with start address
- To translate logical address, add base
- Achieves relocation
- To move process: change base



Protection

- Bound register works with base register
- Is address < bound
 - Yes: add to base
 - No: invalid address, TRAP
- Achieves protection

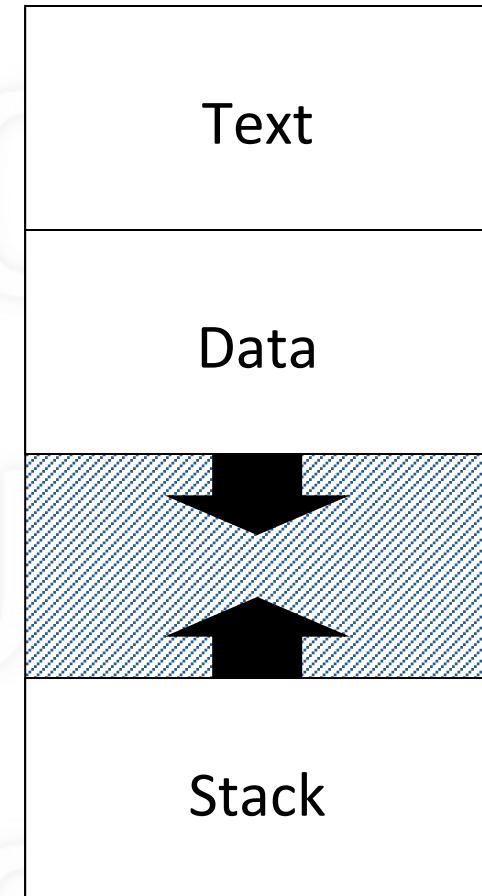


Memory Registers Part of Context

- On Every Context Switch
 - Load base/bound registers for selected process
 - Only kernel does loading of these registers
 - Kernel must be protected from all processes
- Benefit
 - Allows each process to be separately located
 - Protects each process from all others

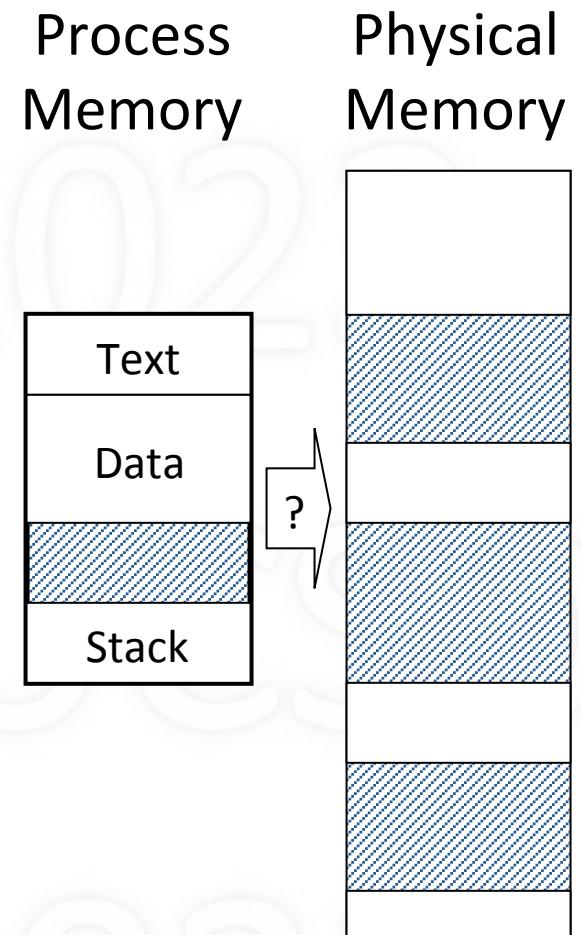
Recall Process Address Space

- Text: program instructions
 - Execute-only, fixed size
- Data: variables (static, heap)
 - Read/write, variable size
 - Dynamic allocation by request
- Stack: activation records (auto)
 - Read/write, variable size
 - Automatic growth/shrinkage



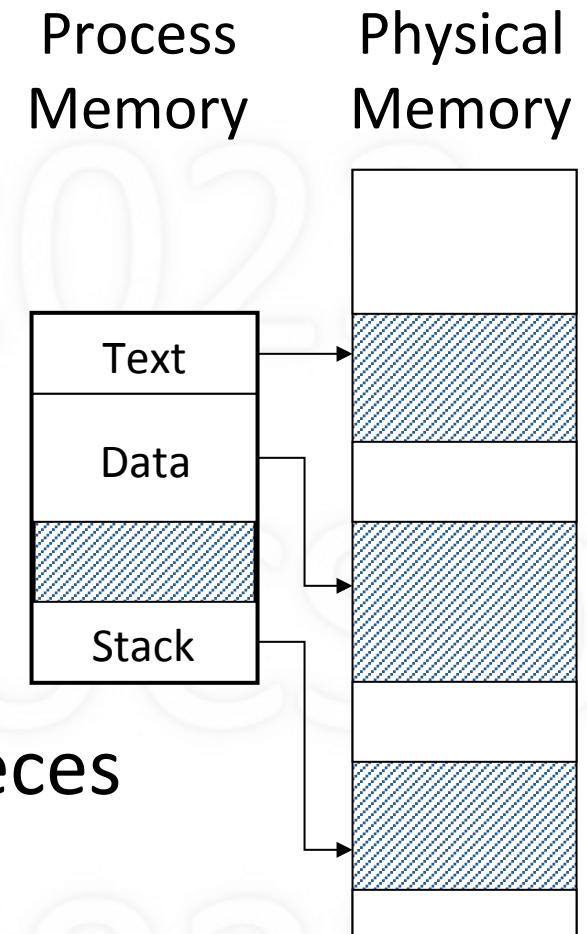
Fitting Process into Memory

- Must find large enough hole
- May not succeed
 - even if enough fragment space
- Even if successful, inefficient
 - Space must be allocated for potential growth areas



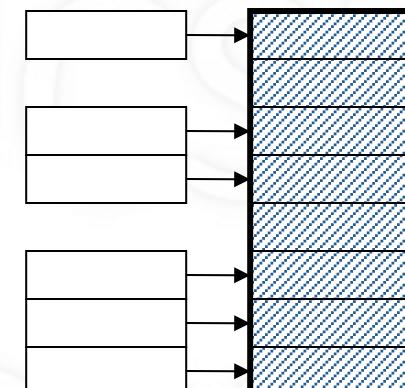
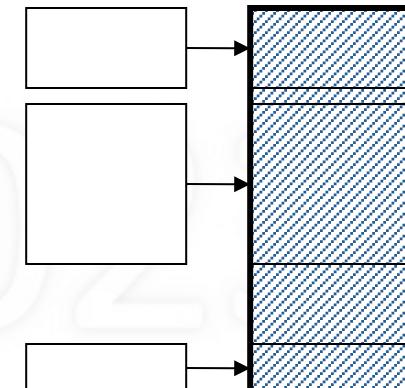
Fitting Process into Memory

- Must find large enough hole
- May not succeed
 - even if enough fragment space
- Even if successful, inefficient
 - Space must be allocated for potential growth areas
- Solution: break process into pieces
 - Distribute into available holes



Two Approaches

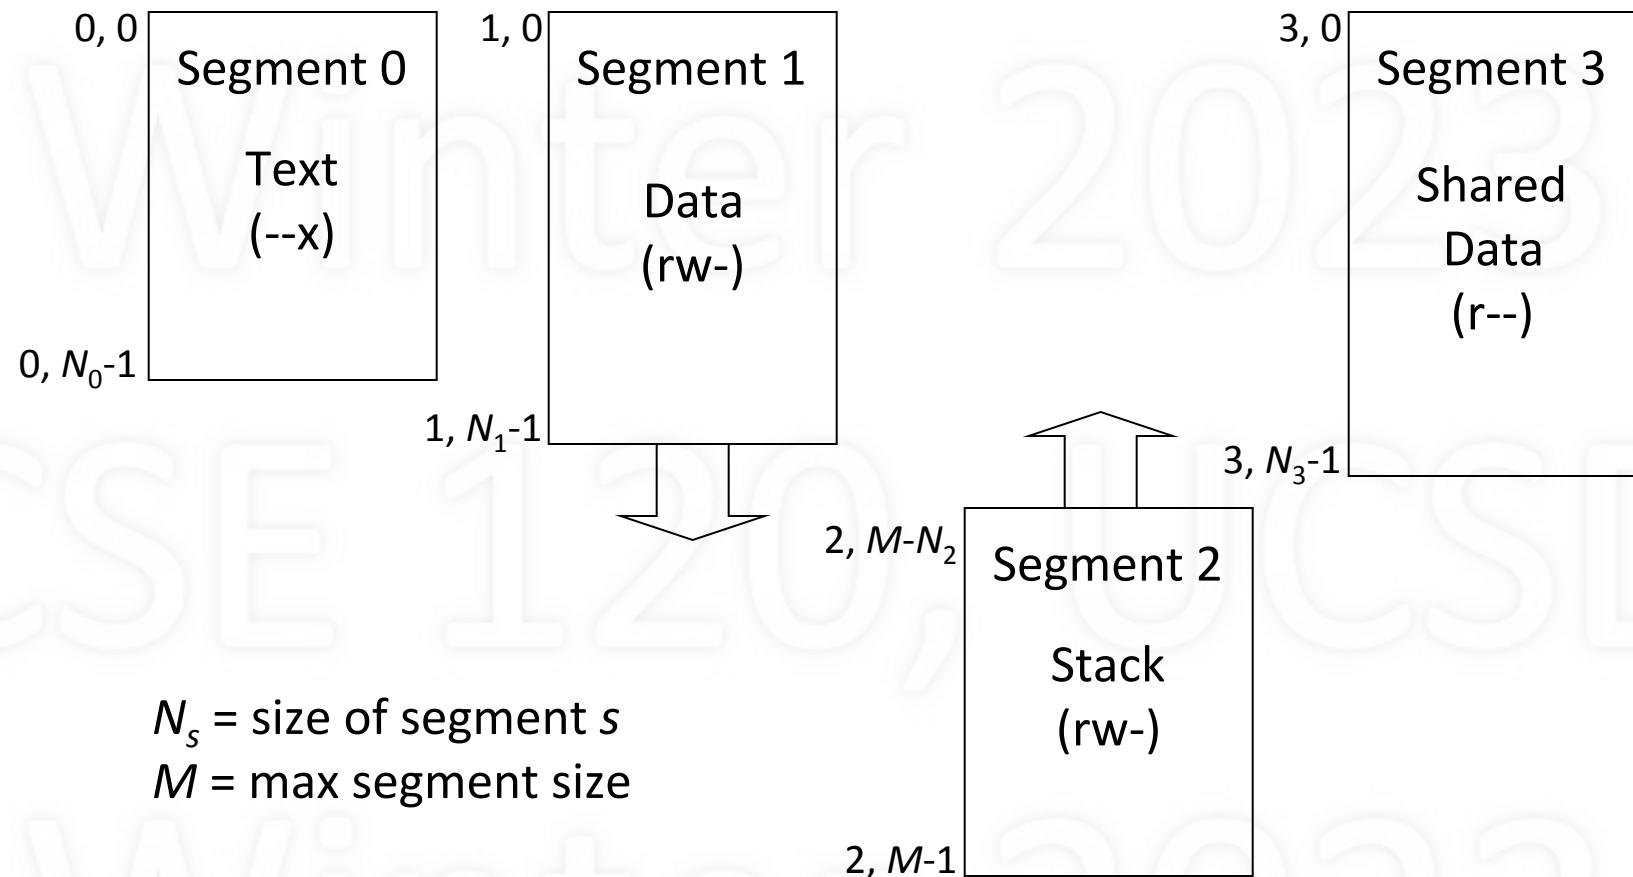
- Segmented address space
 - Partition into segments
 - Segments can be different sizes
- Paged address space
 - Partition into pages
 - Corresponding memory: frames
 - Pages are the same size



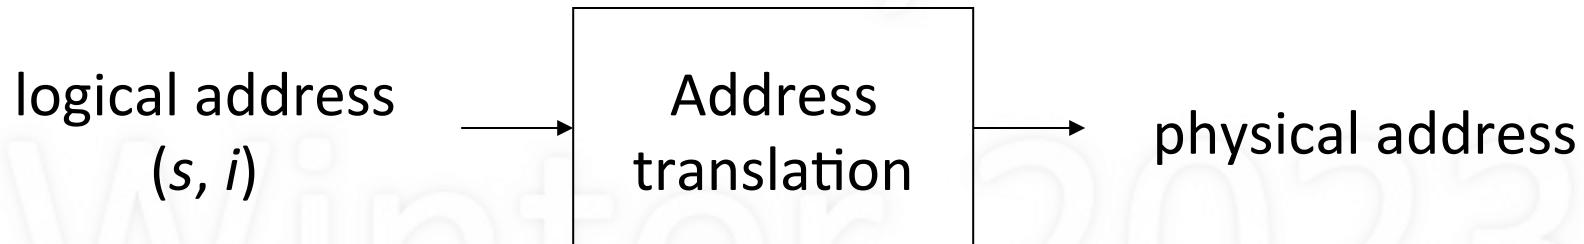
Segmented Address Space

- Address space is a set of segments
- Segment: a linearly addressed memory
 - Typically contains logically-related information
 - Examples: program code, data, stack
- Each segment has an identifier s , and a size N_s
 - s between 0 and $S-1$, $S = \text{max number of segments}$
- Logical addresses are of the form (s, i)
 - Offset i within segment s , i must be less than N_s

Example: Segmented Address Space



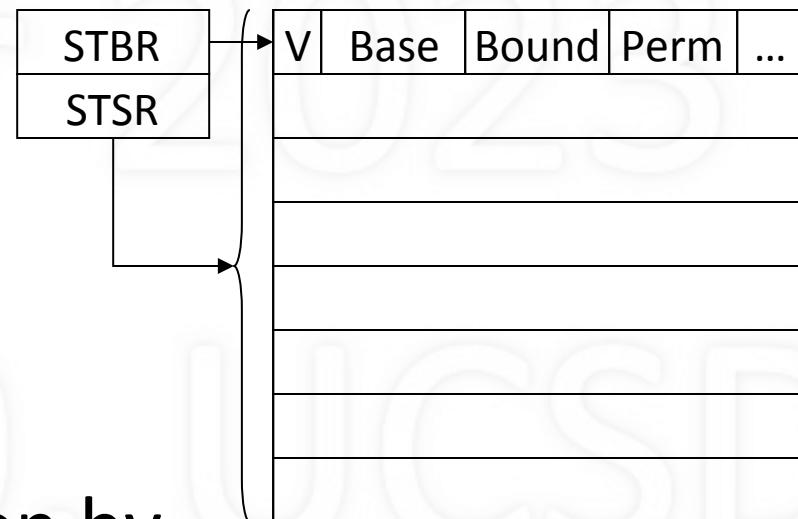
Segment-based Address Translation



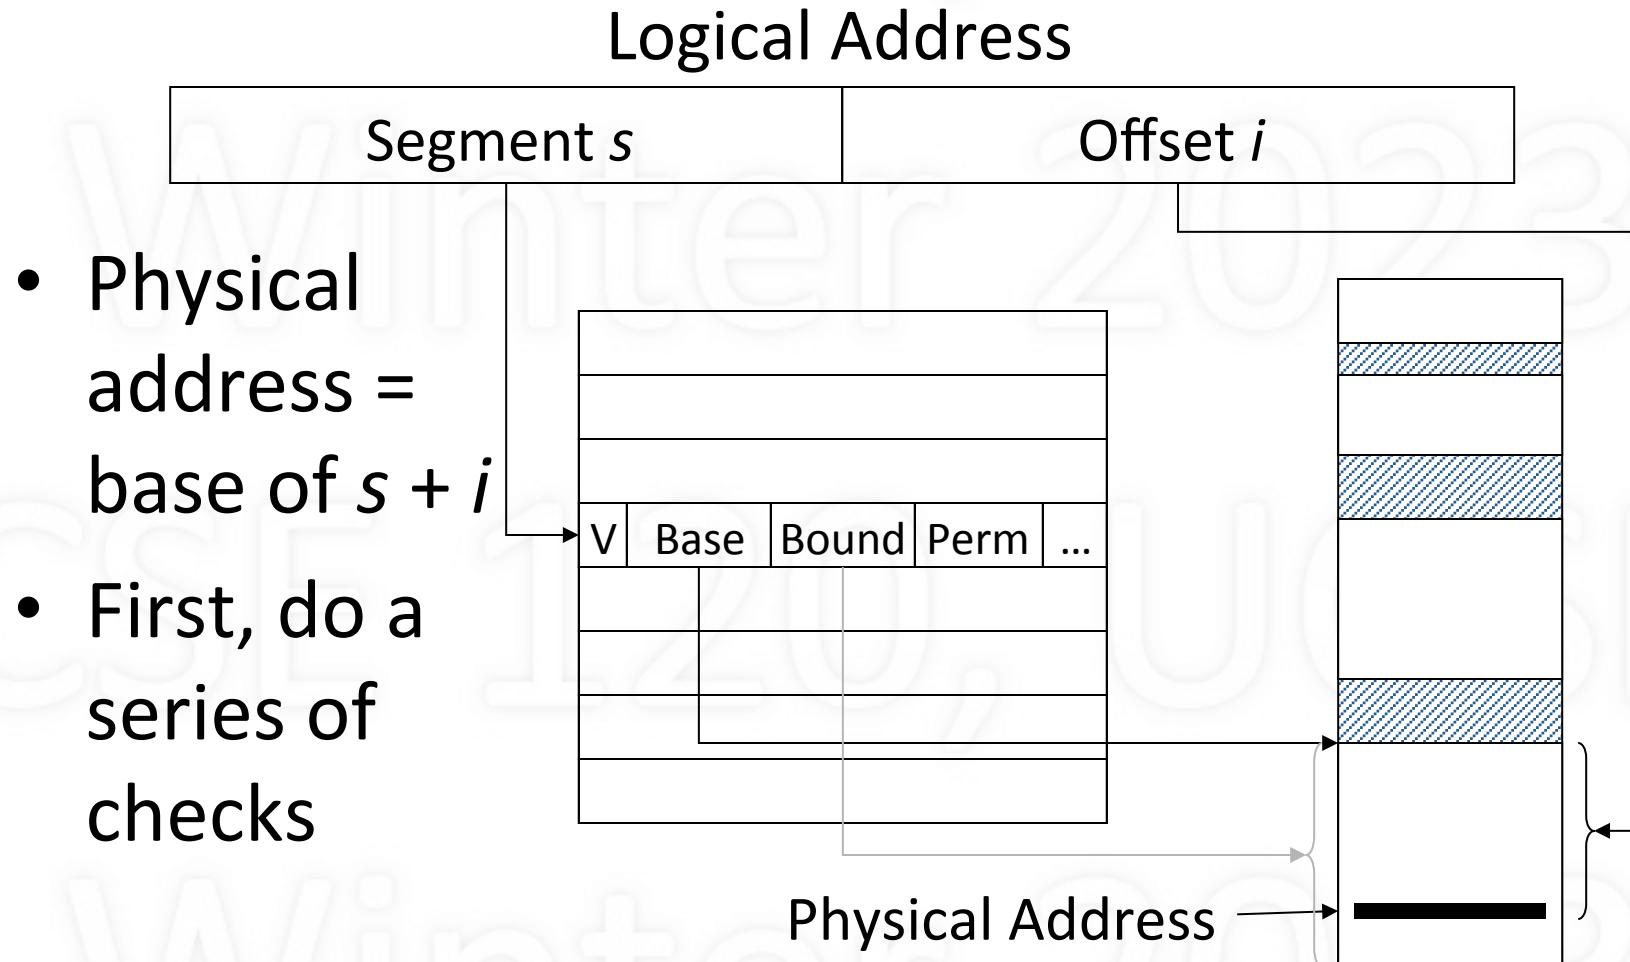
- Problem: how to translate
 - a logical address (s, i) into physical address a ?
- Solution: use a segment (translation) table ST
 - to segment s into base physical address $b = ST(s)$
 - then add b and i
- Summary: physical address = $ST(s) + i$

Segment Table

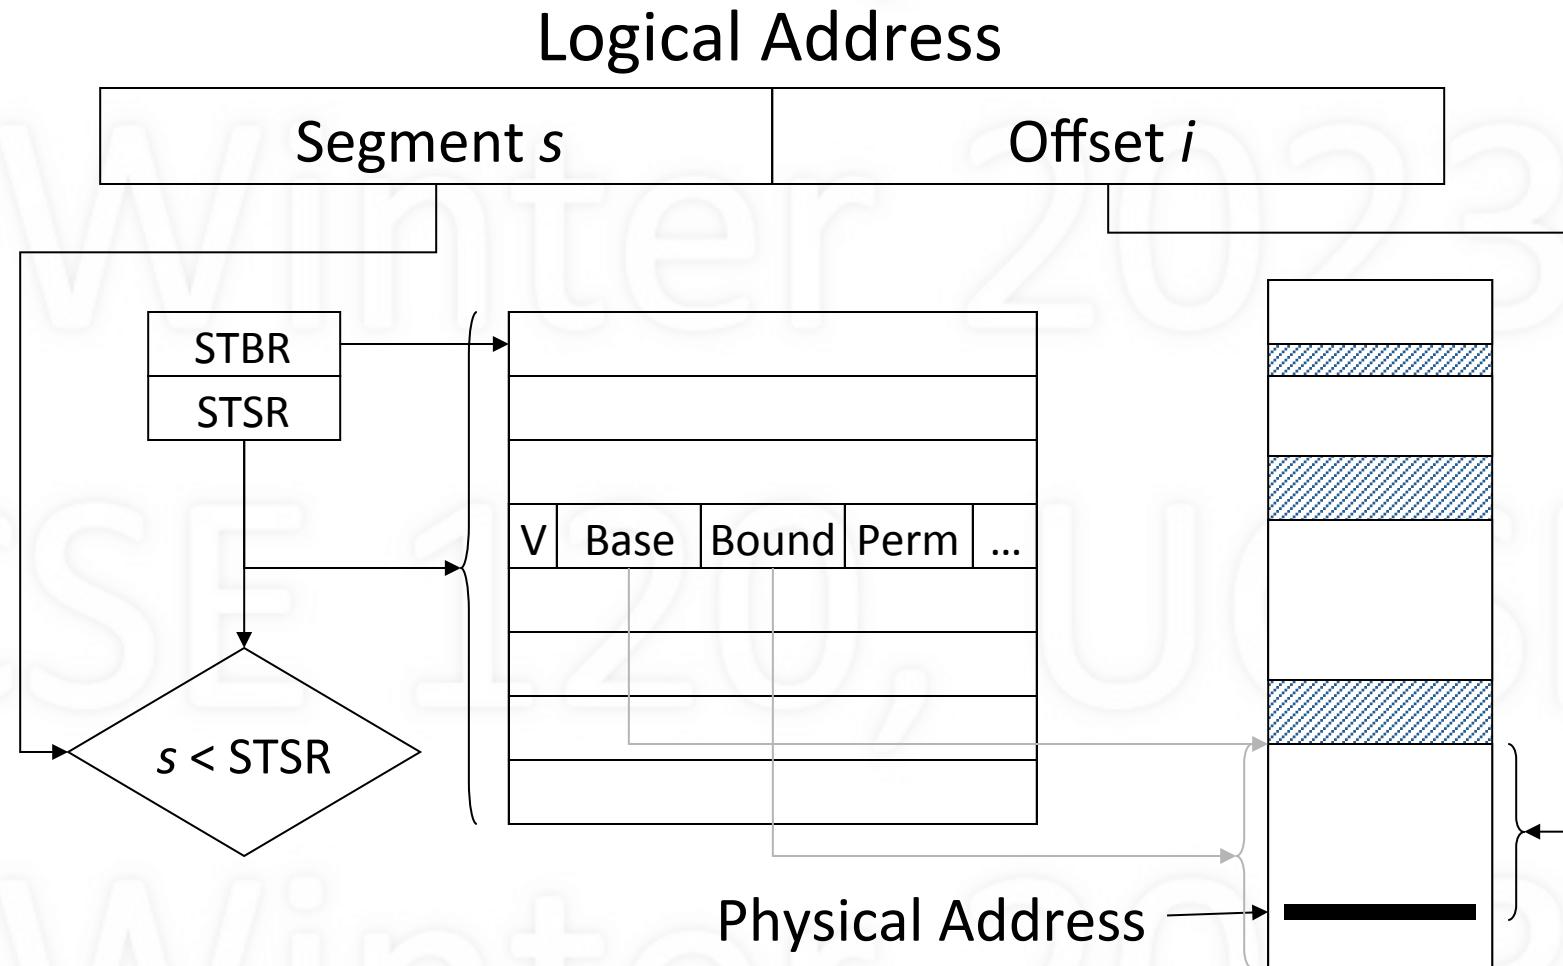
- One per process (typically)
- Table entry elements
 - V: valid bit
 - Base: segment location
 - Bound: segment size
 - Perm: permissions
- Location in memory given by
 - Segment table base register (hardware)
 - Segment table size register (hardware)



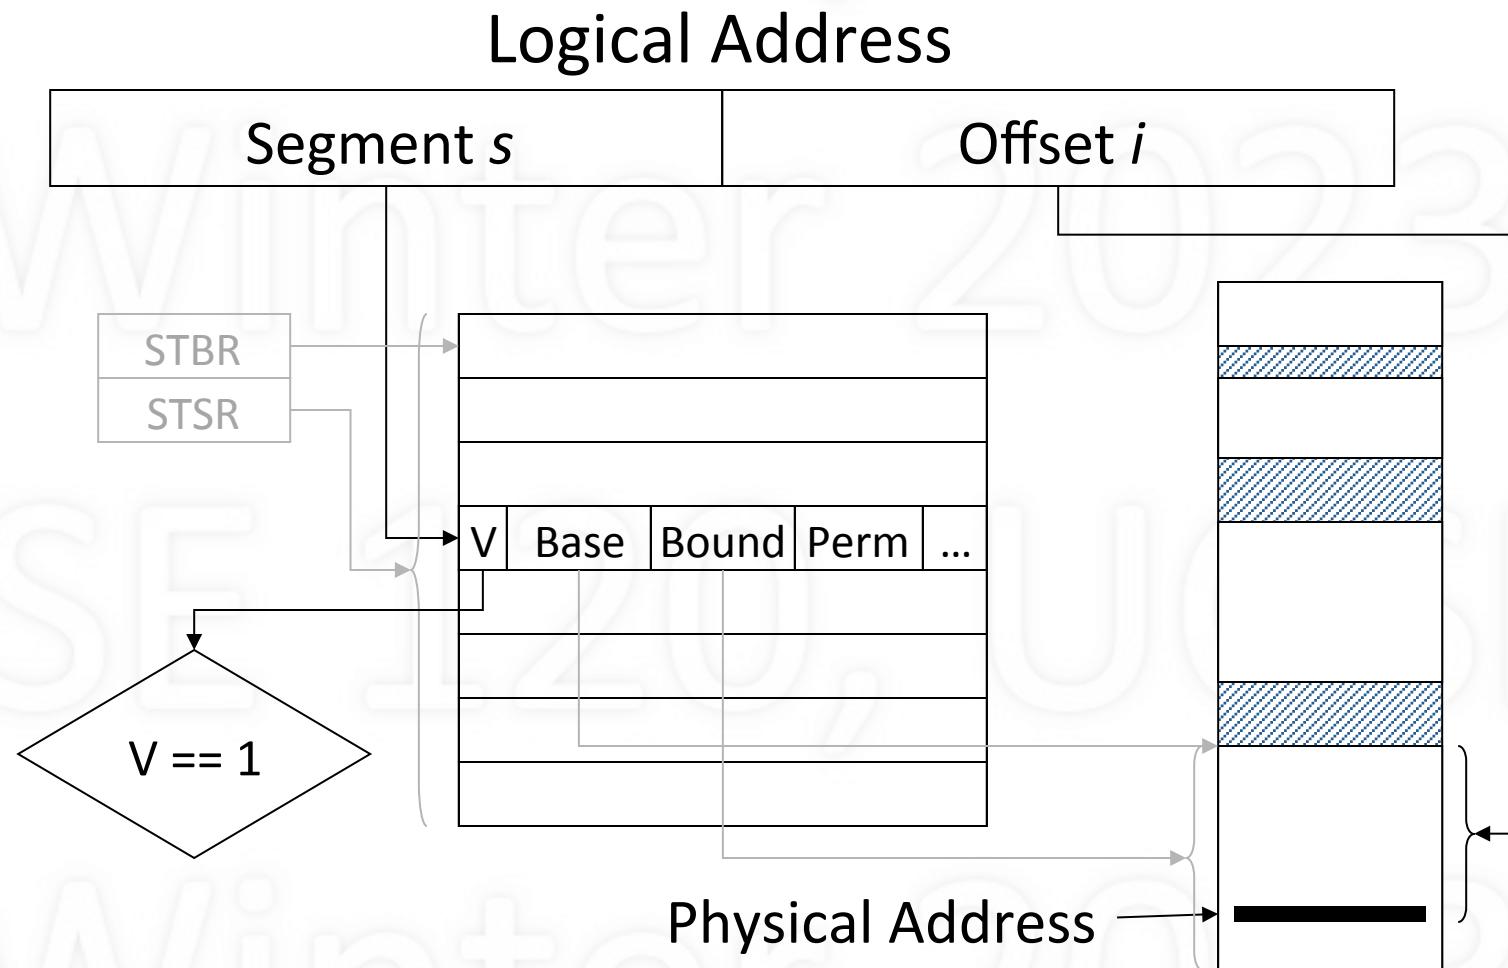
Address Translation



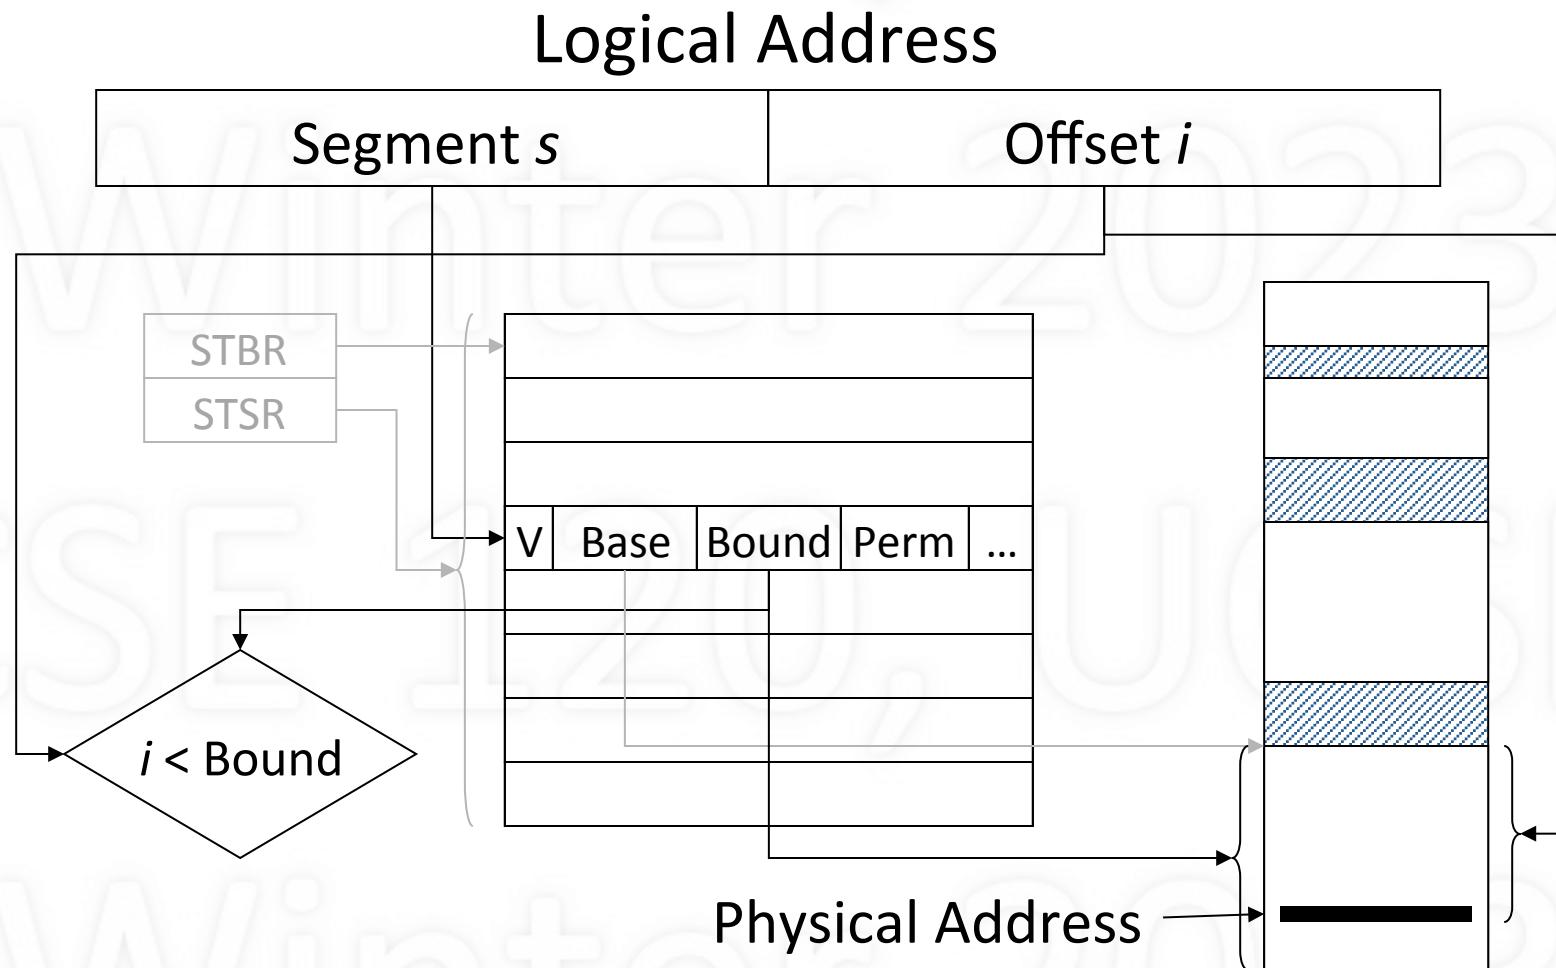
Check if Segment s is within Range



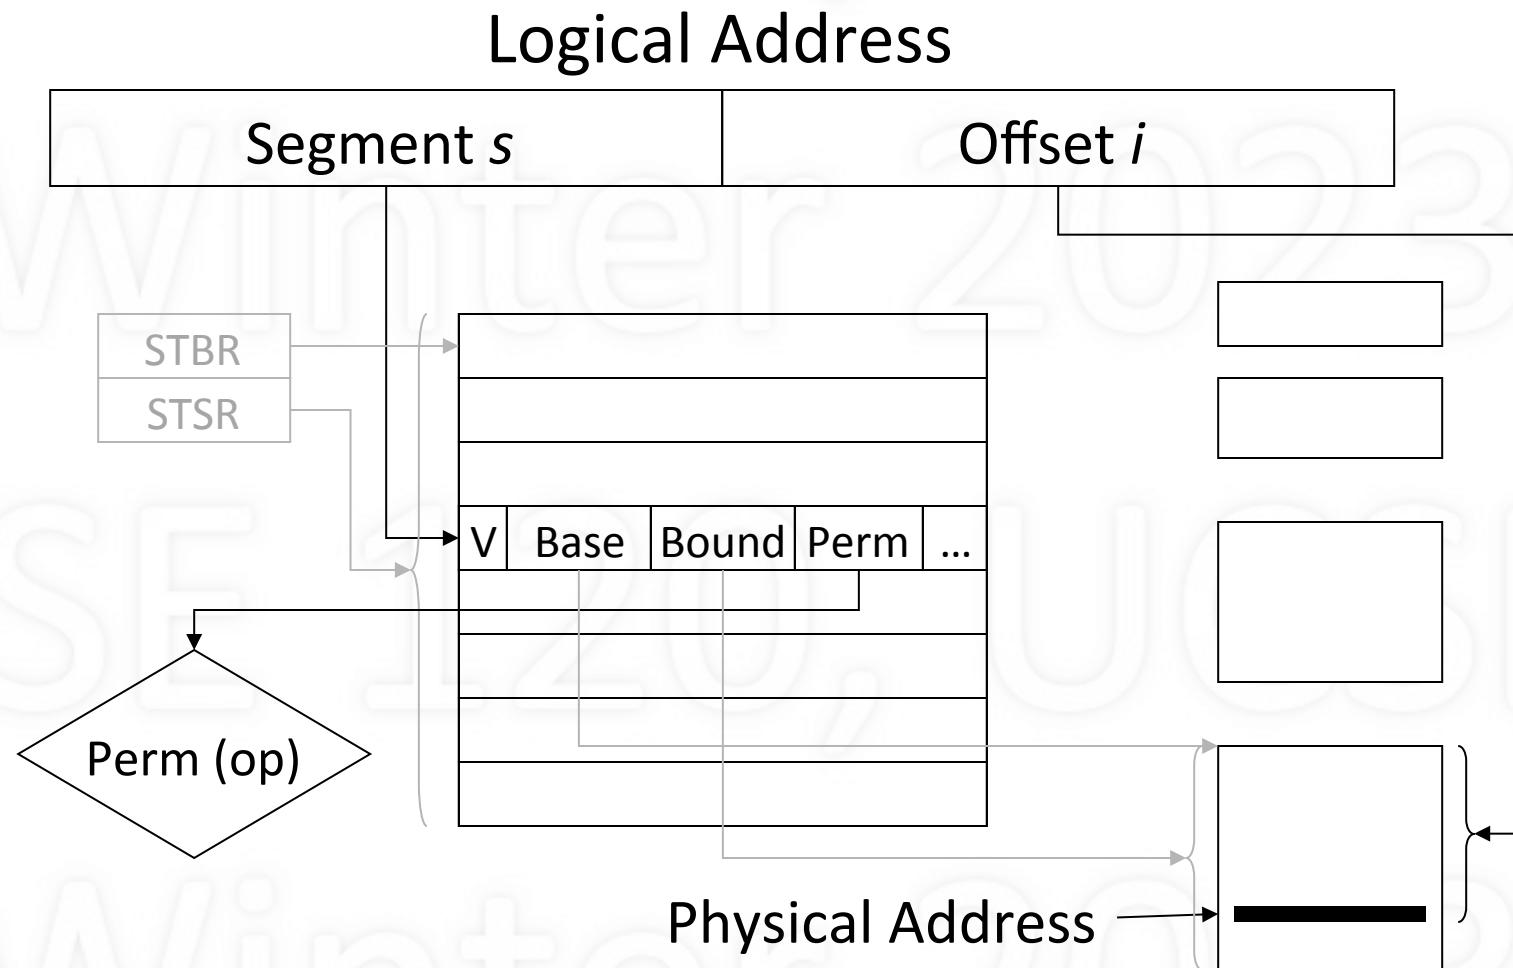
Check if Segment Entry s is Valid



Check if Offset i is within Bounds

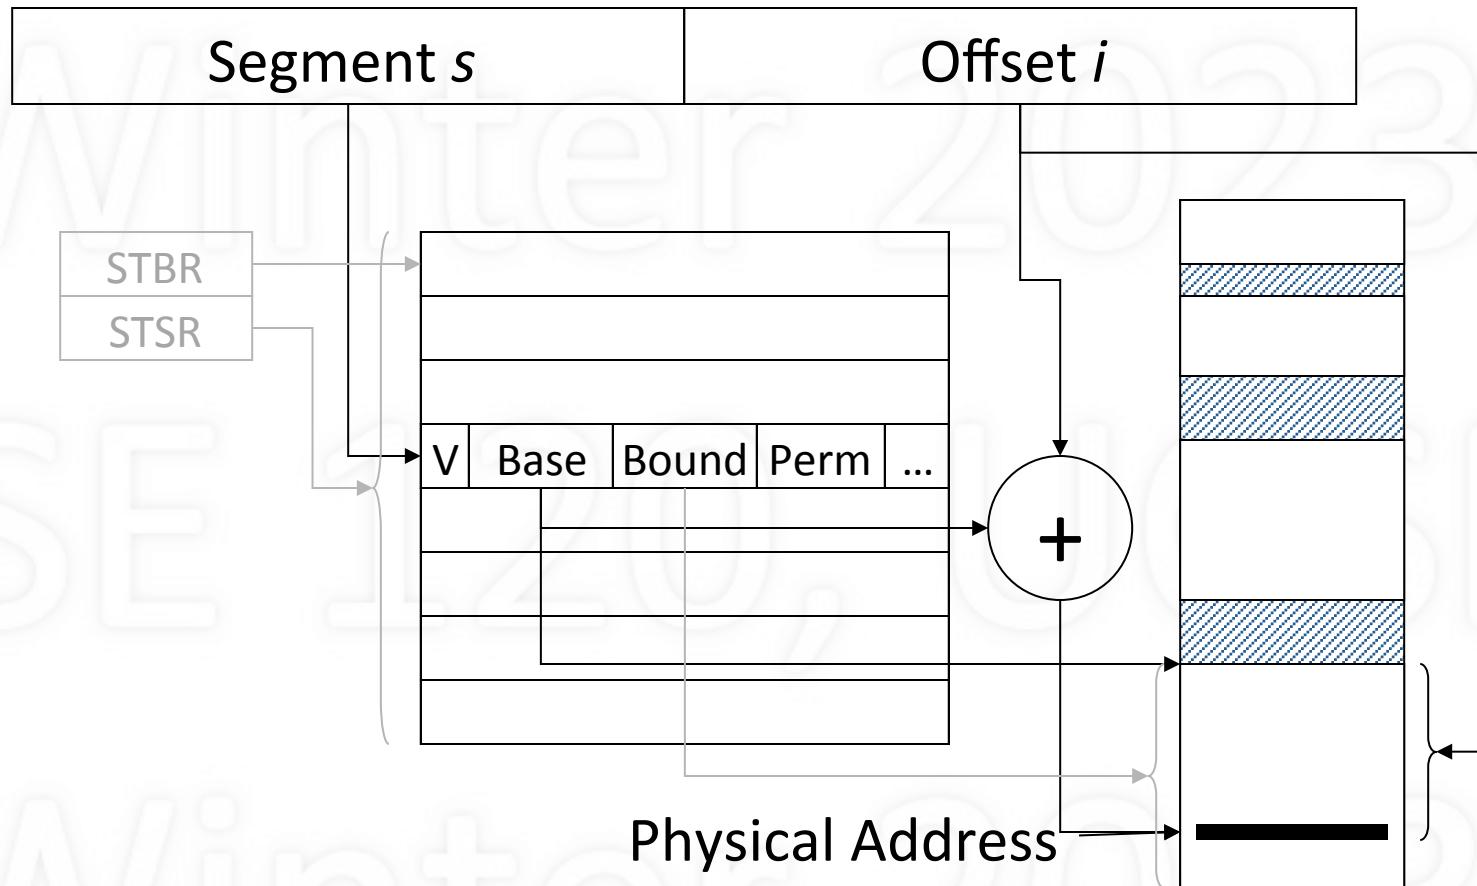


Check if Operation is Permitted

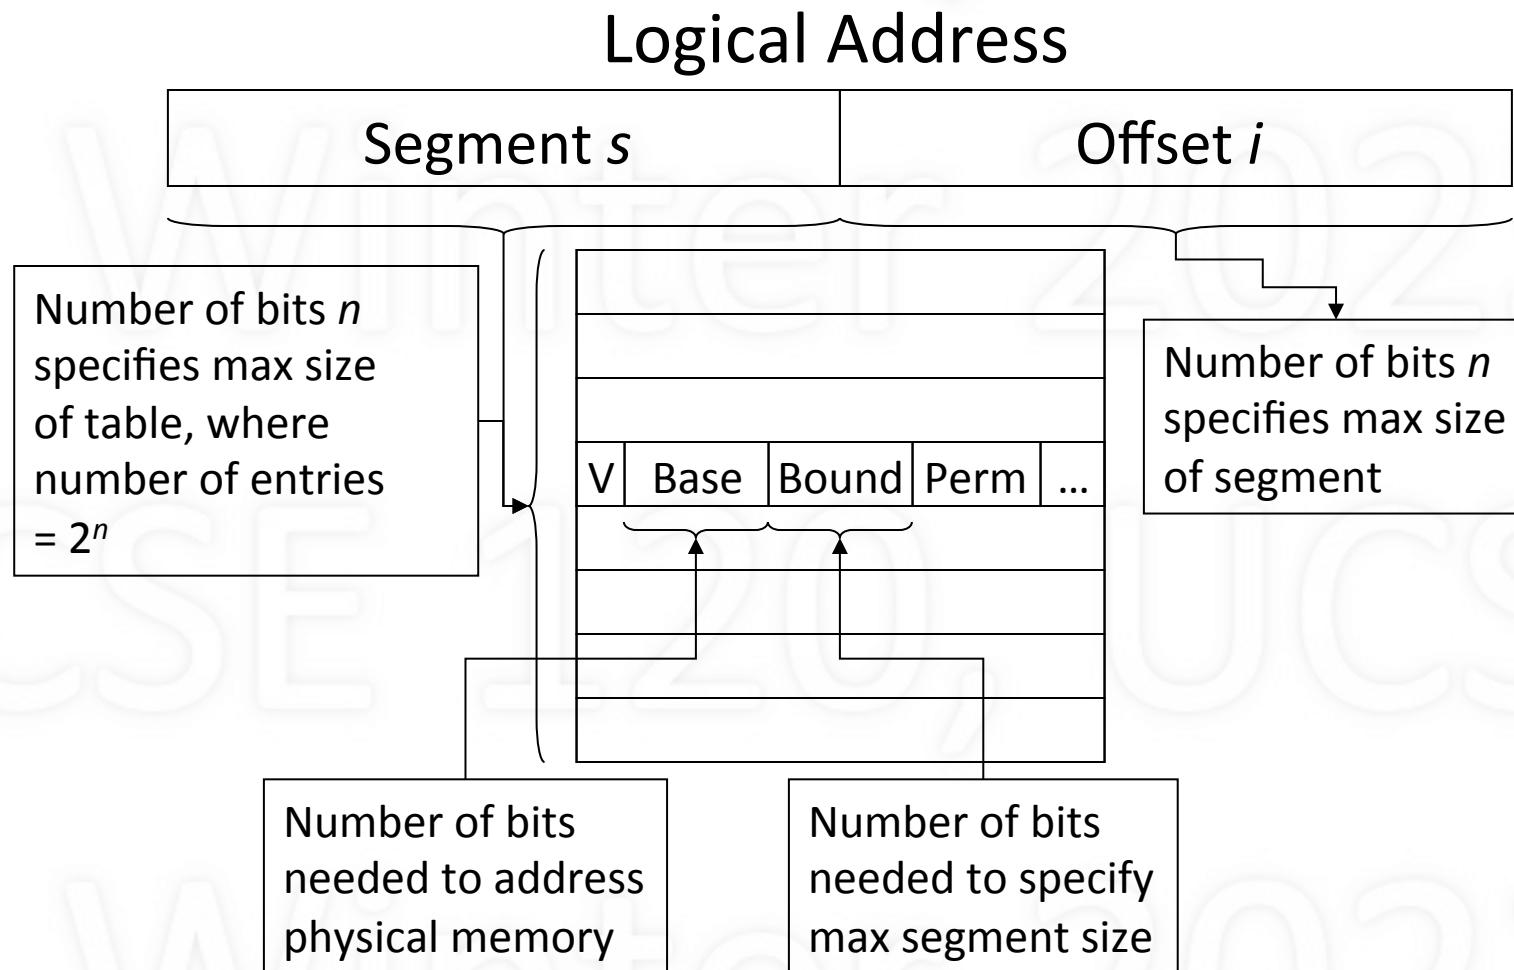


Translate Address

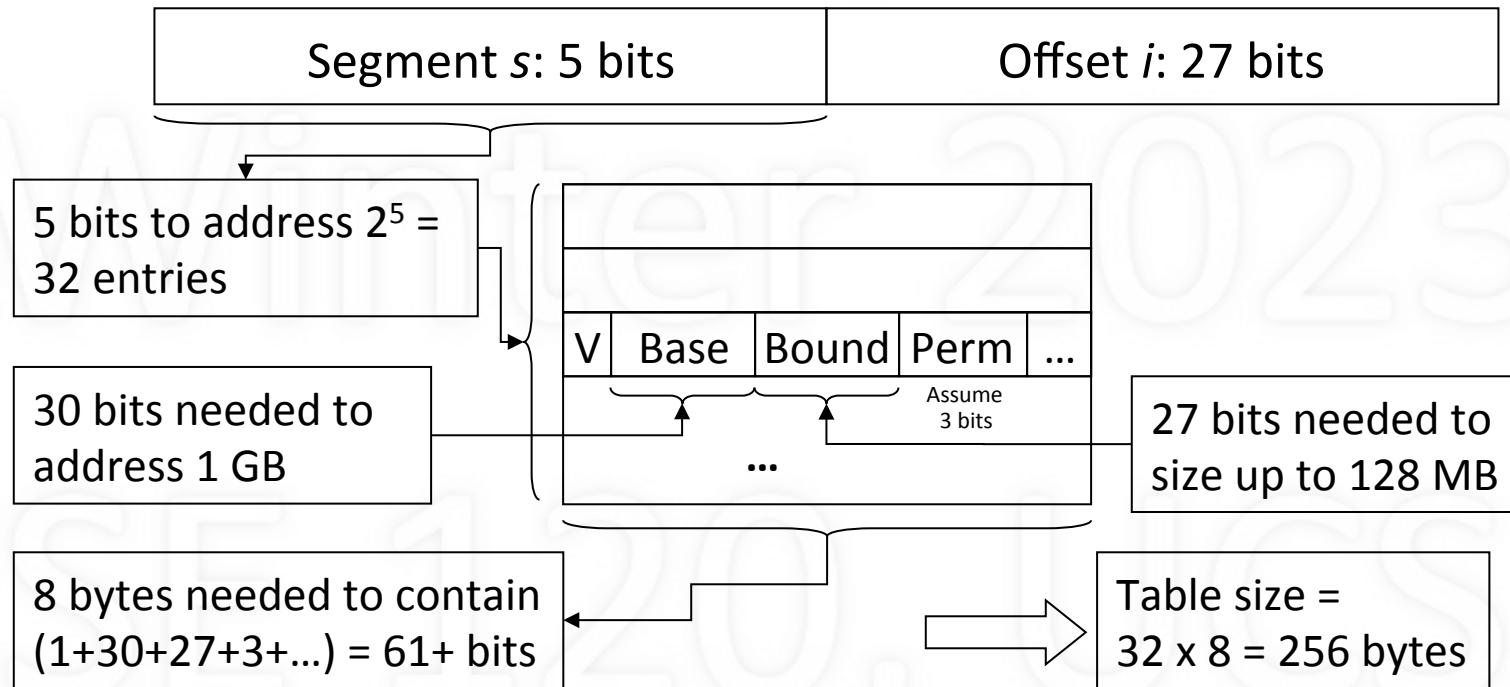
Logical Address



Sizing the Segment Table



Example of Sizing the Segment Table



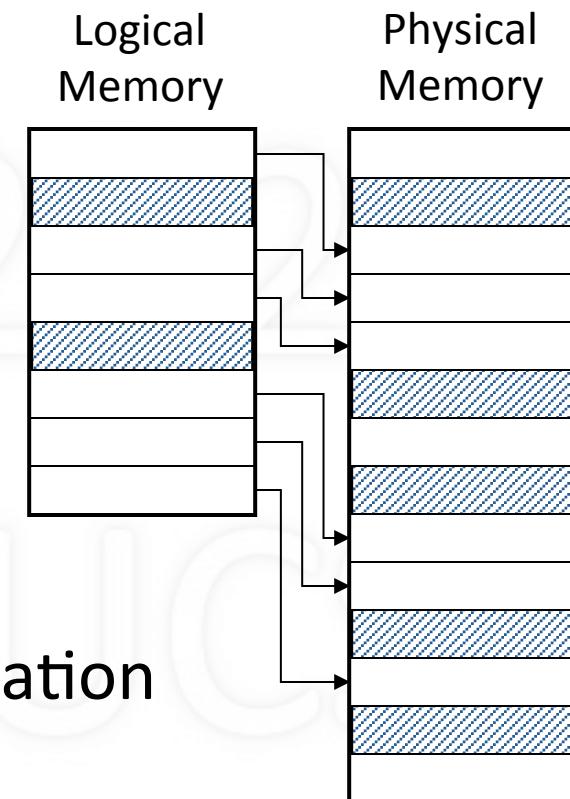
- Given 32 bit logical, 1 GB physical memory (max)
 - 5 bit segment number, 27 bit offset

Pros and Cons of Segmentation

- Pro: Each segment can be
 - located independently
 - separately protected
 - grown/shrunk independently
- Pro: Segments can be shared by processes
- Con: Variable-size allocation
 - Difficult to find holes in physical memory
 - External fragmentation

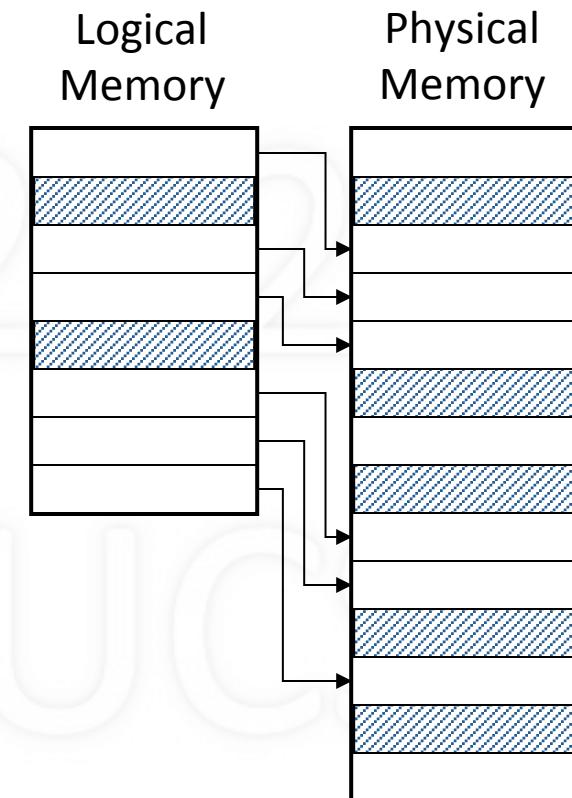
Paged Address Space

- Logical (process) memory
 - Linear sequence of pages
- Physical memory
 - Linear sequence of frames
- Pages and frames
 - Frame: a physical unit of information
 - A page fits exactly into a frame
 - Fixed size, all pages/frames same size



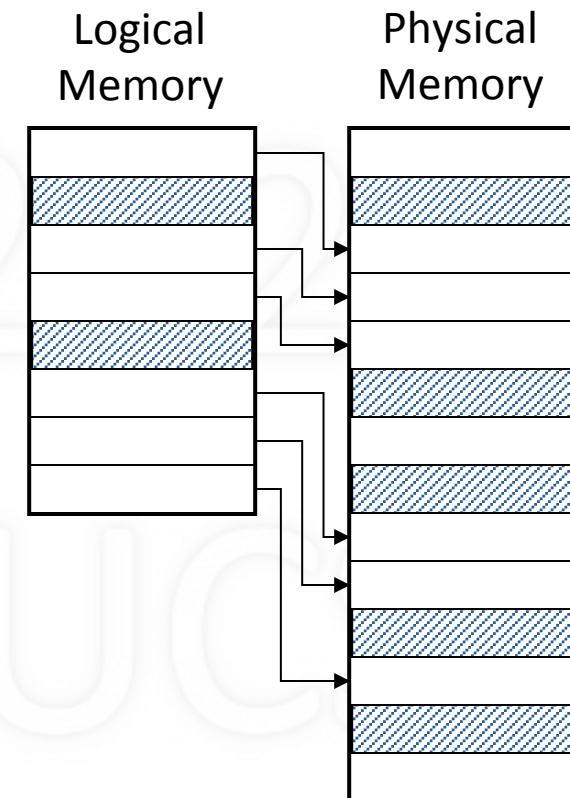
Page-based Logical Addressing

- Form of logical address: (p, i)
 - p is page number, 0 to $N_L - 1$
 - i is offset within page
 - Note: i is less than page size
 - no need to check
- Size of logical address space
 - N_L = max number of pages
 - $N_L \times$ page size = size of logical address space

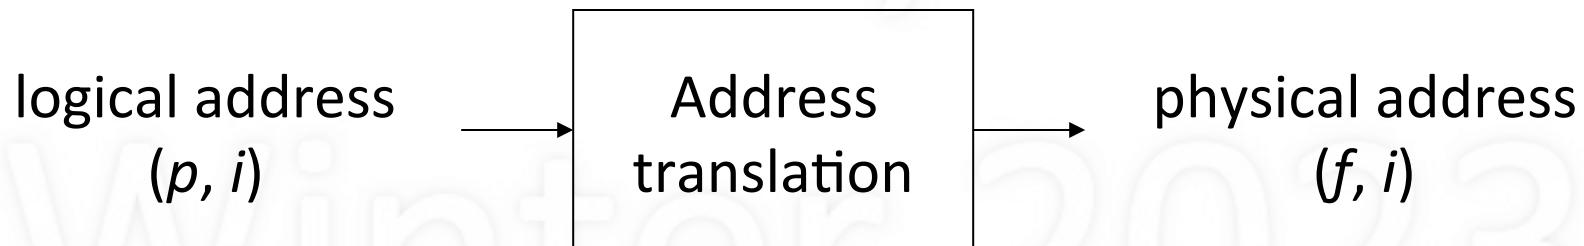


Frame-based Physical Addressing

- Form of physical address: (f, i)
 - f is frame number, 0 to $N_p - 1$
 - i is offset within frame
 - Note: i is less than frame size
 - since page size = frame size
- Size of physical address space
 - N_p = max number of frames
 - $N_p \times$ frame size = size of physical address space



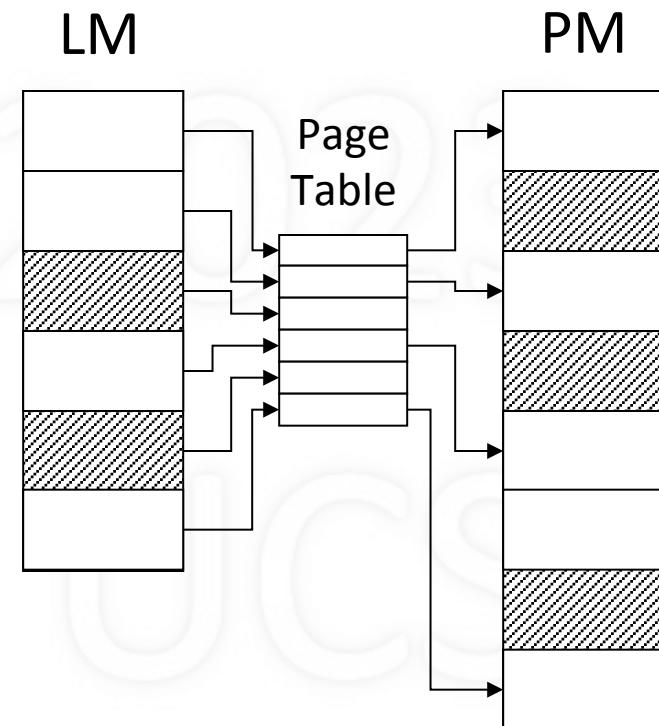
Page-based Address Translation



- Problem: how to translate
 - logical address (p, i) into physical address (f, i) ?
- Solution: use a page (translation) table PT
 - to translate page p into frame $f = PT(p)$
 - then concatenate f and i
- Summary: physical address = $PT(p) \ | \ | \ i$

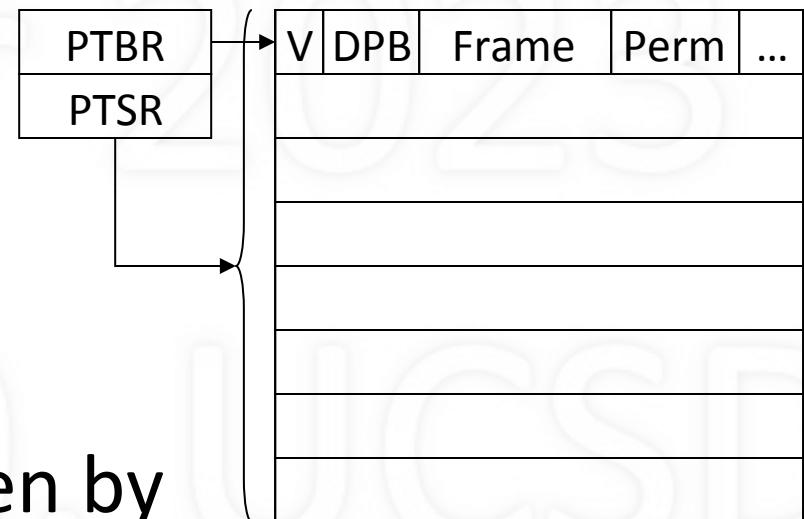
Logical Pages to Physical Frames

- Each page of logical memory corresponds to entry in page table
- Page table “maps” logical page into frame of physical memory



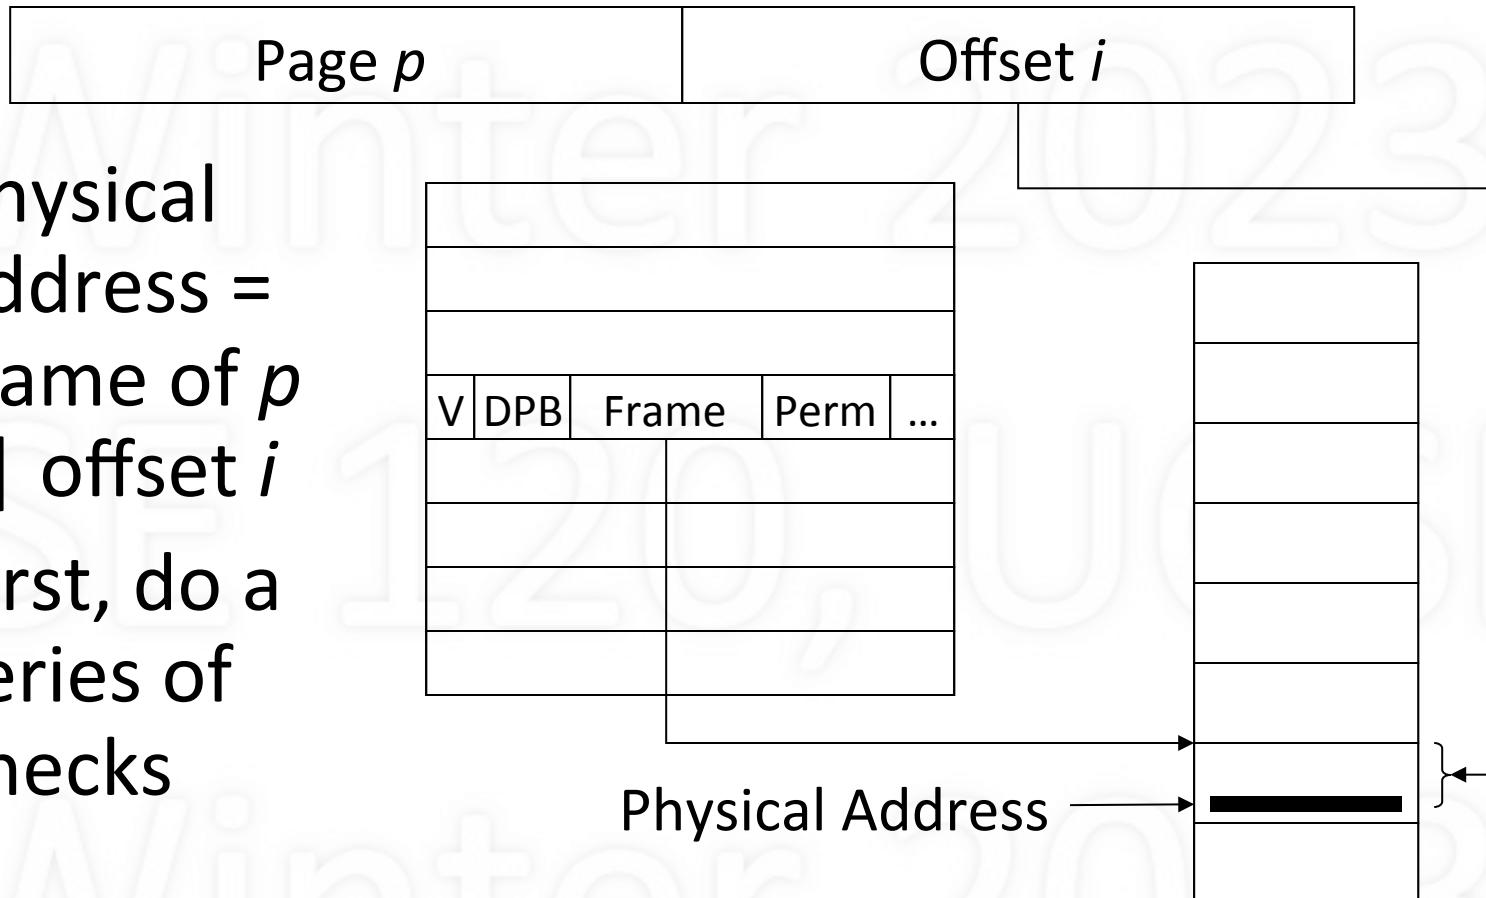
Page Table

- One per process (typically)
- Table entry elements
 - V: valid bit
 - Demand paging bits
 - Frame: page location
- Location in memory given by
 - Page table base register (hardware)
 - Page table size register (hardware)



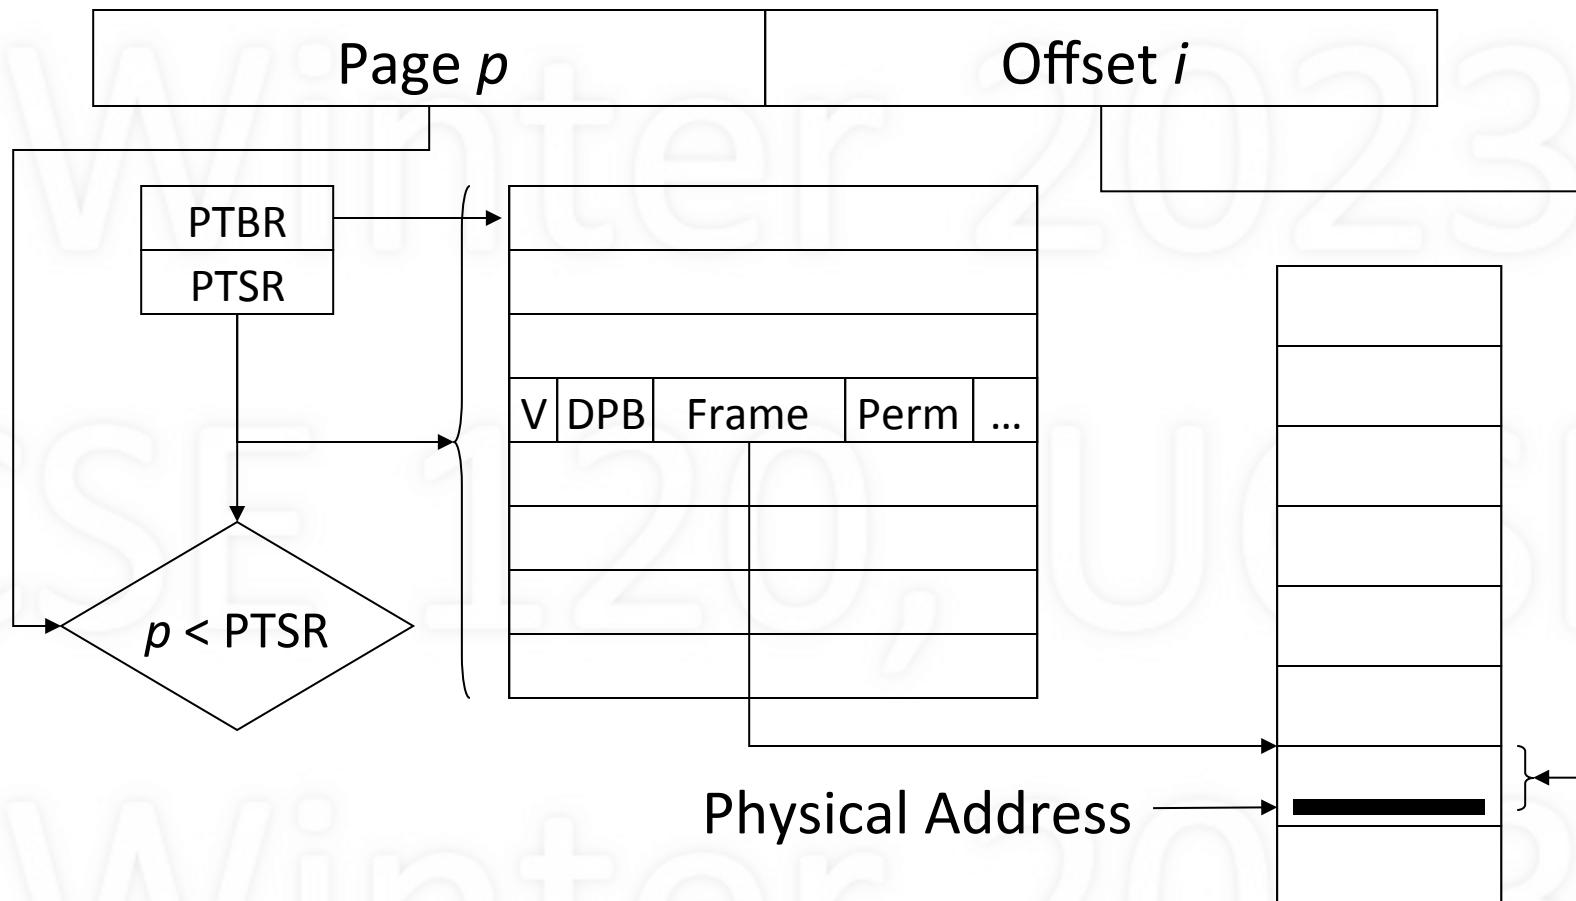
Address Translation

Logical Address

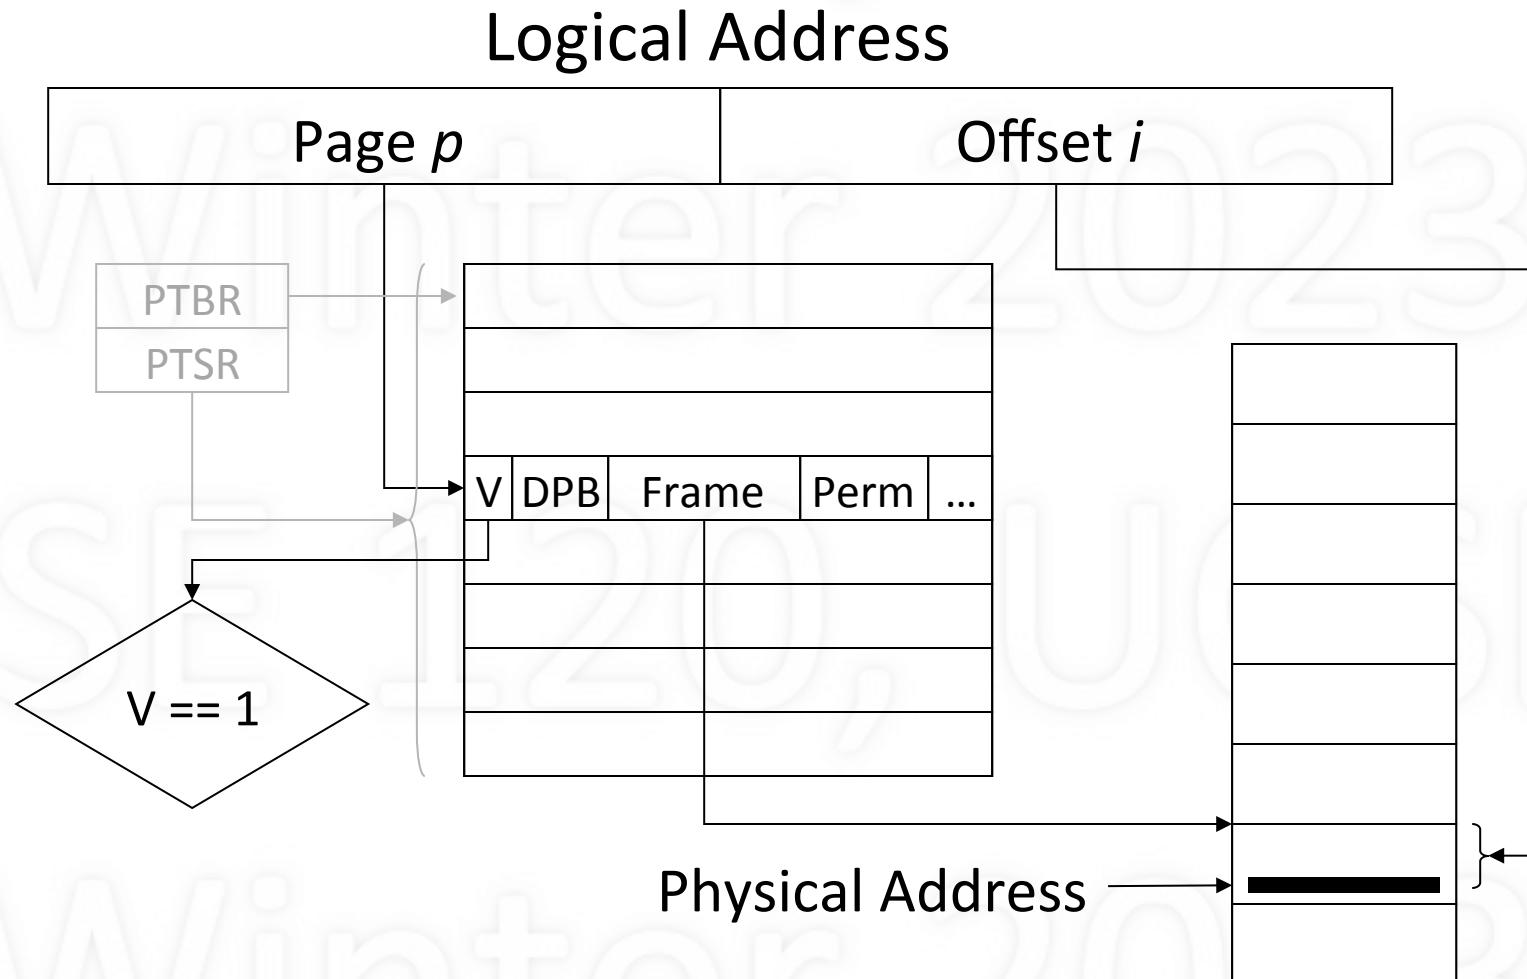


Check if Page p is Within Range

Logical Address

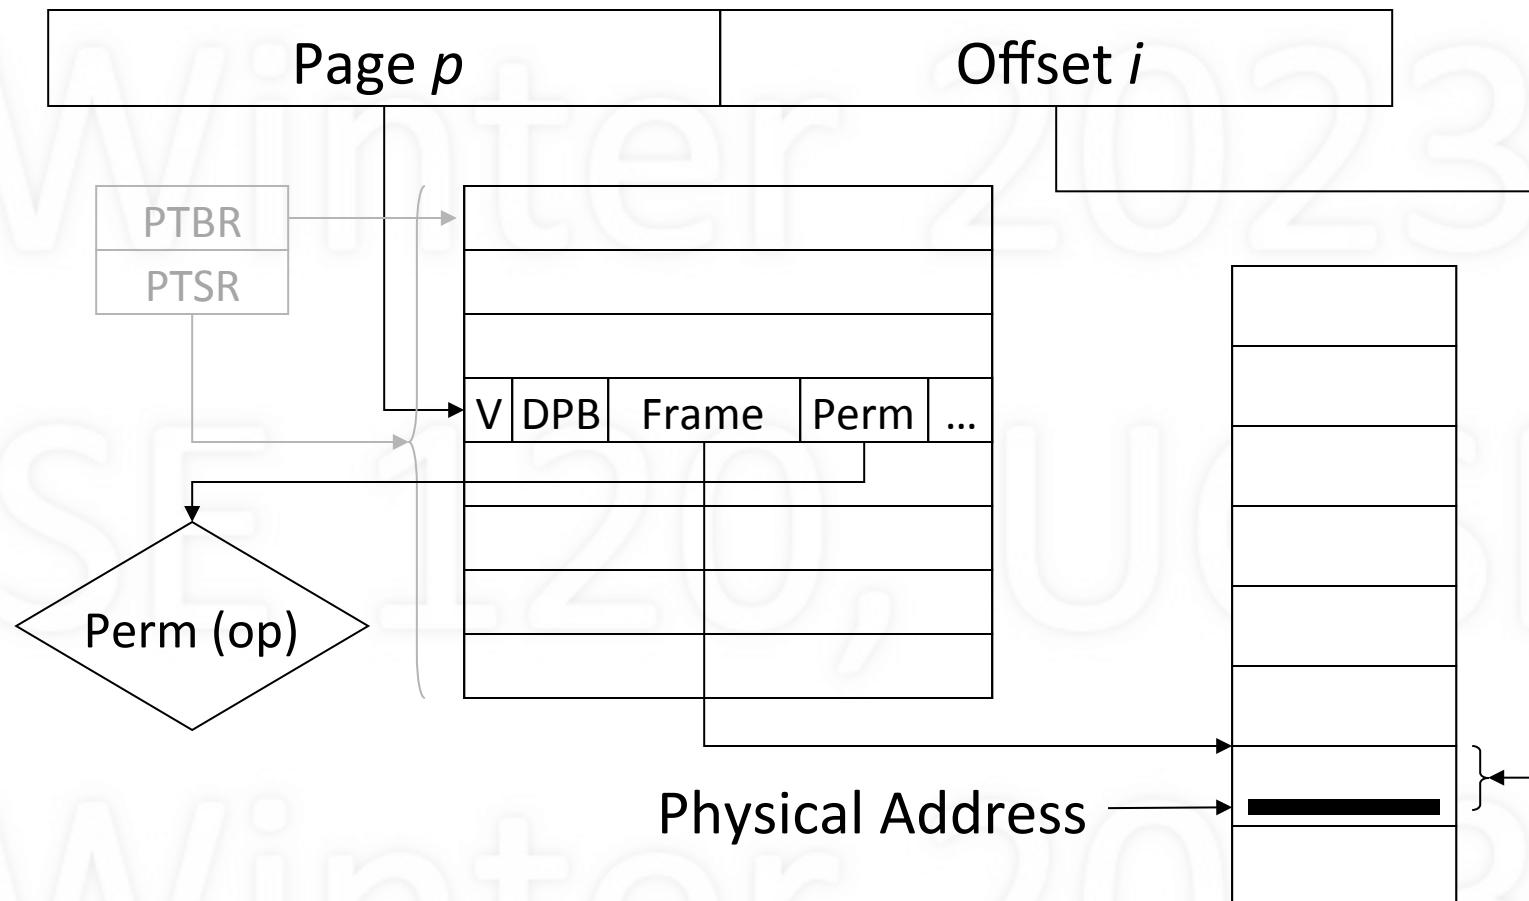


Check if Page Table Entry p is Valid



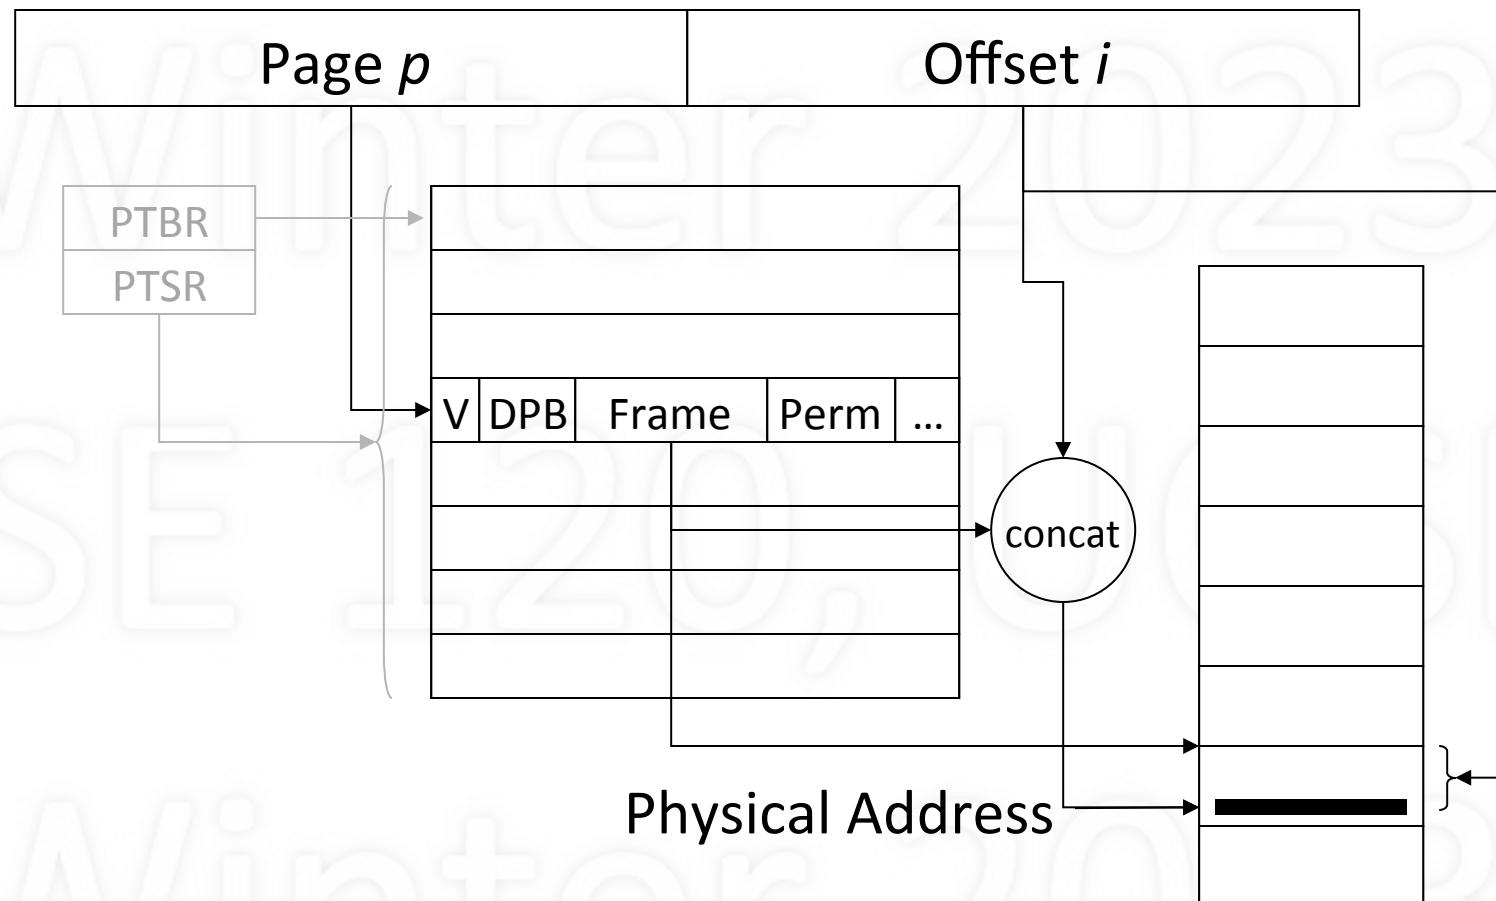
Check if Operation is Permitted

Logical Address

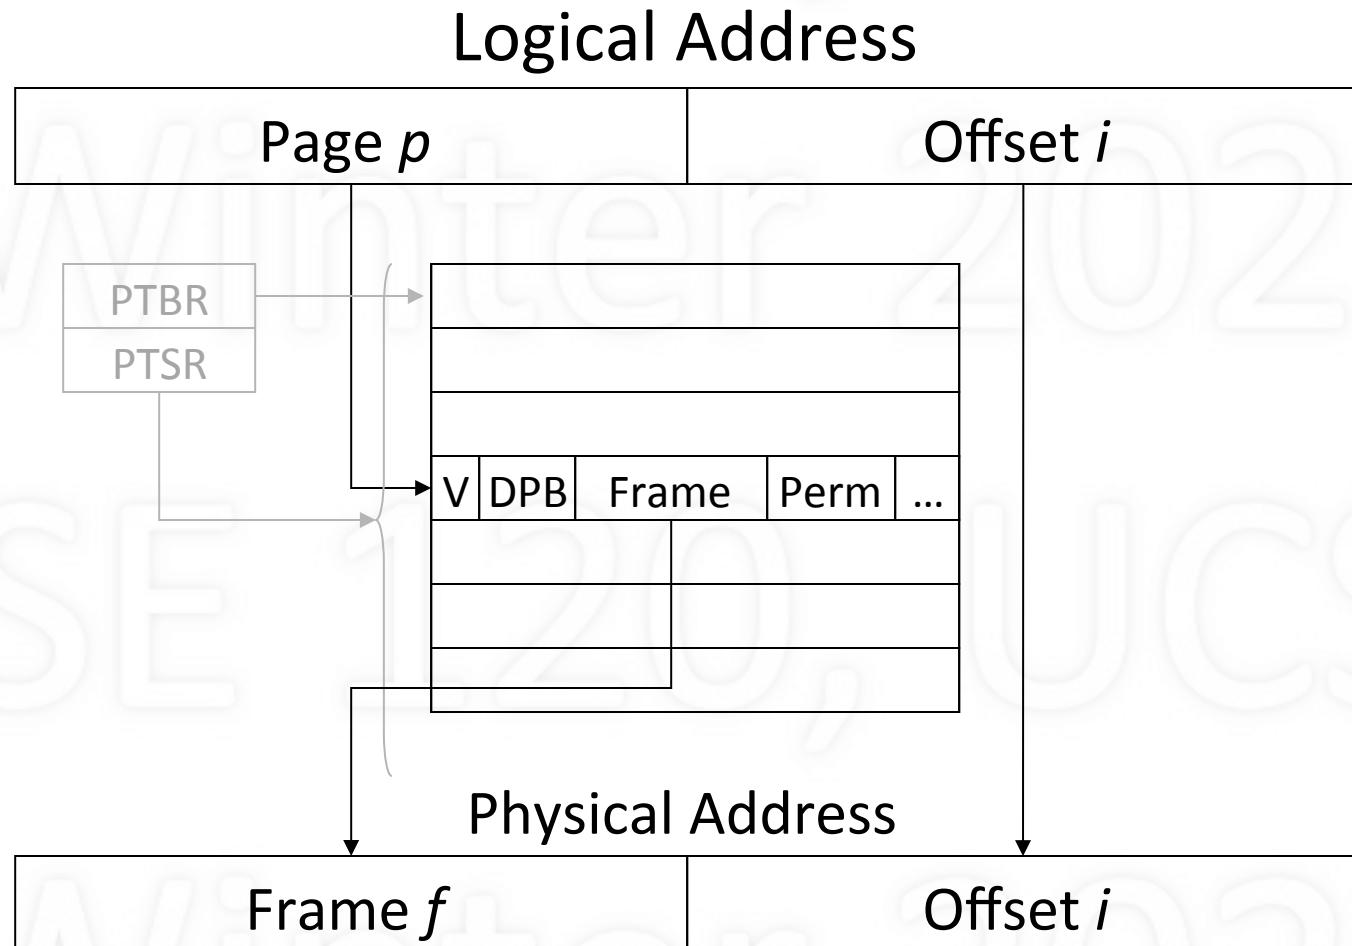


Translate Address

Logical Address

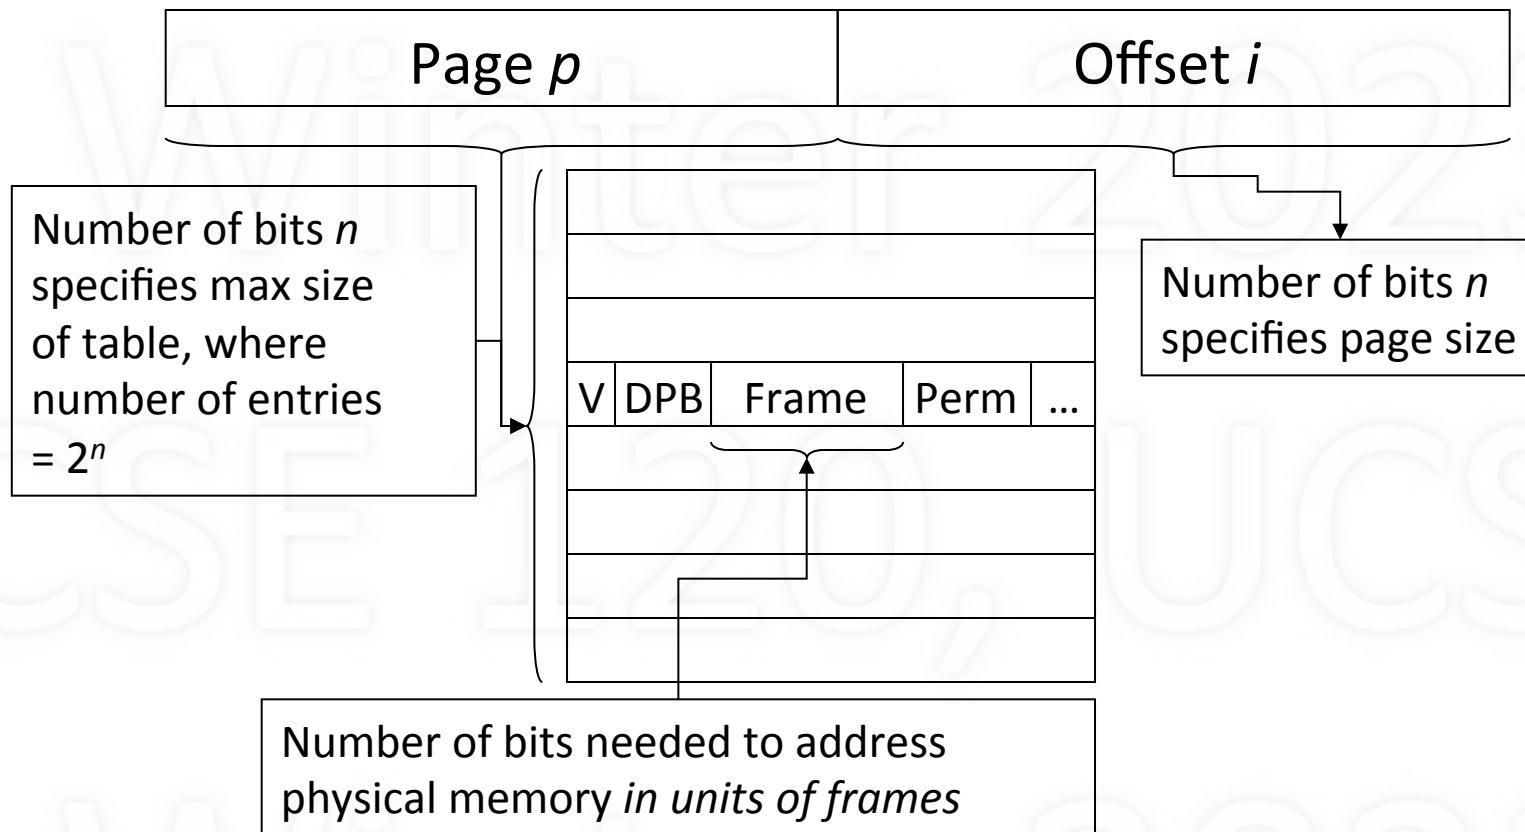


Physical Address by Concatenation

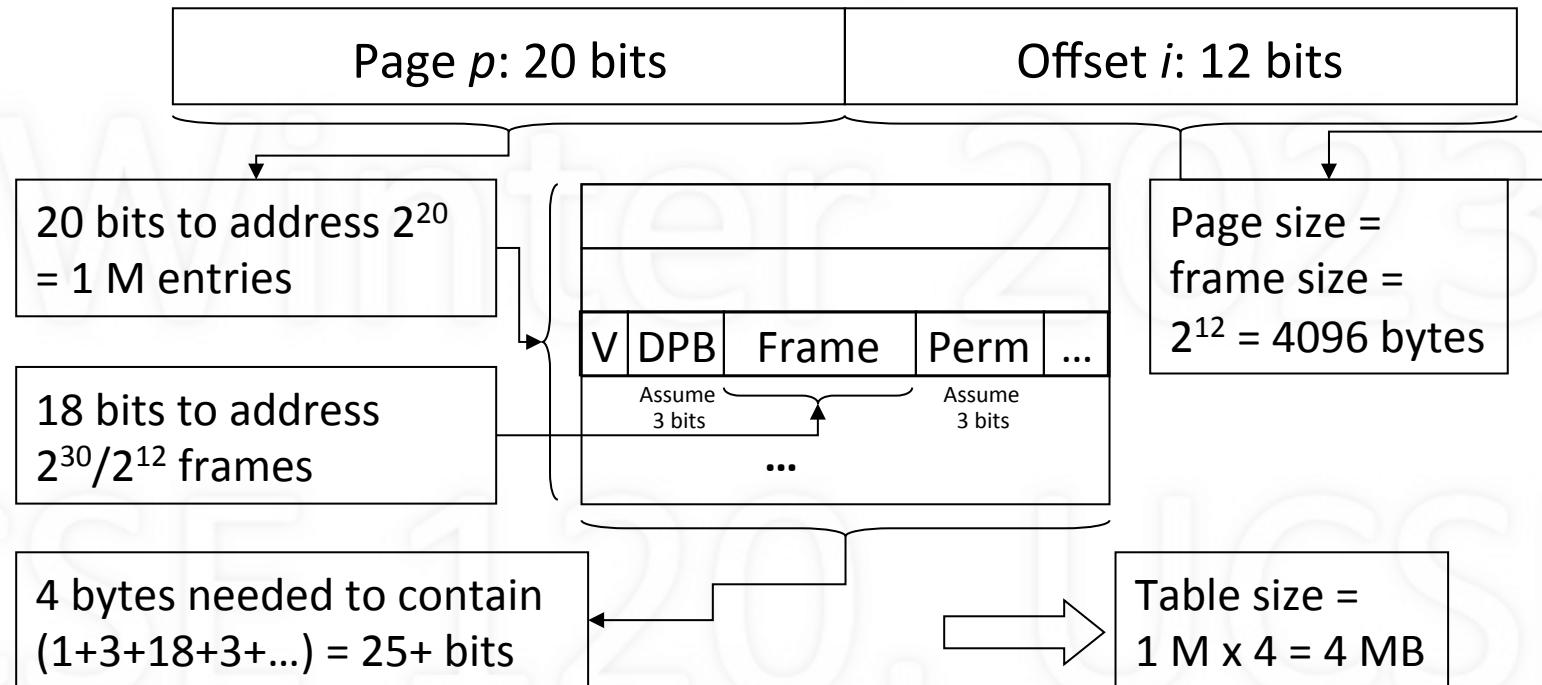


Sizing the Page Table

Logical Address



Example of Sizing the Page Table



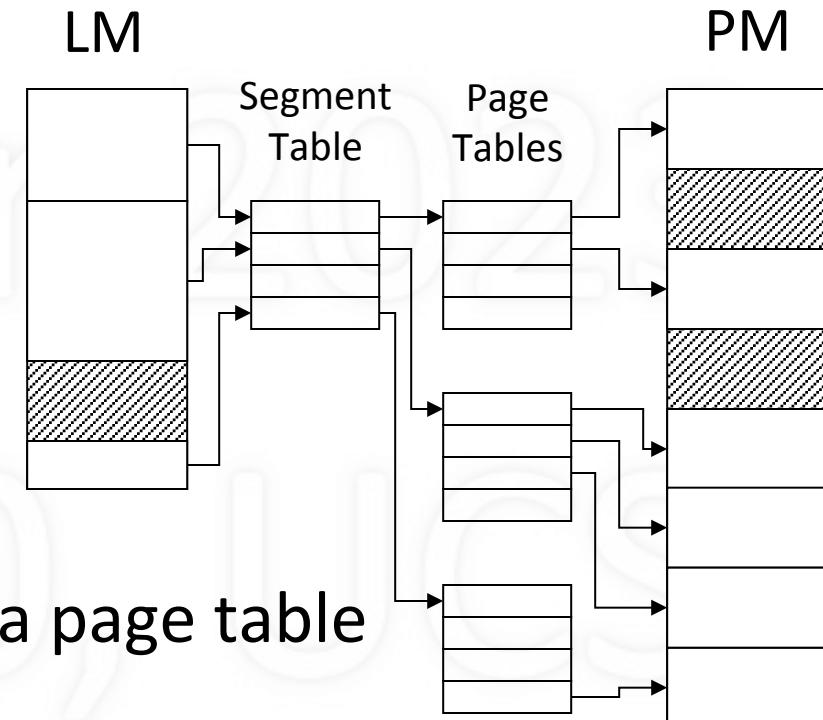
- Given 32 bit logical, 1 GB physical memory (max)
 - 20 bit page number, 12 bit offset

Segments vs. Pages

- Segment is good “logical” unit of information
 - Can be sized to fit any contents
 - Makes sense to share (e.g., code, data)
 - Can be protected according to contents
- Page is good “physical” unit of information
 - Simple memory management
- Why not have best of both
 - Segments, each a set of pages

Combining Segments and Pages

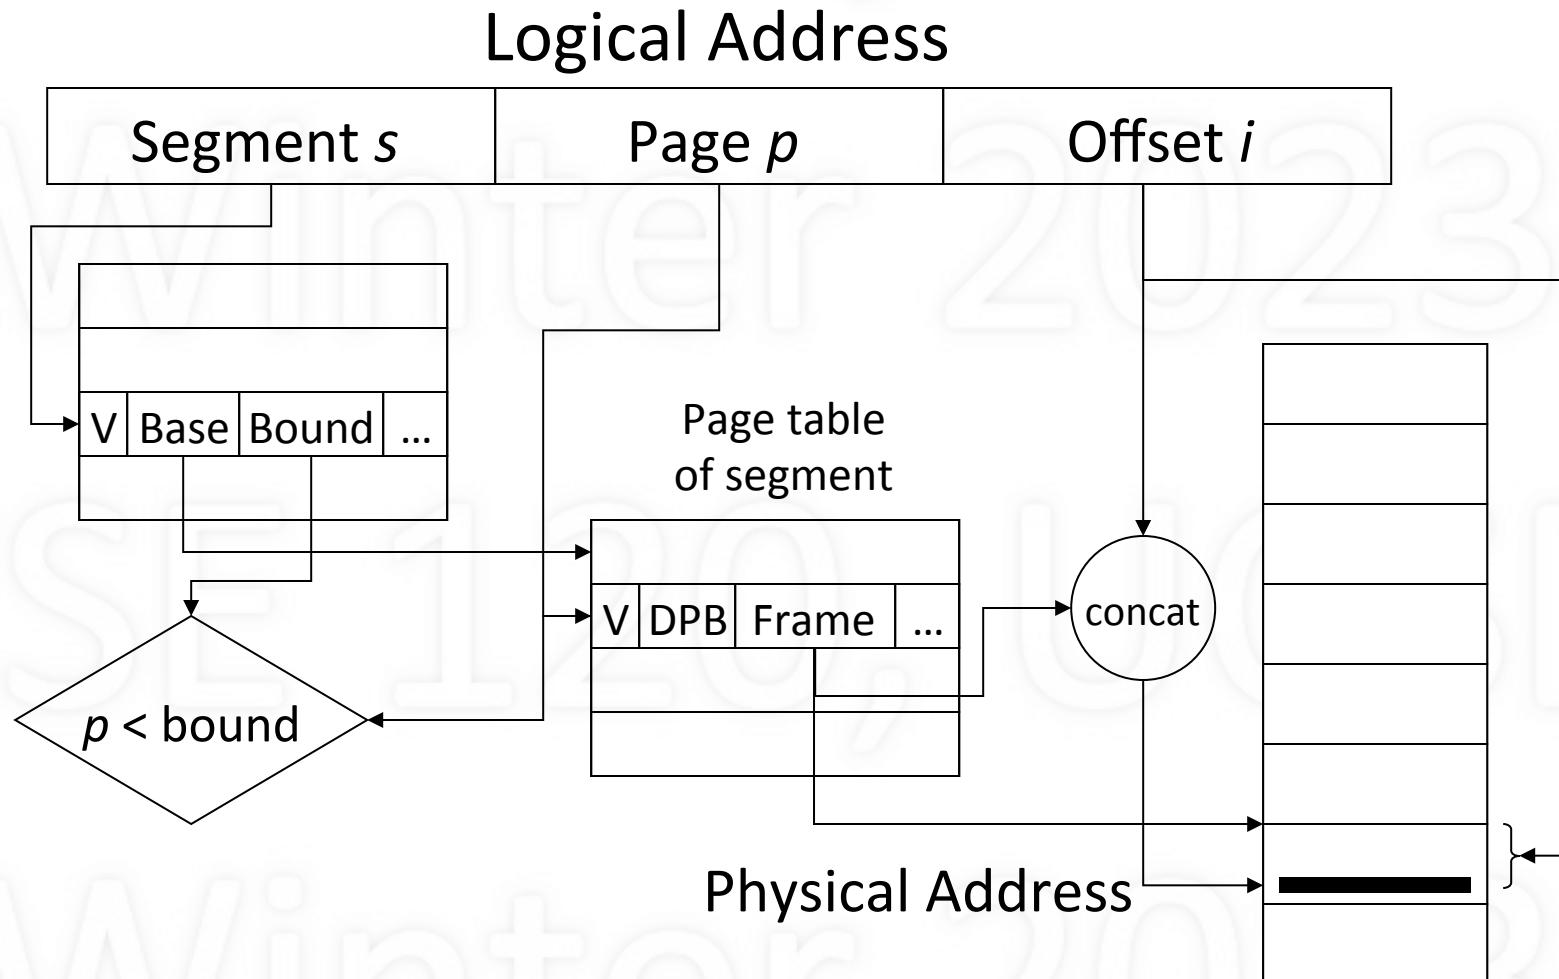
- Logical memory
 - composed of segments
- Each segment
 - composed of pages
- Segment table
 - Maps each segment to a page table
- Page tables
 - Maps each page to physical page frames



Address Translation

- Logical address: [segment s , page p , offset i]
- Do various checks
 - $s < \text{STSR}$, $V == 1$, $p < \text{bound}$, perm (op)
 - May get a segmentation violation
- Use s to index segment table to get page table
- Use p to index page table to get frame f
- Physical address = *concatenate* (f, i)

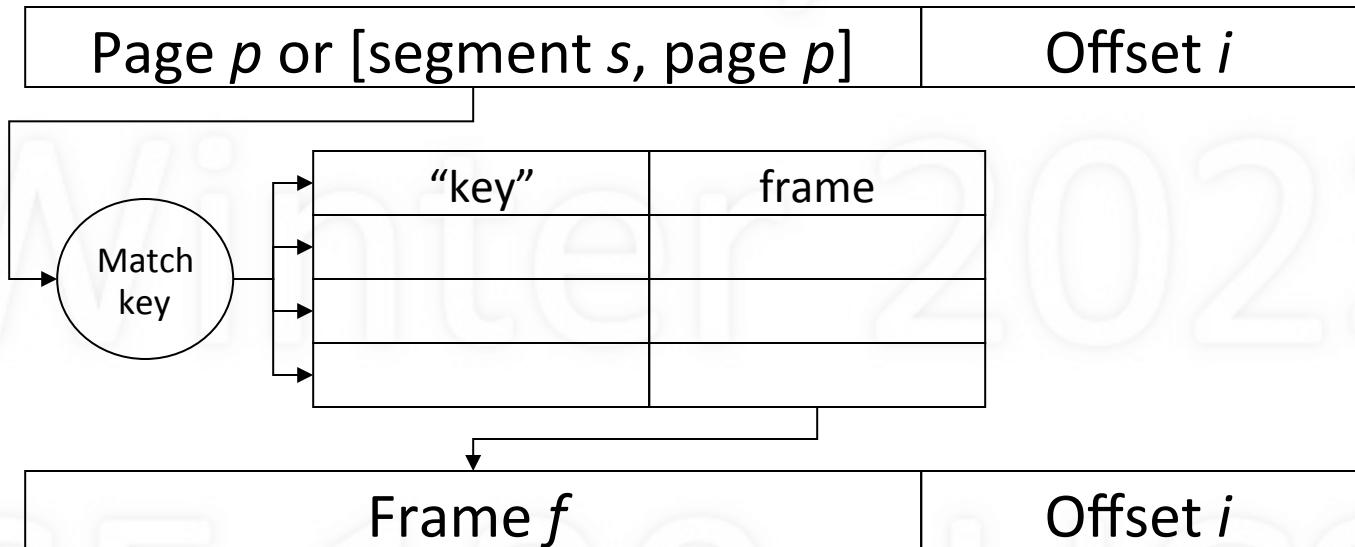
Segment/Page Address Translation



Cost of Translation

- Each lookup costs another memory reference
 - For each reference, additional references required
 - Slows machine down by factor of 2 or more
- Take advantage of locality of reference
 - Most references are to a small number of pages
 - Keep translations of these in high-speed memory
- Problem: don't know which pages till accessed

TLB: Translation Look-aside Buffer



- Fast memory keeps most recent translations
- If key matches, get frame number
- else wait for normal translation (in parallel)

Translation Cost with TLB

- Cost is determined by
 - Speed of memory: ~ 100 nsec
 - Speed of TLB: ~ 5 nsec
 - Hit ratio: fraction of refs satisfied by TLB, ~99%
- Speed with no address translation: 100 nsec
- Speed with address translation
 - TLB miss: 200 nsec (100% slowdown)
 - TLB hit: 105 nsec (5% slowdown)
 - Average: $105 \times 0.99 + 200 \times 0.01 \approx 106$ nsec

TLB Design Issues

- The larger the TLB
 - the higher the hit rate
 - the slower the response
 - the greater the expense
- TLB has a major effect on performance!
 - Must be flushed on context switches
 - Alternative: tagging entries with PIDs

Summary

- Logical memory
 - The Addressing Problem
 - The Protection Problem
- Organization of logical address space
 - Segmented
 - Paged
- Logical-to-physical address translation
- High cost of translation: locality, TLB

Textbook

- OSP: Chapter 8
- OSC: Chapter 9 (Main Memory)
 - Lecture-related: 9.1-9.3, 9.8
 - Recommended: 9.4, 9.6-9.7